

1

Writing Basic SQL Statements

Copyright © Oracle Corporation, 1998. All rights reserved.

ORACLE®

Schedule:	Timing	Topic
	40 minutes	Lecture
	25 minutes	Practice
	65 minutes	Total

Objectives

At the end of this lesson, you should be able to:

- **List the capabilities of SQL SELECT statements**
- **Execute a basic SELECT statement**
- **Differentiate between SQL statements and SQL*Plus commands**

Lesson Aim

To extract data from the database you need to use the structured query language (SQL) SELECT statement. You may need to restrict the columns that are displayed. This lesson describes all the SQL statements that you need to perform these actions.

You may want to create SELECT statements that can be used time and time again. This lesson also covers the use of SQL*Plus commands to execute SQL statements.

Capabilities of SQL SELECT Statements

Selection

Table 1

Projection

Table 1

Join

Table 1



Table 2

Basic SELECT Statement

```
SELECT    [DISTINCT] {*, column [alias],...}  
FROM      table;
```

- **SELECT** identifies *what* columns
- **FROM** identifies *which* table

Writing SQL Statements

- **SQL statements are not case sensitive.**
- **SQL statements can be on one or more lines.**
- **Keywords cannot be abbreviated or split across lines.**
- **Clauses are usually placed on separate lines.**
- **Tabs and indents are used to enhance readability.**

Writing SQL Statements

By following simple rules and guidelines given below, you can construct valid statements that are both easy to read and easy to edit:

- SQL statements are not case sensitive, unless indicated.
- SQL statements can be entered on one or many lines.
- Keywords cannot be split across lines or abbreviated.
- Clauses are usually placed on separate lines for readability and ease of editing.
- Tabs and indents can be used to make code more readable.
- Keywords typically are entered in uppercase; all other words, such as table names and columns, are entered in lowercase.
- Within SQL*Plus, a SQL statement is entered at the SQL prompt, and the subsequent lines are numbered. This is called the *SQL buffer*. Only one statement can be current at any time within the buffer.

Executing SQL Statements

- Place a semicolon (;) at the end of last clause.
- Place a slash on the last line in the buffer.
- Place a slash at the SQL prompt.
- Issue a SQL*Plus RUN command at the SQL prompt.

Selecting All Columns

```
SQL> SELECT *  
2 FROM dept;
```

DEPTNO	DNAME	LOC
10	ACCOUNTING	NEW YORK
20	RESEARCH	DALLAS
30	SALES	CHICAGO
40	OPERATIONS	BOSTON

Selecting Specific Columns

```
SQL> SELECT deptno, loc  
2 FROM dept;
```

DEPTNO	LOC
10	NEW YORK
20	DALLAS
30	CHICAGO
40	BOSTON

Column Label Defaults

- **Default justification**
 - **Left: Date and character data**
 - **Right: Numeric data**
- **Default display: Uppercase**

1-8

Copyright © Oracle Corporation, 1998. All rights reserved.

ORACLE®

Column Heading Defaults

Character column heading and data as well as date column heading and data are left-justified within a column width. Number headings and data are right-justified.

```
SQL> SELECT ename, hiredate, sal  
2 FROM emp;
```

ENAME	HIREDATE	SAL
KING	17-NOV-81	5000
BLAKE	01-MAY-81	2850
CLARK	09-JUN-81	2450
JONES	02-APR-81	2975
MARTIN	28-SEP-81	1250
ALLEN	20-FEB-81	1600

Character and date column headings can be truncated, but number headings cannot be truncated. The column labels appear in uppercase by default. You can override the column label display with an alias. Column aliases are covered later in this lesson.

Arithmetic Expressions

Create expressions on NUMBER and DATE data types by using arithmetic operators.

Operator	Description
+	Add
-	Subtract
*	Multiply
/	Divide

Using Arithmetic Operators

```
SQL> SELECT ename, sal, sal+300  
2 FROM emp;
```

ENAME	SAL	SAL+300
-----	-----	-----
KING	5000	5300
BLAKE	2850	3150
CLARK	2450	2750
JONES	2975	3275
MARTIN	1250	1550
ALLEN	1600	1900
...		

14 rows selected.

Operator Precedence

*	/	+	—
----------	----------	----------	----------

- **Multiplication and division take priority over addition and subtraction.**
- **Operators of the same priority are evaluated from left to right.**
- **Parentheses are used to force prioritized evaluation and to clarify statements.**

Operator Precedence

```
SQL> SELECT ename, sal, 12*sal+100  
2 FROM emp;
```

ENAME	SAL	12*SAL+100
-----	-----	-----
KING	5000	60100
BLAKE	2850	34300
CLARK	2450	29500
JONES	2975	35800
MARTIN	1250	15100
ALLEN	1600	19300
...		

14 rows selected.

Using Parentheses

```
SQL> SELECT ename, sal, 12*(sal+100)
2 FROM emp;
```

ENAME	SAL	12*(SAL+100)
KING	5000	61200
BLAKE	2850	35400
CLARK	2450	30600
JONES	2975	36900
MARTIN	1250	16200
...		

14 rows selected.

Defining a Null Value

- A null is a value that is unavailable, unassigned, unknown, or inapplicable.
- A null is not the same as zero or a blank space.

```
SQL> SELECT  ename, job, comm  
2 FROM      emp;
```

ENAME	JOB	COMM
-----	-----	-----
KING	PRESIDENT	
BLAKE	MANAGER	
...		
TURNER	SALESMAN	0
...		
14 rows selected.		

Null Values in Arithmetic Expressions

Arithmetic expressions containing a null value evaluate to null.

```
SQL> select  ename NAME, 12*sal+comm  
2   from    emp  
3  WHERE    ename='KING' ;
```

NAME	12*SAL+COMM
-----	-----
KING	

Defining a Column Alias

- **Renames a column heading**
- **Is useful with calculations**
- **Immediately follows column name; optional AS keyword between column name and alias**
- **Requires double quotation marks if it contains spaces or special characters or is case sensitive**

Column Aliases

When displaying the result of a query, SQL*Plus normally uses the name of the selected column as the column heading. In many cases, this heading may not be descriptive and hence is difficult to understand. You can change a column heading by using a column alias.

Specify the alias after the column in the SELECT list using a space as a separator. By default, alias headings appear in uppercase. If the alias contains spaces, special characters (such as # or \$), or is case sensitive, enclose the alias in double quotation marks (“ ”).

Class Management Note

Within a SQL statement, a column alias can be used in both the SELECT clause and the ORDER BY clause. You cannot use column aliases in the WHERE clause. Both alias features comply with the ANSI SQL 92 standard.

Demo: *l1alias.sql*

Using Column Aliases

```
SQL> SELECT ename AS name, sal salary
2 FROM emp;
```

```
NAME                SALARY
-----
...
```

```
SQL> SELECT ename "Name",
2           sal*12 "Annual Salary"
3 FROM emp;
```

```
Name                Annual Salary
-----
...
```

Concatenation Operator

- **Concatenates columns or character strings to other columns**
- **Is represented by two vertical bars (||)**
- **Creates a resultant column that is a character expression**

Concatenation Operator

You can link columns to other columns, arithmetic expressions, or constant values to create a character expression by using the concatenation operator (||). Columns on either side of the operator are combined to make a single output column.

Using the Concatenation Operator

```
SQL> SELECT  ename||job AS "Employees"  
2  FROM      emp;
```

```
Employees  
-----  
KINGPRESIDENT  
BLAKEMANAGER  
CLARKMANAGER  
JONESMANAGER  
MARTINSALESMAN  
ALLENSALESMAN  
...  
14 rows selected.
```

Literal Character Strings

- A literal is a character, expression, or number included in the **SELECT** list.
- Date and character literal values must be enclosed within single quotation marks.
- Each character string is output once for each row returned.

Literal Character Strings

A literal is any character, expression, or number included in the **SELECT** list that is not a column name or a column alias. It is printed for each row returned. Literal strings of free-format text can be included in the query result and are treated the same as a column in the **SELECT** list.

Date and character literals *must* be enclosed within single quotation marks (' '); number literals must not.

Using Literal Character Strings

```
SQL> SELECT ename || ' ' || 'is a' || ' ' || job  
2          AS "Employee Details"  
3 FROM      emp;
```

```
Employee Details  
-----  
KING is a PRESIDENT  
BLAKE is a MANAGER  
CLARK is a MANAGER  
JONES is a MANAGER  
MARTIN is a SALESMAN  
...  
14 rows selected.
```

Duplicate Rows

The default display of queries is all rows, including duplicate rows.

```
SQL> SELECT deptno  
2 FROM emp;
```

```
DEPTNO  
-----  
      10  
      30  
      10  
      20  
...  
14 rows selected.
```

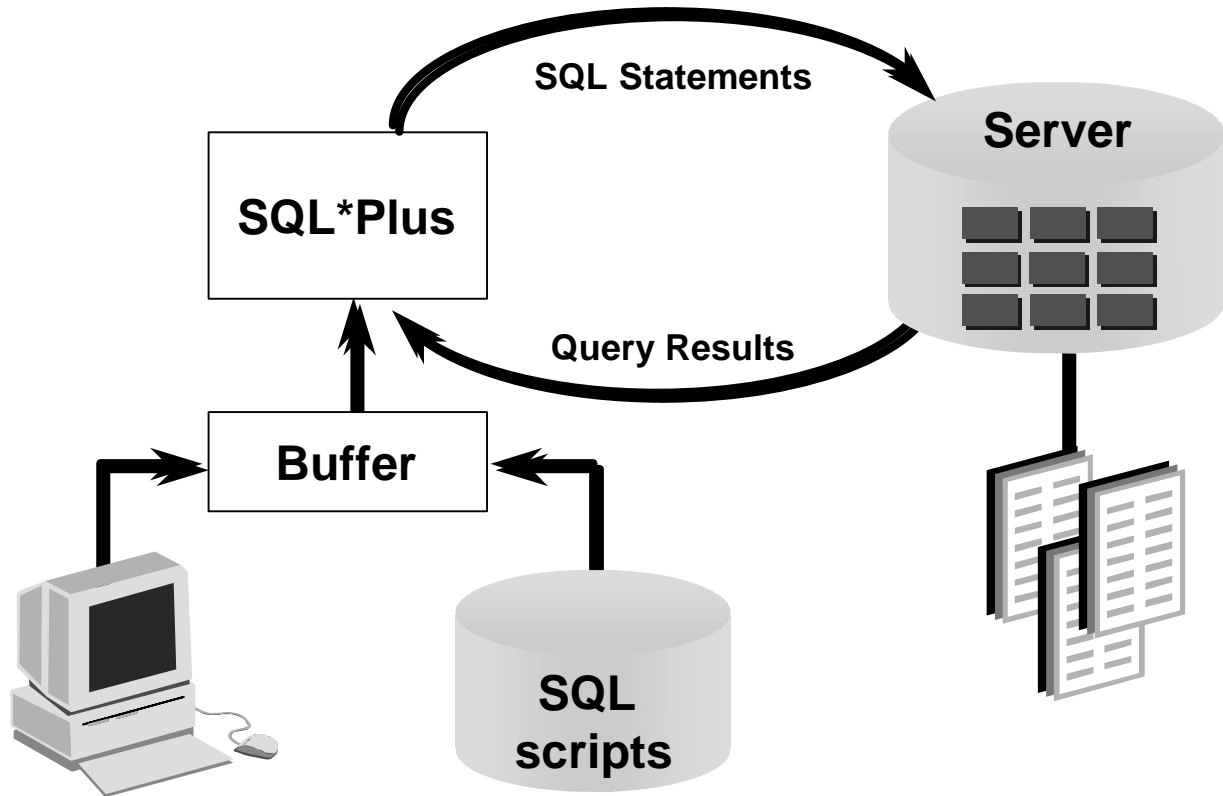
Eliminating Duplicate Rows

Eliminate duplicate rows by using the **DISTINCT** keyword in the **SELECT** clause.

```
SQL> SELECT DISTINCT deptno  
2 FROM emp;
```

DEPTNO
10
20
30

SQL and SQL*Plus Interaction



1-24

Copyright © Oracle Corporation, 1998. All rights reserved.

ORACLE®

SQL Statements Versus SQL*Plus Commands

SQL

- A language
- ANSI standard
- Keyword cannot be abbreviated
- Statements manipulate data and table definitions in the database

SQL
statements



SQL
buffer



SQL*Plus

- An environment
- Oracle proprietary
- Keywords can be abbreviated
- Commands do not allow manipulation of values in the database

SQL*Plus
commands



SQL*Plus
buffer



Overview of SQL*Plus

- **Log in to SQL*Plus.**
- **Describe the table structure.**
- **Edit your SQL statement.**
- **Execute SQL from SQL*Plus.**
- **Save SQL statements to files and append SQL statements to files.**
- **Execute saved files.**
- **Load commands from file to buffer to edit.**

1-26

Copyright © Oracle Corporation, 1998. All rights reserved.

ORACLE®

SQL*Plus

SQL*Plus is an environment in which you can do the following:

- Execute SQL statements to retrieve, modify, add, and remove data from the database
- Format, perform calculations on, store, and print query results in the form of reports
- Create script files to store SQL statements for repetitive use in the future

SQL*Plus commands can be divided into the following main categories:

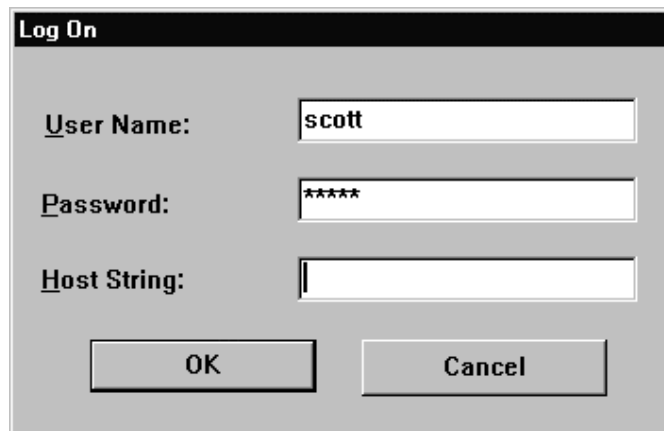
Category	Purpose
Environment	Affects the general behavior of SQL statements for the session
Format	Formats query results
File manipulation	Saves, loads, and runs script files
Execution	Sends SQL statements from SQL buffer to Oracle8 Server
Edit	Modifies SQL statements in the buffer
Interaction	Allows you to create and pass variables to SQL statements, print SQL statements, and print messages to the screen
Miscellaneous	Has various commands to connect to the database, manipulate the SQL*Plus environment, and display column definitions

Class Management Note (For page 1-27)

Snippet: Establishing a Database Session

Logging In to SQL*Plus

- From Windows environment:

A screenshot of the 'Log On' dialog box in a Windows environment. The dialog box has a title bar that says 'Log On'. It contains three input fields: 'User Name:' with the text 'scott', 'Password:' with six asterisks '*****', and 'Host String:' which is empty. At the bottom of the dialog box are two buttons: 'OK' and 'Cancel'.

- From command line:

```
sqlplus [username[/password  
[ @database]]]
```

Displaying Table Structure

Use the SQL*Plus DESCRIBE command to display the structure of a table.

```
DESC[RIBE] tablename
```

Displaying Table Structure

```
SQL> DESCRIBE dept
```

Name	Null?	Type
-----	-----	----
DEPTNO	NOT NULL	NUMBER(2)
DNAME		VARCHAR2(14)
LOC		VARCHAR2(13)

SQL*Plus Editing Commands

- **A[PPEND] *text***
- **C[HANGE] / *old* / *new***
- **C[HANGE] / *text* /**
- **CL[EAR] BUFF[ER]**
- **DEL**
- **DEL *n***
- **DEL *m n***

1-30

Copyright © Oracle Corporation, 1998. All rights reserved.

ORACLE®

SQL*Plus Editing Commands

SQL*Plus commands are entered one line at a time and are not stored in the SQL buffer.

Command	Description
A[PPEND] <i>text</i>	Adds text to the end of the current line
C[HANGE] / <i>old</i> / <i>new</i>	Changes <i>old</i> text to <i>new</i> in the current line
C[HANGE] / <i>text</i> /	Deletes <i>text</i> from the current line
CL[EAR] BUFF[ER]	Deletes all lines from the SQL buffer
DEL	Deletes current line

- If you press [Return] before completing a command, SQL*Plus prompts you with a line number.
- You terminate the SQL buffer by either entering one of the terminator characters (semicolon or slash) or by pressing [Return] twice. You then see the SQL prompt.

SQL*Plus Editing Commands

- I[NPUT]
- I[NPUT] *text*
- L[IST]
- L[IST] *n*
- L[IST] *m n*
- R[UN]
- *n*
- *n text*
- **0** *text*

1-31

Copyright © Oracle Corporation, 1998. All rights reserved.

ORACLE®

SQL*Plus Editing Commands (continued)

Command	Description
I[NPU T]	Inserts an indefinite number of lines
I[NPUT] <i>text</i>	Inserts a line consisting of <i>text</i>
L[IST]	Lists all lines in the SQL buffer
L[IST] <i>n</i>	Lists one line (specified by <i>n</i>)
L[IST] <i>m n</i>	Lists a range of lines (<i>m</i> to <i>n</i>)
R[UN]	Displays and runs the current SQL statement in the buffer
<i>n</i>	Specifies the line to make the current line
<i>n text</i>	Replaces line <i>n</i> with <i>text</i>
0 <i>text</i>	Inserts a line before line 1

You can enter only one SQL*Plus command per SQL prompt. SQL*Plus commands are not stored in the buffer. To continue a SQL*Plus command on the next line, end the current line with a hyphen (-).



Class Management Note

Show students the use of the commonly used editing commands, like A[PPEND], C[HANGE], DEL, L[IST], and R[UN].

SQL*Plus File Commands

- **SAVE *filename***
- **GET *filename***
- **START *filename***
- **@ *filename***
- **EDIT *filename***
- **SPOOL *filename***
- **EXIT**

SQL*Plus File Commands

SQL statements communicate with the Oracle Server. SQL*Plus commands control the environment, format query results, and manage files. You can use the commands identified in the following table:

Command	Description
SAV[E] <i>filename</i> [.ext] [REP[LACE]APP[END]]	Saves current contents of SQL buffer to a file. Use APPEND to add to an existing file; use REPLACE to overwrite an existing file. The default extension is .sql.
GET <i>filename</i> [.ext]	Writes the contents of a previously saved file to the SQL buffer. The default extension for the filename is .sql.
STA[RT] <i>filename</i> [.ext]	Runs a previously saved command file.
@ <i>filename</i>	Runs a previously saved command file (same as START).
ED[IT]	Invokes the editor and saves the buffer contents to a file named <i>afiedt.buf</i> .
ED[IT] [<i>filename</i> [.ext]]	Invokes editor to edit contents of a saved file.
SPO[OL] [<i>filename</i> [.ext]] OFF[OUT]	Stores query results in a file. OFF closes the spool file. OUT closes the spool file and sends the file results to the system printer.
EXIT	Leaves SQL*Plus.

Summary

```
SELECT    [DISTINCT] {*,column[alias],...}  
FROM      table;
```

Use SQL*Plus as an environment to:

- **Execute SQL statements**
- **Edit SQL statements**

Practice Overview

- **Selecting all data from different tables.**
- **Describing the structure of tables.**
- **Performing arithmetic calculations and specifying column names.**
- **Using SQL*Plus editor.**

Practice Overview

This is the first of many practices. The solutions (if you require them) can be found in Appendix A. Practices are intended to introduce all topics covered in the lesson. Questions 2-4 are paper-based.

In any practice, there may be “if you have time” or “if you want extra challenge” questions. Do these only if you have completed all other questions within the allocated time and would like a further challenge to your skills.

Take the practice slowly and precisely. You can experiment with saving and running command files. If you have any questions at any time, attract the instructor’s attention.



Paper-Based Questions

For questions 2-4 circle either True or False.

1

Restricting and Sorting Data

Copyright © Oracle Corporation, 1998. All rights reserved. ORACLE®

Schedule:	Timing	Topic
	45 minutes	Lecture
	30 minutes	Practice
	75 minutes	Total

Objectives

At the end of this lesson, you should be able to:

- **Limit the rows retrieved by a query**
- **Sort the rows retrieved by a query**

Lesson Aim


While retrieving data from the database, you may need to restrict the rows of data that are displayed or specify the order in which the rows are displayed. This lesson explains the SQL statements that you will use to perform these actions.

Limiting Rows Using a Selection

EMP

EMPNO	ENAME	JOB	...	DEPTNO
7839	KING	PRESIDENT		10
7698	BLAKE	MANAGER		30
7782	CLARK	MANAGER		10
7566	JONES	MANAGER		20
...				

**“...retrieve all
employees
in department 10”**



EMP

EMPNO	ENAME	JOB	...	DEPTNO
7839	KING	PRESIDENT		10
7782	CLARK	MANAGER		10
7934	MILLER	CLERK		10

Limiting Rows Using a Selection

In the above example, assume that you want to display all the employees in department 10. The highlighted set of rows with a value of 10 in DEPTNO column are the only ones returned. This method of restriction is the basis of the WHERE clause in SQL.

Limiting Rows Selected

- Restrict the rows returned by using the **WHERE** clause.

```
SELECT          [DISTINCT] {*, column [alias], ...}  
FROM            table  
[WHERE          condition(s)];
```

- The **WHERE** clause follows the **FROM** clause.

Limiting Rows Selected

You can restrict the rows returned from the query by using the WHERE clause. A WHERE clause contains a condition that must be met, and it directly follows the FROM clause.

In the syntax:

WHERE	restricts the query to rows that meet a condition
<i>condition</i>	is composed of column names, expressions, constants, and comparison operator

The WHERE clause can compare values in columns, literal values, arithmetic expressions, or functions. The WHERE clause consists of three elements:

- Column name
- Comparison operator
- Column name, constant, or list of values

Using the WHERE Clause

```
SQL> SELECT ename, job, deptno
      2 FROM emp
      3 WHERE job= 'CLERK';
```

ENAME	JOB	DEPTNO
JAMES	CLERK	30
SMITH	CLERK	20
ADAMS	CLERK	20
MILLER	CLERK	10

Using the WHERE clause

In the example, the SELECT statement retrieves the name, job title, and the department number of all employees whose job title is CLERK.

Note that the job title CLERK has been specified in uppercase to ensure that the match is made with the job column in the EMP table. Character strings are case sensitive.

Character Strings and Dates

- Character strings and date values are enclosed in single quotation marks
- Character values are case-sensitive and date values are format-sensitive
- Default date format is 'DD-MON-YY'

```
SQL> SELECT  ename, job, deptno
2  FROM      emp
3  WHERE     ename = 'JAMES';
```

Character Strings and Dates

Character strings and dates in the WHERE clause must be enclosed in single quotation marks (' '). Number constants, however, must not.

All character searches are case sensitive. In the following example, no rows are returned because the EMP table stores all the data in uppercase.

```
SQL> SELECT  ename, empno, job, deptno
2  FROM      emp
3  WHERE     job='clerk';
```

Oracle stores dates in an internal numeric format, representing the century, year, month, day, hours, minutes, and seconds. The default date display is DD-MON-YY.

Note: Changing default date format will be covered in lesson 3.

Number values are not enclosed within quotation marks.

Class Management Note

Some students may ask how to override the case sensitivity. Later in the course, we will cover the use of single-row functions such as UPPER and LOWER to override the case sensitivity.

Comparison Operators

Operator	Meaning
=	Equal to
>	Greater than
>=	Greater than or equal to
<	Less than
<=	Less than or equal to
<>	Not equal to

1-45

Copyright © Oracle Corporation, 1998. All rights reserved.

ORACLE®

Comparison Operators

Comparison operators are used in conditions that compare one expression to another. They are used in the WHERE clause in the following format.

Syntax

```
... WHERE expr operator value
```

Examples

```
... WHERE hiredate='01-JAN-95'
```

```
... WHERE sal>=1500
```

```
... WHERE ename='SMITH'
```

Class Management Note

Remind students that the *expr* cannot be an alias.

Using the Comparison Operators

```
SQL> SELECT ename, sal, comm
2 FROM emp
3 WHERE sal<=comm;
```

ENAME	SAL	COMM
MARTIN	1250	1400

Using the Comparison Operators

In the example, the SELECT statement retrieves name, salary, and commission from the EMP table, where the employee salary is less than or equal to their commission amount. Note that there is no explicit value supplied to the WHERE clause. The two values being compared are taken from the SAL and COMM columns in the EMP table.

Class Management Note

Rows that have a null value in the COMM column result in a null value for the comparison expression and are effectively not part of the result.

Other Comparison Operators

Operator	Meaning
BETWEEN ...AND...	Between two values (inclusive)
IN(list)	Match any of a list of values
LIKE	Match a character pattern
IS NULL	Is a null value

Using the BETWEEN Operator

Use the BETWEEN operator to display rows based on a range of values.

```
SQL> SELECT  ename, sal
2  FROM      emp
3  WHERE     sal BETWEEN 1000 AND 1500;
```

ENAME	SAL		
-----	-----		
MARTIN	1250	Lower	Higher
TURNER	1500	limit	limit
WARD	1250		
ADAMS	1100		
MILLER	1300		

1-48

Copyright © Oracle Corporation, 1998. All rights reserved.

ORACLE®

The BETWEEN Operator

You can display rows based on a range of values using the BETWEEN operator. The range that you specify contains a lower range and an upper range.

The SELECT statement above returns rows from the EMP table for any employee whose salary is between \$1000 and \$1500.

Values specified with the BETWEEN operator are inclusive. You must specify the lower limit first.



Class Management Note

Emphasize that the values specified with the BETWEEN operator in the example are inclusive. Point out that Turner who earns \$1500 (higher limit) is included in the output.

Demo: *l2betw.sql*

Using the IN Operator

Use the IN operator to test for values in a list.

```
SQL> SELECT  empno, ename, sal, mgr
2  FROM      emp
3  WHERE     mgr IN (7902, 7566, 7788);
```

EMPNO	ENAME	SAL	MGR
7902	FORD	3000	7566
7369	SMITH	800	7902
7788	SCOTT	3000	7566
7876	ADAMS	1100	7788

1-49

Copyright © Oracle Corporation, 1998. All rights reserved.

ORACLE®

The IN Operator

To test for values in a specified list, use the IN operator.

The above example displays employee number, name, salary, and manager's employee number of all the employees whose manager's employee number is 7902, 7566, or 7788.

The IN operator can be used with any datatype. The following example returns a row from the EMP table for any employee whose name is included in the list of names in the WHERE clause.

```
SQL> SELECT  empno, ename, mgr, deptno
2  FROM      emp
3  WHERE     ename IN ('FORD' , 'ALLEN');
```

If characters or dates are used in the list, they must be enclosed in single quotation marks (' ').



Class Management Note

Demo: *l2in.sql*

Using the LIKE Operator

- Use the LIKE operator to perform wildcard searches of valid search string values.
- Search conditions can contain either literal characters or numbers.
 - (%) denotes zero or many characters
 - (_) denotes one character

```
SQL> SELECT   ename
      2  FROM    emp
      3  WHERE   ename LIKE 'S%';
```

1-50

Copyright © Oracle Corporation, 1998. All rights reserved.

ORACLE®

The LIKE Operator

You may not always know the exact value to search for. You can select rows that match a character pattern by using the LIKE operator. The character pattern matching operation is referred to as a **wildcard** search. Two symbols can be used to construct the search string.

Symbol	Description
%	Represents any sequence of zero or more characters
_	Represents any single character

– The SELECT statement above returns the employee name from the EMP table for any employee whose name begins with an “S”. Note the uppercase “S.” Names beginning with an “s” will not be returned.

The LIKE operator can be used as a shortcut for some BETWEEN comparisons. The following example displays names and hiredates of all employees who joined in 81.

```
SQL> SELECT   ename, hiredate
      2  FROM    emp
      3  WHERE   hiredate LIKE '%81';
```

Using the LIKE Operator

- You can combine pattern matching characters.

```
SQL> SELECT  ename
      2  FROM    emp
      3  WHERE   ename LIKE '_A%';
```

```
ENAME
-----
JAMES
WARD
```

- You can use the ESCAPE identifier to search for “%” or “_”.

1-51

Copyright © Oracle Corporation, 1998. All rights reserved.

ORACLE®

Combining Wildcard Characters

The % and _ symbols can be used in any combination with literal characters. The example on the slide displays the names of all employees whose name has an “A” as the second character.

The ESCAPE Option

When you need to have an exact match for the actual “%” and “_” characters, use the ESCAPE option. This option specifies what the ESCAPE character is. To display the names of employees whose name contains “A_B”, use the following SQL statement:

```
SQL> SELECT  name
      2  FROM    emp
      3  WHERE   name LIKE '%A\_%B' ESCAPE '\';
```

The ESCAPE option identifies the backslash (\) as the escape character. In the pattern, the escape character precedes the underscore (_). This causes the Oracle8 Server to interpret the underscore literally.

Using the IS NULL Operator

Test for null values with the IS NULL operator

```
SQL> SELECT  ename, mgr
      2 FROM    emp
      3 WHERE   mgr IS NULL;
```

ENAME	MGR

KING	

The IS NULL Operator

The IS NULL operator tests for values that are null. A null value means the value is unavailable, unassigned, unknown, or inapplicable. Therefore, you cannot test with (=) because a null value cannot be equal or unequal to any value. The example above retrieves the name and manager of all employees who do not have a manager.

For example, to display name, job title, and commission for all employees who are not entitled to get a commission, use the following SQL statement:

```
SQL> SELECT  ename, job, comm
      2 FROM    emp
      3 WHERE   comm IS NULL;
```

ENAME	JOB	COMM

KING	PRESIDENT	
BLAKE	MANAGER	
CLARK	MANAGER	
...		

Logical Operators

Operator	Meaning
AND	Returns TRUE if <i>both</i> component conditions are TRUE
OR	Returns TRUE if <i>either</i> component condition is TRUE
NOT	Returns TRUE if the following condition is FALSE

Logical Operators

A logical operator combines the result of two component conditions to produce a single result based on them or to invert the result of a single condition. Three logical operators are available in SQL:

- AND
- OR
- NOT

All the examples so far have specified only one condition in the WHERE clause. You can use several conditions in one WHERE clause using the AND and OR operators.

Using the AND Operator

AND requires both conditions to be TRUE.

```
SQL> SELECT empno, ename, job, sal
2  FROM emp
3  WHERE sal >= 1100
4  AND job = 'CLERK';
```

EMPNO	ENAME	JOB	SAL
7876	ADAMS	CLERK	1100
7934	MILLER	CLERK	1300

The AND Operator

In the example, both conditions must be true for any record to be selected. Therefore, an employee who has a job title of CLERK *and* earns more than \$1100 will be selected.

All character searches are case sensitive. No rows are returned if CLERK is not in uppercase. Character strings must be enclosed in quotation marks.



AND Truth Table

The following table shows the results of combining two expressions with AND:

AND	TRUE	FALSE	UNKNOWN
TRUE	TRUE	FALSE	UNKNOWN
FALSE	FALSE	FALSE	FALSE
UNKNOWN	UNKNOWN	FALSE	UNKNOWN

Class Management Note

Demo: *l2and.sql*

Using the OR Operator

OR requires either condition to be TRUE.

```
SQL> SELECT empno, ename, job, sal
2   FROM emp
3   WHERE sal >= 1100
4   OR      job = 'CLERK';
```

EMPNO	ENAME	JOB	SAL
7839	KING	PRESIDENT	5000
7698	BLAKE	MANAGER	2850
7782	CLARK	MANAGER	2450
7566	JONES	MANAGER	2975
7654	MARTIN	SALESMAN	1250
...			

14 rows selected.

1-55

Copyright © Oracle Corporation, 1998. All rights reserved.

ORACLE®

The OR Operator

In the example, either condition can be true for any record to be selected. Therefore, an employee who has a job title of CLERK *or* earns more than \$1100 will be selected.

OR Truth Table

The following table shows the results of combining two expressions with OR:

OR	TRUE	FALSE	UNKNOWN
TRUE	TRUE	TRUE	TRUE
FALSE	TRUE	FALSE	UNKNOWN
UNKNOWN	TRUE	UNKNOWN	UNKNOWN

Class Management Note

Demo: *l2or.sql*

Using the NOT Operator

```
SQL> SELECT ename, job
2   FROM emp
3  WHERE job NOT IN ( 'CLERK', 'MANAGER', 'ANALYST' );
```

ENAME	JOB
KING	PRESIDENT
MARTIN	SALESMAN
ALLEN	SALESMAN
TURNER	SALESMAN
WARD	SALESMAN

1-56

Copyright © Oracle Corporation, 1998. All rights reserved.

ORACLE®

The NOT Operator

The example above displays name and job title of all the employees whose job title *is not* CLERK, MANAGER, or ANALYST.

NOT Truth Table

The following table shows the result of applying the NOT operator to a condition:

NOT	TRUE	FALSE	UNKNOWN
	FALSE	TRUE	UNKNOWN

Note: The NOT operator can also be used with other SQL operators such as BETWEEN, LIKE, and NULL.

```
... WHERE sal NOT BETWEEN 1000 AND 1500
... WHERE ename NOT LIKE '%A%'
... WHERE comm IS NOT NULL
```

Rules of Precedence

Order Evaluated	Operator
1	All comparison operators
2	NOT
3	AND
4	OR

Override rules of precedence by using parentheses.

Rules of Precedence

```
SQL> SELECT ename, job, sal
  2 FROM emp
  3 WHERE job='SALESMAN'
  4 OR   → job='PRESIDENT'
  5 AND → sal>1500;
```

ENAME	JOB	SAL
KING	PRESIDENT	5000
MARTIN	SALESMAN	1250
ALLEN	SALESMAN	1600
TURNER	SALESMAN	1500
WARD	SALESMAN	1250

Example of Precedence of AND Operator

In the example, there are two conditions:

- The first condition is that job is PRESIDENT *and* salary is greater than 1500.
- The second condition is that job is SALESMAN.

Therefore, the SELECT statement reads as follows:

“Select the row if an employee is a PRESIDENT *and* earns more than \$1500 *or* if the employee is a SALESMAN.”

Class Management Note

Demo: *l2sal1.sql*

Rules of Precedence

Use parentheses to force priority.

```
SQL> SELECT      ename, job, sal
  2  FROM        emp
  3  WHERE      ( job='SALESMAN'
  4  OR  →      job='PRESIDENT' )
  5  AND        sal>1500;
```

ENAME	JOB	SAL
KING	PRESIDENT	5000
ALLEN	SALESMAN	1600

Using Parentheses

In the example, there are two conditions:

- The first condition is that job is PRESIDENT *or* SALESMAN.
- The second condition is that salary is greater than 1500.

Therefore, the SELECT statement reads as follows:

“Select the row if an employee is a PRESIDENT or a SALESMAN and if the employee earns more than \$1500.”

Class Management Note

Demo: *l2sal2.sql*

ORDER BY Clause

- Sort rows with the ORDER BY clause
 - ASC: ascending order, default
 - DESC: descending order
- The ORDER BY clause comes last in the SELECT statement.

```
SQL> SELECT      ename, job, deptno, hiredate
  2  FROM          emp
  3  ORDER BY hiredate;
```

ENAME	JOB	DEPTNO	HIREDATE
-----	-----	-----	-----
SMITH	CLERK	20	17-DEC-80
ALLEN	SALESMAN	30	20-FEB-81
...			
14 rows selected.			

1-60

Copyright © Oracle Corporation, 1998. All rights reserved.

ORACLE®

The ORDER BY Clause

The order of rows returned in a query result is undefined. The ORDER BY clause can be used to sort the rows. If used, you must place the ORDER BY clause last. You can specify an expression or an alias to sort.

Syntax

```
SELECT      expr
FROM        table
[WHERE      condition (s)]
where: ORDER BY {column, expr} [ASC|DESC];
[ORDER BY  ASC orders the rows in ascending order. This is the default order.
DESC orders the rows in descending order.
```

If the ORDER BY clause is not used, the sort order is undefined, and the Oracle8 Server may not fetch rows in the same order for the same query twice. Use the ORDER BY clause to display the rows in a specific order.



Sorting in Descending Order

```
SQL> SELECT      ename, job, deptno, hiredate
  2  FROM        emp
  3  ORDER BY hiredate DESC;
```

ENAME	JOB	DEPTNO	HIREDATE
ADAMS	CLERK	20	12-JAN-83
SCOTT	ANALYST	20	09-DEC-82
MILLER	CLERK	10	23-JAN-82
JAMES	CLERK	30	03-DEC-81
FORD	ANALYST	20	03-DEC-81
KING	PRESIDENT	10	17-NOV-81
MARTIN	SALESMAN	30	28-SEP-81
...			

14 rows selected.

Default Ordering of Data

The default sort order is ascending:

- Numeric values are displayed with the lowest values first—for example, 1-999.
- Date values are displayed with the earliest value first—for example, 01-JAN-92 before 01-JAN-95.
- Character values are displayed in alphabetical order—for example, A first and Z last.
- Null values are displayed last for ascending sequences and first for descending sequences.

Reversing the Default Order

To reverse the order in which rows are displayed, specify the keyword **DESC** after the column name in the **ORDER BY** clause. The example above sorts the result by the most recently hired employee.

Sorting by Column Alias

```
SQL> SELECT    empno, ename, sal*12 annsal
  2  FROM      emp
  3  ORDER BY  annsal;
```

EMPNO	ENAME	ANNSAL
7369	SMITH	9600
7900	JAMES	11400
7876	ADAMS	13200
7654	MARTIN	15000
7521	WARD	15000
7934	MILLER	15600
7844	TURNER	18000

...

14 rows selected.

Sorting By Column Aliases

You can use a column alias in the ORDER BY clause. The above example sorts the data by annual salary.

Sorting by Multiple Columns

- The order of ORDER BY list is the order of sort.

```
SQL> SELECT  ename, deptno, sal
2  FROM      emp
3  ORDER BY  deptno, sal DESC;
```

ENAME	DEPTNO	SAL
-----	-----	-----
KING	10	5000
CLARK	10	2450
MILLER	10	1300
FORD	20	3000
...		

14 rows selected.

- You can sort by a column that is not in the SELECT list.

1-63

Copyright © Oracle Corporation, 1998. All rights reserved.

ORACLE®

Sorting by Multiple Columns

You can sort query results by more than one column. The sort limit is the number of columns in the given table.

In the ORDER BY clause, specify the columns, and separate the column names using commas. If you want to reverse the order of a column, specify DESC after its name. You can order by columns that are not included in the SELECT clause.

Example

Display name and salary of all employees. Order the result by department number and then descending order by salary.

```
SQL> SELECT  ename, salary
2  FROM      emp
3  ORDER BY  deptno, sal DESC;
```

Class Management Note

Show that the DEPTNO column is sorted in ascending order and the SAL column in descending order.

Summary

```
SELECT      [DISTINCT] {*, column [alias], ...}  
FROM        table  
[WHERE      condition(s)]  
[ORDER BY   {column, expr, alias} [ASC|DESC]];
```

Summary

In this lesson, you have learned about restricting and sorting rows returned by the SELECT statement. You have also learned how to implement various operators.

Practice Overview

- **Selecting data and changing the order of rows displayed**
- **Restricting rows by using the WHERE clause**
- **Using the double-quotation-marks in column aliases**

Practice Overview

This practice gives you a variety of exercises using the WHERE clause and the ORDER BY clause.

Practice 2

1. Create a query to display the name and salary of employees earning more than \$2850. Save your SQL statement to a file named *p2q1.sql*. Run your query.

ENAME	SAL
-----	-----
KING	5000
JONES	2975
FORD	3000
SCOTT	3000

2. Create a query to display the employee name and department number for employee number 7566.

ENAME	DEPTNO
-----	-----
JONES	20

3. Modify *p2q1.sql* to display the name and salary for all employees whose salary is not in the range of \$1500 and \$2850. Resave your SQL statement to a file named *p2q3.sql*. Rerun your query.

ENAME	SAL
-----	-----
KING	5000
JONES	2975
MARTIN	1250
JAMES	950
WARD	1250
FORD	3000
SMITH	800
SCOTT	3000
ADAMS	1100
MILLER	1300

10 rows selected.

1

Single-Row Functions

Copyright © Oracle Corporation, 1998. All rights reserved.

ORACLE®

Schedule:	Timing	Topic
	55 minutes	Lecture
	30 minutes	Practice
	85 minutes	Total

Objectives

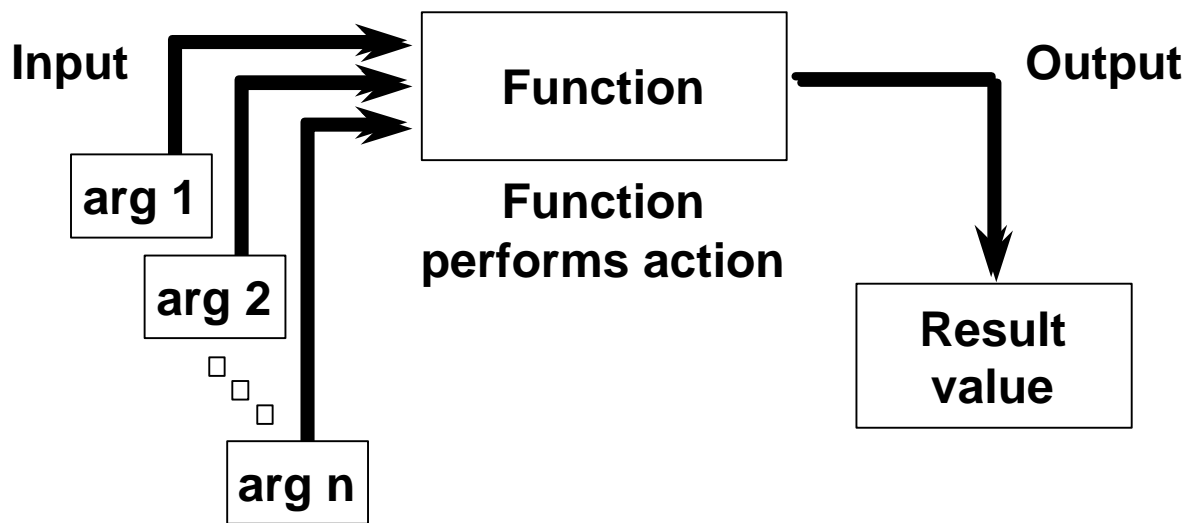
At the end of this lesson, you should be able to:

- **Describe various types of functions available in SQL**
- **Use character, number, and date functions in SELECT statements**
- **Describe the use of conversion functions**

Lesson Aim

Functions make the basic query block more powerful and are used to manipulate data values. This is the first of two lessons that explore functions. You will focus on single-row character, number, and date functions, as well as those functions that convert data from one type to another—for example, character data to numeric.

SQL Functions



1-73

Copyright © Oracle Corporation, 1998. All rights reserved.

ORACLE®

SQL Functions

Functions are a very powerful feature of SQL and can be used to do the following:

- Perform calculations on data
- Modify individual data items
- Manipulate output for groups of rows
- Format dates and numbers for display
- Convert column datatypes


SQL functions accept argument(s) and return value(s).

Note: Most of the functions described in this lesson are specific to Oracle's version of SQL. More on Oracle's version of SQL is covered in the course, *Oracle SQL Specifics*.

Class Management Note

This lesson does not discuss all functions in great detail. Present the most common functions without a long explanation.

Two Types of SQL Functions



1-74

Copyright © Oracle Corporation, 1998. All rights reserved.

ORACLE®

SQL Functions

There are two distinct types of functions:

- Single-row functions
- Multiple-row functions

Single-Row Functions

These functions operate on single-rows only and return one result per row. There are different types of single-row functions. This lesson covers those listed below.

- Character
- Number
- Date
- Conversion

Multiple-Row Functions

These functions manipulate groups of rows to give one result per group of rows.

For more information, see

Server SQL Reference, Release 8.0 for the complete list of available functions and syntax.

Oracle 8



Single-Row Functions

- Manipulate data items
- Accept arguments and return one value
- Act on each row returned
- Return one result per row
- May modify the datatype
- Can be nested

```
function_name (column|expression, [arg1, arg2,...])
```

Single-Row Functions

Single-row functions are used to manipulate data items. They accept one or more arguments and return one value for each row returned by the query. An argument can be one of the following:

- A user-supplied constant
- A variable value
- A column name
- An expression

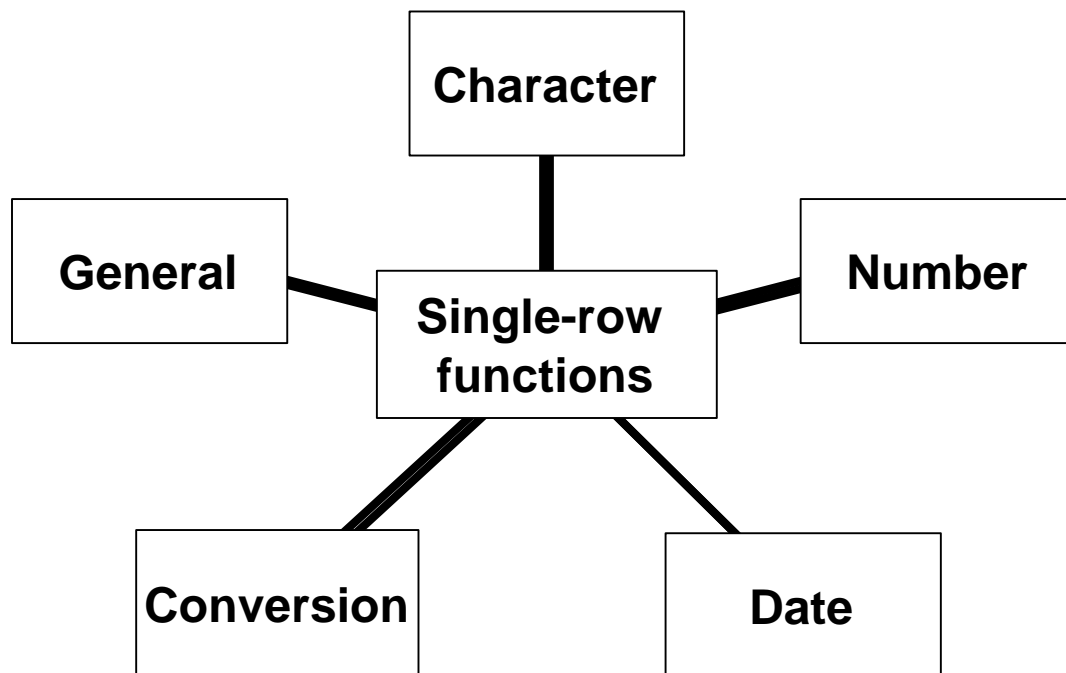
Features of Single-Row Functions

- They act on each row returned in the query.
- They return one result per row.
- They may return a data value of a different type than that referenced.
- They may expect one or more arguments.
- You can use them in SELECT, WHERE, and ORDER BY clauses. You can nest them.

In the syntax:

<i>function_name</i>	is the name of the function
<i>column</i>	is any named database column
<i>expression</i>	is any character string or calculated expression
<i>arg1, arg2</i>	is any argument to be used by the function

Single-Row Functions

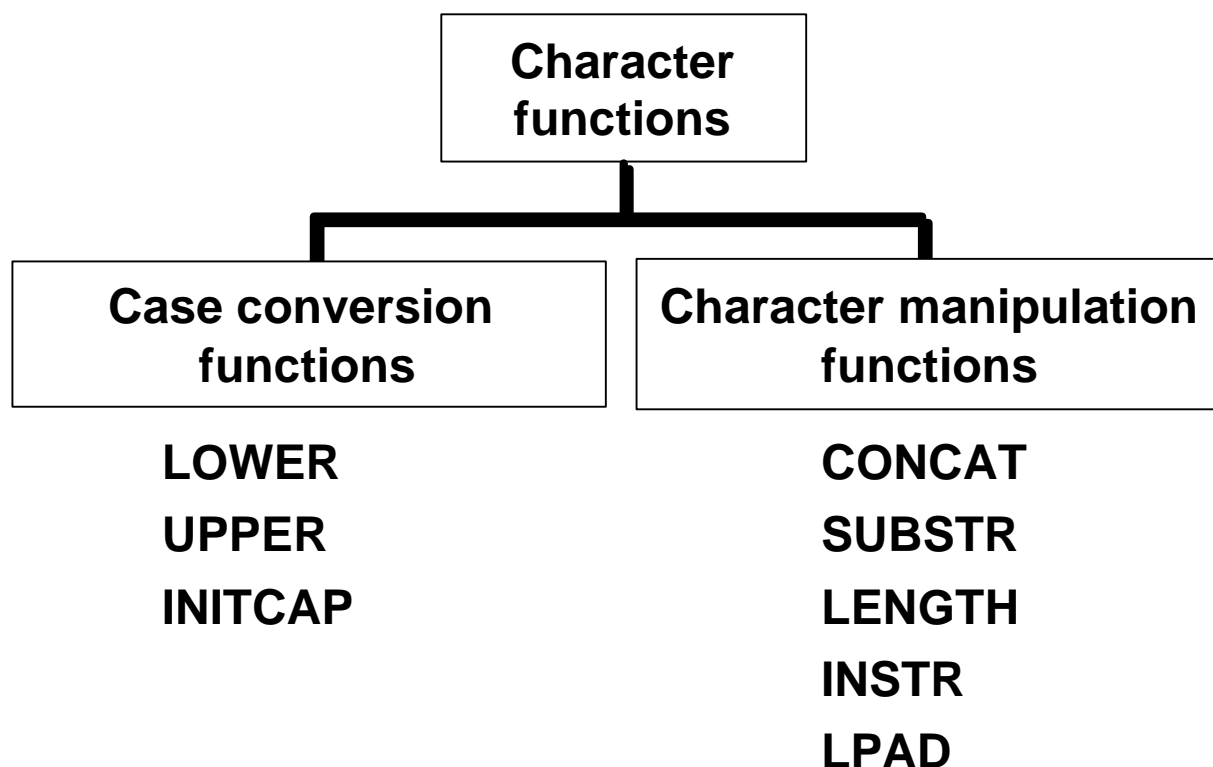


Single-Row Functions

This lesson covers the following single-row functions:

- Character functions—Accept character input and can return both character and number values.
- Number functions—Accept numeric input and return numeric values.
- Date functions—Operate on values of the date datatype. All date functions return a value of date datatype except the MONTHS_BETWEEN function, which returns a number.
- Conversion functions—Convert a value from one datatype to another.
- General functions
 - NVL function
 - DECODE function

Character Functions



1-77

Copyright © Oracle Corporation, 1998. All rights reserved.

ORACLE®

Character Functions

Single-row character functions accept character data as input and can return both character and number values. Character functions can be divided into:

- Case conversion functions
- Character manipulation functions

Function	Purpose
LOWER(<i>column</i> / <i>expression</i>)	Converts alpha character values to lowercase
UPPER(<i>column</i> / <i>expression</i>)	Converts alpha character values to uppercase
INITCAP(<i>column</i> / <i>expression</i>)	Converts alpha character values to uppercase for the first letter of each word, all other letters in lowercase
CONCAT(<i>column1</i> / <i>expression1</i> , <i>column2</i> / <i>expression2</i>)	Concatenates the first character value to the second character value. Equivalent to concatenation operator ()
SUBSTR(<i>column</i> / <i>expression</i> , <i>m</i> [, <i>n</i>])	Returns specified characters from character value starting at character position <i>m</i> , <i>n</i> characters long. If <i>m</i> is negative, the count starts from the end of the character value. If <i>n</i> is omitted, all characters to the end of the string are returned
LENGTH(<i>column</i> / <i>expression</i>)	Returns the number of characters in value
INSTR(<i>column</i> / <i>expression</i> , <i>m</i>)	Returns the numeric position of a named character
LPAD(<i>column</i> / <i>expression</i> , <i>n</i> , 'string')	Pads the character value right-justified to a total width of <i>n</i> character positions

Note: This list is a subset of the available character functions.

For more information, see

Oracle8 Server SQL Reference, Release 8.0, "Character Functions."

Case Conversion Functions

Convert case for character strings

Function	Result
LOWER('SQL Course')	sql course
UPPER('SQL Course')	SQL COURSE
INITCAP('SQL Course')	Sql Course

1-78

Copyright © Oracle Corporation, 1998. All rights reserved.

ORACLE®

Case Conversion Functions

LOWER, UPPER, and INITCAP are the three case conversion functions.

- LOWER—Converts mixed case or uppercase character string to lowercase
- UPPER—Converts mixed case or lowercase character string to uppercase
- INITCAP—Converts first letter of each word to uppercase and remaining letters to lowercase

```
SQL> SELECT 'The job title for ' || INITCAP(ename) || ' is '
```

```
2      || LOWER(job) AS "EMPLOYEE DETAILS"
```

```
3 FROM emp;
```

```
EMPLOYEE DETAILS
```

```
-----  
The job title for King is president  
The job title for Blake is manager  
The job title for Clark is manager  
...  
14 rows selected.
```

Using Case Conversion Functions

Display the employee number, name, and department number for employee Blake.

```
SQL> SELECT  empno, ename, deptno
  2  FROM      emp
  3  WHERE      ename = 'blake';
no rows selected
```

```
SQL> SELECT  empno, ename, deptno
  2  FROM      emp
  3  WHERE      LOWER(ename) = 'blake';
```

EMPNO	ENAME	DEPTNO
-----	-----	-----
7698	BLAKE	30

1-79

Copyright © Oracle Corporation, 1998. All rights reserved.

ORACLE®

Case Conversion Functions

The example above displays the employee number, name, and department number of employee BLAKE.

The WHERE clause of the first SQL statement specifies the employee name as 'blake'. Since all the data in the EMP table is stored in uppercase, the name 'blake' does not find a match in the EMP table and as a result no rows are selected.

The WHERE clause of the second SQL statement specifies that the employee name in the EMP table be converted to lowercase and then be compared to 'blake'. Since both the names are in lowercase now, a match is found and one row is selected. The WHERE clause can be rewritten in the following manner to produce the same result:

The name in the output appears as it was stored in the database. To display the name with the first letter capitalized, use the INITCAP function in the SELECT statement.

```
... WHERE      ename = 'BLAKE'
```

```
SQL> SELECT  empno, INITCAP(ename), deptno
  2  FROM      emp
  3  WHERE      LOWER(ename) = 'blake';
```

Character Manipulation Functions

Manipulate character strings

Function	Result
CONCAT('Good', 'String')	GoodString
SUBSTR('String',1,3)	Str
LENGTH('String')	6
INSTR('String', 'r')	3
LPAD(sal,10,'*')	*****5000

Character Manipulation Functions

CONCAT, SUBSTR, LENGTH, INSTR, and LPAD are the five character manipulation functions covered in this lesson.

- **CONCAT**—Joins values together. You are limited to using two parameters with CONCAT.
- **SUBSTR**—Extracts a string of determined length.
- **LENGTH**—Shows the length of a string as a numeric value.
- **INSTR**—Finds numeric position of a named character.
- **LPAD**—Pads the character value right-justified.

Note: RPAD character manipulation function pads the character value left justified.

Using the Character Manipulation Functions

```
SQL> SELECT  ename, CONCAT (ename, job), LENGTH(ename),  
2          INSTR(ename, 'A')  
3 FROM      emp  
4 WHERE     SUBSTR(job,1,5) = 'SALES';
```

ENAME	CONCAT (ENAME , JOB)	LENGTH (ENAME)	INSTR (ENAME , 'A')
MARTIN	MARTINSALESMAN	6	2
ALLEN	ALLENSALESMAN	5	1
TURNER	TURNERSALESMAN	6	0
WARD	WARDSALESMAN	4	2

Character Manipulation Functions

The above example displays employee name and job joined together, length of the employee name, and the numeric position of the letter A in the employee name, for all employees who are in sales.

Example

Modify the above SQL statement to display the data for those employees whose names end with an N.

```
SQL> SELECT  ename, CONCAT(ename, job), LENGTH(ename),  
            INSTR(ename, 'A')  
2 FROM      emp  
3 WHERE     SUBSTR(ename, -1, 1) = 'N';
```

ENAME	CONCAT (ENAME , JOB)	LENGTH (ENAME)	INSTR (ENAME , 'A')
MARTIN	MARTINSALESMAN	6	2
ALLEN	ALLENSALESMAN	5	1

Number Functions

- **ROUND:** Rounds value to specified decimal

ROUND(45.926, 2) ➡ 45.93

- **TRUNC:** Truncates value to specified decimal

TRUNC(45.926, 2) ➡ 45.92

- **MOD:** Returns remainder of division

MOD(1600, 300) ➡ 100

1-82

Copyright © Oracle Corporation, 1998. All rights reserved.

ORACLE®

Number Functions

Number functions accept numeric input and return numeric values. This section describes some of the number functions.

Function	Purpose
ROUND (<i>column expression, n</i>)	Rounds the column, expression, or value to <i>n</i> decimal places. If <i>n</i> is omitted, no decimal places. If <i>n</i> is negative, numbers to left of the decimal point are rounded.
TRUNC (<i>column expression,n</i>)	Truncates the column, expression, or value to <i>n</i> decimal places, or if <i>n</i> is omitted, no decimal places. If <i>n</i> is negative, numbers left of the decimal point are truncated to zero.
MOD (<i>m,n</i>)	Returns the remainder of <i>m</i> divided by <i>n</i> .

Note: This list is a subset of the available number functions.

For more information, see

Oracle8 Server SQL Reference, Release 8.0, "Number Functions."



Using the ROUND Function

Display the value 45.923 rounded to the hundredth, no, and ten decimal places.

```
SQL> SELECT ROUND(45.923,2), ROUND(45.923,0),  
2          ROUND(45.923,-1)  
3 FROM     SYS.DUAL;
```

ROUND(45.923,2)	ROUND(45.923,0)	ROUND(45.923,-1)
45.92	46	50

ROUND Function

The ROUND function rounds the column, expression, or value to *n* decimal places. If the second argument is 0 or is missing, the value is rounded to zero decimal places. If the second argument is 2, the value is rounded to two decimal places, or to the hundredths. Conversely, if the second argument is -2, the value is rounded to two decimal places to the left, or to the hundreds.

The ROUND function can also be used with date functions. You will see examples later in this lesson.

The SYS.DUAL is a dummy table that maintains some system information that is sometimes needed. You will see more about this later.



Using the TRUNC Function

Display the value 45.923 truncated to the hundredth, no, and ten decimal places.

```
SQL> SELECT TRUNC(45.923,2), TRUNC(45.923),  
2          TRUNC(45.923,-1)  
3 FROM     SYS.DUAL;
```

TRUNC(45.923,2)	TRUNC(45.923)	TRUNC(45.923,-1)
-----	-----	-----
45.92	45	40

TRUNC Function

The TRUNC function truncates the column, expression, or value to *n* decimal places.

The TRUNC function works with similar arguments as the ROUND function. If the second argument is 0 or is missing, the value is truncated to zero decimal places. If the second argument is 2, the value is truncated to two decimal places, or to the hundredths. Conversely, if the second argument is -2, the value is truncated to two decimal places to the left, or to the hundreds.

Like the ROUND function, the TRUNC function can also be used with date functions.

Using the MOD Function

Calculate the remainder of the ratio of salary to commission for all employees whose job title is a salesman.

```
SQL> SELECT  ename, sal, comm, MOD(sal, comm)
2  FROM      emp
3  WHERE     job = 'SALESMAN';
```

ENAME	SAL	COMM	MOD (SAL , COMM)
MARTIN	1250	1400	1250
ALLEN	1600	300	100
TURNER	1500	0	1500
WARD	1250	500	250

MOD Function

The MOD function finds the remainder of value1 divided by value2. The example above calculates the remainder of the ratio of salary to commission for all employees whose job title is a salesman.

Working with Dates

- Oracle stores dates in an internal numeric format: Century, year, month, day, hours, minutes, seconds.
- The default date format is DD-MON-YY.
- SYSDATE is a function returning date and time.
- DUAL is a dummy table used to view SYSDATE.

Oracle Date Format

Oracle stores dates in an internal numeric format, representing the century, year, month, day, hours, minutes, and seconds.

The default display and input format for any date is DD-MON-YY. Valid Oracle dates are between January 1, 4712 B.C., and December 31, 9999 A.D.

SYSDATE

SYSDATE is a date function that returns the current date and time. You can use SYSDATE just as you would use any other column name. For example, you can display the current date by selecting SYSDATE from a table. It is customary to select SYSDATE from a dummy table called DUAL.

DUAL

The DUAL table is owned by the user SYS and can be accessed by all users. It contains one column, DUMMY, and one row with the value X. The DUAL table is useful when you want to return a value once only—for instance, the value of a constant, pseudocolumn, or expression that is not derived from a table with user data.

Example

Display the current date using the DUAL table.

```
SQL> SELECT SYSDATE
2 FROM SYS.DUAL;
```

Arithmetic with Dates

- Add or subtract a number to or from a date for a resultant *date* value.
- Subtract two dates to find the *number* of days between those dates.
- Add *hours* to a date by dividing the number of hours by 24.

Arithmetic with Dates

Since the database stores dates as numbers, you can perform calculations using arithmetic operators such as addition and subtraction. You can add and subtract number constants as well as dates.

You can perform the following operations:

Operation	Result	Description
date + number	Date	Adds a number of days to a date
date - number	Date	Subtracts a number of days from a date
date - date	Number of days	Subtracts one date from another
date + number/24	Date	Adds a number of hours to a date

Using Arithmetic Operators with Dates

```
SQL> SELECT ename, (SYSDATE-hiredate)/7 WEEKS  
2 FROM emp  
3 WHERE deptno = 10;
```

ENAME	WEEKS
KING	830.93709
CLARK	853.93709
MILLER	821.36566

Arithmetic with Dates (continued)

The example on the slide displays the name and the number of weeks employed for all employees in department 10. It subtracts the current date (SYSDATE) from the date on which the employee was hired and divides the result by 7 to calculate the number of weeks that a worker has been employed.

Note: SYSDATE is a SQL function that returns the current date and time. Your results may differ from the example.

Class Management Note

If an older date is subtracted from a more current date, the difference is a negative number.

Date Functions

FUNCTION	DESCRIPTION
MONTHS_BETWEEN	Number of months between two dates
ADD_MONTHS	Add calendar months to date
NEXT_DAY	Next day of the date specified
LAST_DAY	Last day of the month
ROUND	Round date
TRUNC	Truncate date

Date Functions

Date functions operate on Oracle dates. All date functions return a value of DATE datatype except MONTHS_BETWEEN, which returns a numeric value.

- MONTHS_BETWEEN(*date1*, *date2*)—Finds the number of months between *date1* and *date2*. The result can be positive or negative. If *date1* is later than *date2*, the result is positive; if *date1* is earlier than *date2*, the result is negative. The noninteger part of the result represents a portion of the month.
- ADD_MONTHS(*date*, *n*)—Adds *n* number of calendar months to *date*. *n* must be an integer and can be negative.
- NEXT_DAY(*date*, '*char*')—Finds the date of the next specified day of the week ('*char*') following *date*. *char* may be a number representing a day or a character string.
- LAST_DAY(*date*)—Finds the date of the last day of the month that contains *date*.
- ROUND(*date*[, '*fmt*'])—Returns *date* rounded to the unit specified by the format model *fmt*. If the format model *fmt* is omitted, *date* is rounded to the nearest date.
- TRUNC(*date*[, '*fmt*'])—Returns *date* with the time portion of the day truncated to the unit specified by the format model *fmt*. If the format model *fmt* is omitted, *date* is truncated to the nearest day.

This list is a subset of the available date functions. The format models are covered later in this chapter. Examples of format models are month or year.

Using Date Functions

- **MONTHS_BETWEEN ('01-SEP-95','11-JAN-94')**
→ **19.6774194**
- **ADD_MONTHS ('11-JAN-94',6)** → **'11-JUL-94'**
- **NEXT_DAY ('01-SEP-95','FRIDAY')** → **'08-SEP-95'**
- **LAST_DAY('01-SEP-95')** → **'30-SEP-95'**

1-90

Copyright © Oracle Corporation, 1998. All rights reserved.

ORACLE®

Date Functions (continued)

For all employees employed for fewer than 200 months, display the employee number, hire date, number of months employed, six month review date, first friday after hire date, and the last day of the month when hired.

```
SQL> SELECT empno, hiredate,
2 MONTHS BETWEEN(SYSDATE, hiredate) TENURE,
3 ADD MONTHS(hiredate, 6) REVIEW,
4 NEXT_DAY(hiredate, 'FRIDAY'), LAST_DAY(hiredate)
5 FROM emp
6 WHERE MONTHS BETWEEN (SYSDATE, hiredate)<200;
```

```

EMPNO  HIREDATE      TENURE  REVIEW      NEXT_DAY(   LAST_DAY(
-----
      7839 17-NOV-81  192.24794 17-MAY-82  20-NOV-81  30-NOV-81
      7698 01-MAY-81  198.76407 01-NOV-81  08-MAY-81  31-MAY-81
...
11 rows selected.

```

Using Date Functions

- **ROUND('25-JUL-95','MONTH') ➡ 01-AUG-95**
- **ROUND('25-JUL-95','YEAR') ➡ 01-JAN-96**
- **TRUNC('25-JUL-95','MONTH') ➡ 01-JUL-95**
- **TRUNC('25-JUL-95','YEAR') ➡ 01-JAN-95**

Date Functions (continued)

The ROUND and TRUNC functions can be used for number and date values. When using these functions with dates, they round or truncate to the specified format model. Therefore, you can round dates to the nearest year or month.

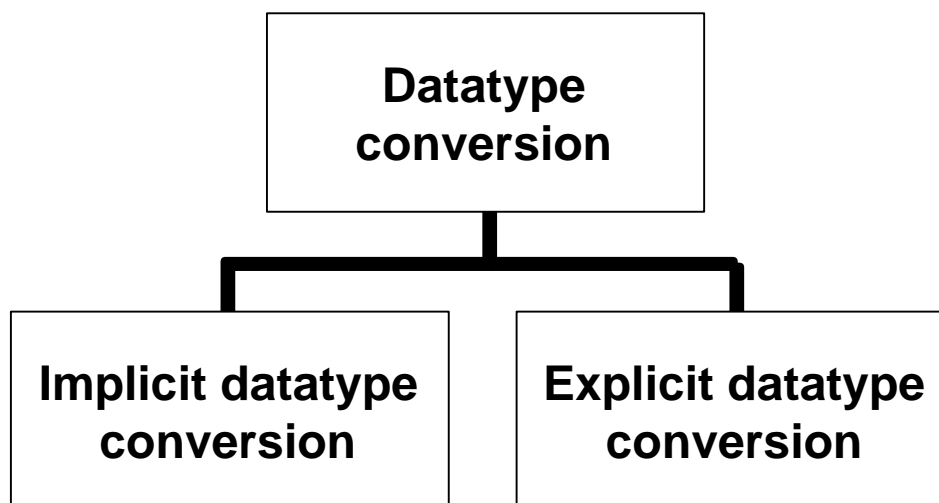
Example

Compare the hire dates for all employees who started in 1987. Display the employee number, hire date, and month started using the ROUND and TRUNC functions.

```
SQL> SELECT empno, hiredate,  
2          ROUND(hiredate, 'MONTH'), TRUNC(hiredate, 'MONTH')  
3 FROM emp  
4 WHERE hiredate like '%87';
```

EMPNO	HIREDATE	ROUND(HIR	TRUNC(HIR
7788	19-APR-87	01-MAY-87	01-APR-87
7876	23-MAY-87	01-JUN-87	01-MAY-87

Conversion Functions



Conversion Functions

In addition to Oracle datatypes, columns of tables in an Oracle8 database can be defined using ANSI, DB2, and SQL/DS datatypes. However, Oracle8 Server internally converts such datatypes to Oracle8 datatypes.

In some cases, Oracle8 Server allows data of one datatype where it expects data of a different datatype. This is allowed when Oracle8 Server can automatically convert the data to the expected datatype. This datatype conversion can be done *implicitly* by Oracle8 Server or *explicitly* by the user.

Implicit datatype conversions work according to the rules explained in next two slides.

Explicit datatype conversions are done by using the conversion functions. Conversion functions convert a value from one datatype to another. Generally, the form of the function names follow the convention *datatype TO datatype*. The first datatype is the input datatype; the last datatype is the output.

Note: Though implicit datatype conversion is available, it is recommended that you do explicit datatype conversion to ensure reliability of your SQL statements.

Implicit Datatype Conversion

For assignments, Oracle can automatically convert

From	To
VARCHAR2 or CHAR	NUMBER
VARCHAR2 or CHAR	DATE
NUMBER	VARCHAR2
DATE	VARCHAR2

1-93

Copyright © Oracle Corporation, 1998. All rights reserved.

ORACLE®

Implicit Datatype Conversion

For assignments, Oracle8 Server can automatically convert the following:

- VARCHAR2 or CHAR to NUMBER
- VARCHAR2 or CHAR to DATE
- NUMBER to VARCHAR2
- DATE to VARCHAR2

The assignment succeeds if Oracle8 Server can convert the datatype of the value used in the assignment to that of the assignment's target.

Implicit Datatype Conversion

For expression evaluation, Oracle can automatically convert

From	To
VARCHAR2 or CHAR	NUMBER
VARCHAR2 or CHAR	DATE

Implicit Datatype Conversion


For expression evaluation, Oracle8 Server can automatically convert the following:

- VARCHAR2 or CHAR to NUMBER
- VARCHAR2 or CHAR to DATE

In general, Oracle8 Server uses the rule for expression when a datatype conversion is needed in places not covered by a rule for assignment conversions.

Note: CHAR to NUMBER conversions succeed only if the character string represents a valid number. CHAR to DATE conversions succeed only if the character string has the default format DD-MON-YY.

Explicit Datatype Conversion



1-95

Copyright © Oracle Corporation, 1998. All rights reserved.

ORACLE®

Explicit Datatype Conversion

SQL provides three functions to convert a value from one datatype to another.

Function	Purpose
TO_CHAR(<i>number date</i> ,[' <i>fmt</i> '])	Converts a number or date value to a VARCHAR2 character string with format model <i>fmt</i> .
TO_NUMBER(<i>char</i>)	Converts a character string containing digits to a number.
TO_DATE(<i>Char</i> ,[' <i>fmt</i> '])	Converts a character string representing a date to a date value according to the <i>fmt</i> specified. If <i>fmt</i> is omitted, format is DD-MON-YY.

Note: This list is a subset of the available conversion functions.
For more information, see *Oracle8 Server SQL Reference, Release 8.0, "Conversion Functions."*



Class Management Note

An additional conversion function is `CHR(number)` that returns the character having the binary equivalent of *number* as a VARCHAR2 value in the database character set.

TO_CHAR Function with Dates

```
TO_CHAR(date, 'fmt') 
```

The format model:

- **Must be enclosed in single quotation marks and is case sensitive**
- **Can include any valid date format element**
- **Has an *fm* element to remove padded blanks or suppress leading zeros**
- **Is separated from the date value by a comma**

1-96

Copyright © Oracle Corporation, 1998. All rights reserved.

ORACLE®

Displaying a Date in a Specific Format

Previously, all Oracle date values were displayed in the DD-MON-YY format. The TO_CHAR function allows you to convert a date from this default format to one specified by you.

Guidelines

- The format model must be enclosed in single quotation marks and is case sensitive.
- The format model can include any valid date format element. Be sure to separate the date value from the format model by a comma.
- The names of days and months in the output are automatically padded with blanks.
- To remove padded blanks or to suppress leading zeros, use the fill mode *fm* element.
- You can resize the display width of the resulting character field with the SQL*Plus COLUMN command.
- The resultant column width is 80 characters by default.

```
SQL> SELECT empno, TO_CHAR(hiredate, 'MM/YY') Month_Hired
2 FROM emp
3 WHERE ename = 'BLAKE';
```


Date Format Model Elements

YYYY	Full year in numbers
YEAR	Year spelled out
MM	2-digit value for month
MONTH	Full name of the month
DY	3-letter abbreviation of the day of the week
DAY	Full name of the day

Sample Valid Date Format Elements

Element	Description
SCC or CC	Century; S prefixes BC date with -
Years in dates YYYY or SYYYY	Year; S prefixes BC date with -
YYY or YY or Y	Last 3, 2, or 1 digit(s) of year
Y,YYY	Year with comma in this position
IYYY, IYY, IY, I	4, 3, 2, or 1 digit year based on the ISO standard
SYEAR or YEAR	Year spelled out; S prefixes BC date with -
BC or AD	BC/AD indicator
B.C. or A.D.	BC/AD indicator with periods
Q	Quarter of year
MM	Month, two-digit value
MONTH	Name of month padded with blanks to length of 9 characters
MON	Name of month, three-letter abbreviation
RM	Roman numeral month
WW or W	Week of year or month
DDD or DD or D	Day of year, month, or week
DAY	Name of day padded with blanks to length of 9 characters.
DY	Name of day; 3-letter abbreviation.
J	Julian day; the number of days since 31 December 4713 BC.

Date Format Model Elements

- Time elements format the time portion of the date.

HH24:MI:SS AM

15:45:32 PM

- Add character strings by enclosing them in double quotation marks.

DD "of" MONTH

12 of OCTOBER

- Number suffixes spell out numbers.

ddspth

fourteenth

Time Formats

Use the formats listed in the following tables to display time information and literals and to change numerals to spelled numbers.

Element	Description
AM or PM	Meridian indicator
A.M. or P.M.	Meridian indicator with periods
HH or HH12 or HH24	Hour of day or hour (1-12) or hour (0-23)
MI	Minute (0-59)
SS	Second (0-59).
SSSSS	Seconds past midnight (0-86399).

Other Formats

Specifying Suffixes to Influence Number Display

Element	Description
/ . ,	Punctuation is reproduced in the result
“ of the ”	Quoted string is reproduced in the result

Element	Description
TH	Ordinal number (for example, DDTH for 4TH)
SP	Spelled-out number (for example, DDSP for FOUR)
SPTH or THSP	Spelled-out ordinal numbers (for example, DDSPTH for FOURTH)

RR Date Format

Current Year	Specified Date	RR Format	YY Format
1995	27-OCT-95	1995	1995
1995	27-OCT-17	2017	1917
2001	27-OCT-17	2017	2017
2001	27-OCT-95	1995	2095

		If the specified two-digit year is	
		0-49	50-99
If two digits of the current year are	0-49	The return date is in the current century.	The return date is in the century before the current one.
	50-99	The return date is in the century after the current one.	The return date is in the current century.

1-99

Copyright © Oracle Corporation, 1998. All rights reserved.

ORACLE®

The RR Date Format Element

The RR date format is similar to the YY element, but it allows you to specify different centuries. You can use the RR date format element instead of YY, so that the century of the return value varies according to the specified two-digit year and the last two digits of the current year. The table on the slide summarizes the behavior of the RR element.

Current Year	Given Date	Interpreted (RR)	Interpreted (YY)
1994	27-OCT-95	1995	1995
1994	27-OCT-17	2017	1917
2001	27-OCT-17	2017	2017

Class Management Note

RR is available in Oracle7, not Oracle Version 6. NLS parameters can be added to the `init.ora` file to set default date formats and language names and abbreviations. For more information, see *Oracle8 Server SQL Reference*, Release 8.0, “Alter Session” clause.

Demo: (for Date Format Model Elements) *l3hire.sql*

Using TO_CHAR Function with Dates

```
SQL> SELECT ename,  
2          TO_CHAR(hiredate, 'fmDD Month YYYY') HIREDATE  
3 FROM emp;
```

ENAME	HIREDATE
-----	-----
KING	17 November 1981
BLAKE	1 May 1981
CLARK	9 June 1981
JONES	2 April 1981
MARTIN	28 September 1981
ALLEN	20 February 1981
...	

14 rows selected.

1-100

Copyright © Oracle Corporation, 1998. All rights reserved.

ORACLE®

TO_CHAR Function with Dates

The above SQL statement displays the name and hire dates for all the employees. The hire date looks like 17 November 1981.

Example

Modify the above example to display the dates in a format that looks like Seventh of February 1981 08:00:00 AM.

```
SQL> SELECT ename,  
2          TO_CHAR(hiredate, 'fmDdspth "of" Month YYYY fmHH:MI:SSAM')  
3          HIREDATE  
4 FROM emp;
```

ENAME	HIREDATE
-----	-----
KING	Seventeenth of November 1981 12:00:00 AM
BLAKE	First of May 1981 12:00:00 AM
...	

14 rows selected.

Notice that the month follows the format model specified (NTHCAP):

TO_CHAR Function with Numbers

```
TO_CHAR(number, 'fmt')
```

Use these formats with the TO_CHAR function to display a number value as a character.

9	Represents a number
0	Forces a zero to be displayed
\$	Places a floating dollar sign
L	Uses the floating local currency symbol
.	Prints a decimal point
,	Prints a thousand indicator

1-101

Copyright © Oracle Corporation, 1998. All rights reserved.

ORACLE®

TO_CHAR Function with Numbers

When working with number values such as character strings, you should convert those numbers to the character datatype using the TO_CHAR function, which translates a value of NUMBER datatype to VARCHAR2 datatype. This technique is especially useful with concatenation.

Number Format Elements

If you are converting a number to character datatype, you can use the elements listed below.

Element	Description	Example	Result
9	Numeric position (number of 9s determine display width)	999999	1234
0	Display leading zeros	099999	001234
\$	Floating dollar sign	\$999999	\$1234
L	Floating local currency symbol	L999999	FF1234
.	Decimal point in position specified	999999.99	1234.00
,	Comma in position specified	999,999	1,234
MI	Minus signs to right (negative values)	999999MI	1234-
PR	Parenthesize negative numbers	999999PR	<1234>
EEEE	Scientific notation (format must specify four Es)	99.999EEEE	1.234E+03
V	Multiply by 10 <i>n</i> times (<i>n</i> = no. of 9s after V)	9999V99	123400
B	Display zero values as blank, not 0	B9999.99	1234.00

Using TO_CHAR Function with Numbers

```
SQL> SELECT    TO_CHAR(sal, '$99,999') SALARY
      2 FROM      emp
      3 WHERE     ename = 'SCOTT';
```

```
SALARY
-----
$3,000
```

Guidelines

- The Oracle8 Server displays a string of pound signs (#) in place of a whole number whose digits exceed the number of digits provided in the format model.
- The Oracle8 Server rounds the stored decimal value to the number of decimal spaces provided in the format model.

TO_NUMBER and TO_DATE Functions

- Convert a character string to a number format using the TO_NUMBER function

```
TO_NUMBER(char)
```

- Convert a character string to a date format using the TO_DATE function

```
TO_DATE(char[, 'fmt'])
```

1-103

Copyright © Oracle Corporation, 1998. All rights reserved.

ORACLE®

TO_NUMBER and TO_DATE Functions

You may want to convert a character string to either a number or a date. To accomplish this task, you use the TO_NUMBER or TO_DATE functions. The format model you choose will be based on the previously demonstrated format elements.

Example

Display the names and hire dates of all the employees who joined on February 22, 1981.

```
SQL> SELECT ename, hiredate
2 FROM emp
3 WHERE hiredate = TO_DATE('February 22, 1981', 'Month dd, YYYY');
```

ENAME	HIREDATE
WARD	22-FEB-81

NVL Function

Converts null to an actual value

- Datatypes that can be used are date, character, and number.
- Datatypes must match
 - NVL(comm,0)
 - NVL(hiredate,'01-JAN-97')
 - NVL(job,'No Job Yet')

1-104

Copyright © Oracle Corporation, 1998. All rights reserved.

ORACLE®

The NVL Function

To convert a null value to an actual value, use the NVL function.

Syntax

~~NVL (expr1, expr2)~~

where:	<i>expr1</i>	is the source value or expression that may contain null
	<i>expr2</i>	is the target value for converting null

You can use the NVL function to convert any datatype, but the return value is always the same as the datatype of *expr1*.

NVL Conversions for Various Datatypes

Datatype	Conversion Example
NUMBER	NVL(<i>number_column</i> ,9)
DATE	NVL(<i>date_column</i> , '01-JAN-95')
CHAR or VARCHAR2	NVL(<i>character_column</i> , 'Unavailable')

Using the NVL Function

```
SQL> SELECT ename, sal, comm, (sal*12)+NVL(comm,0)
2 FROM emp;
```

ENAME	SAL	COMM	(SAL*12)+NVL(COMM,0)
KING	5000		60000
BLAKE	2850		34200
CLARK	2450		29400
JONES	2975		35700
MARTIN	1250	1400	16400
ALLEN	1600	300	19500
...			

14 rows selected.

1-105

Copyright © Oracle Corporation, 1998. All rights reserved.

ORACLE®

NVL Function

To calculate the annual compensation of all employees, you need to multiply the monthly salary by 12 and then add the commission to it.

```
SQL> SELECT ename, sal, comm, (sal*12)+comm
2 FROM emp;
```

ENAME	JOB	(SAL*12)+COMM
KING	PRESIDENT	
BLAKE	MANAGER	
CLARK	MANAGER	
JONES	MANAGER	
MARTIN	SALESMAN	16400
...		

14 rows selected.

Notice that the annual compensation is calculated only for those employees who earn a commission. If any column value in an expression is null, the result is null. To calculate values for all employees, you must convert the null value to a number before applying the arithmetic operator. In the example on the slide, the NVL function is used to convert null values to zero.

DECODE Function

Facilitates conditional inquiries by doing the work of a CASE or IF-THEN-ELSE statement

```
DECODE(col/expression, search1, result1  
      [, search2, result2, ..., ]  
      [, default])
```

The DECODE Function

The DECODE function decodes an expression in a way similar to the IF-THEN-ELSE logic used in various languages. The DECODE function decodes *expression* after comparing it to each *search* value. If the expression is the same as *search*, *result* is returned.

If the default value is omitted, a null value will be returned where a search value does not match any of the result values.

Using the DECODE Function

```
SQL> SELECT job, sal,  
2          DECODE(job, 'ANALYST'  SAL*1.1,  
3                      'CLERK',    SAL*1.15,  
4                      'MANAGER',  SAL*1.20,  
5                      SAL)  
6          REVISED_SALARY  
7 FROM emp;
```

JOB	SAL	REVISED_SALARY
PRESIDENT	5000	5000
MANAGER	2850	3420
MANAGER	2450	2940
...		

14 rows selected.

1-107

Copyright © Oracle Corporation, 1998. All rights reserved.

ORACLE®

Using the DECODE Function

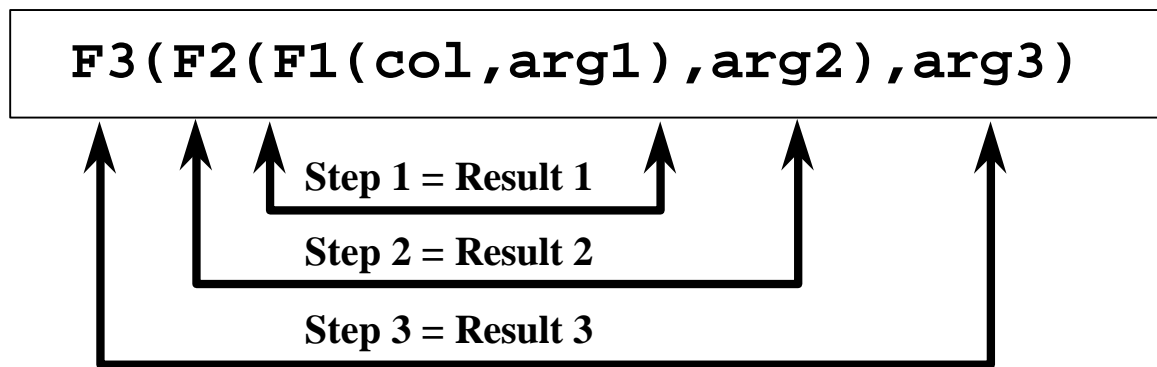
In the SQL statement above, the value of JOB is decoded. If JOB is ANALYST, the salary increase is 10%; if JOB is CLERK, the salary increase is 15%; if JOB is MANAGER, the salary increase is 20%. For all other job roles, there is no increase in salary.

The same statement can be written as an IF-THEN-ELSE statement:

```
IF job = 'ANALYST' THEN sal = sal*1.1  
IF job = 'CLERK' THEN sal = sal*1.15  
IF job = 'MANAGER' THEN sal = sal*1.20  
ELSE sal = sal
```

Nesting Functions

- **Single-row functions can be nested to any level**
- **Nested functions are evaluated from deepest level to the least deep level**



Nesting Functions

Single-row functions can be nested to any depth. Nested functions are evaluated from the innermost level to the outermost level. Some examples follow to show you the flexibility of these functions.

Nesting Functions

```
SQL> SELECT  ename,
2           NVL(TO_CHAR(mgr), 'No Manager')
3 FROM      emp
4 WHERE     mgr IS NULL;
```

```
ENAME          NVL(TO_CHAR(MGR), 'NOMANAGER')
-----
KING           No Manager
```

1-109

Copyright © Oracle Corporation, 1998. All rights reserved.

ORACLE®

Nesting Functions

The example above displays the head of the company, who has no manager. The evaluation of the SQL statement involves two steps:

1. Evaluate the inner function to convert a number value to a character string.
 - Result1 = TO_CHAR(mgr)
2. Evaluate the outer function to replace the null value with a text string.
 - NVL(Result1, 'No Manager')

The entire expression becomes the column heading since no column alias was given.

Example



Display the date of the next Friday that is six months from the hire date. The resultant date should look like Friday, March 12th, 1982. Order the results by hire date.

```
SQL> SELECT  TO_CHAR(NEXT_DAY(ADD_MONTHS
2           (hiredate, 6), 'FRIDAY'),
3           'fmDay, Month ddth, YYYY')
4           "Next 6 Month Review"
5 FROM      emp
6 ORDER BY  hiredate;
```

Class Management Note

Demo: l3nest.sql

Summary

Use functions to:

- **Perform calculations on data**
- **Modify individual data items**
- **Manipulate output for groups of rows**
- **Alter date formats for display**
- **Convert column datatypes**

Single-Row Functions

Single-row functions can be nested to any level. Single-row functions can manipulate

- Character data
 - LOWER, UPPER, INITCAP, CONCAT, SUBSTR, INSTR, LENGTH
- Number data
 - ROUND, TRUNC, MOD
- Date data
 - MONTHS_BETWEEN, ADD_MONTHS, NEXT_DAY, LAST_DAY, ROUND, TRUNC
 - Date values can also use arithmetic operators.
- Conversion functions can convert character, date, and numeric values.
 - TO_CHAR, TO_DATE, TO_NUMBER

SYSDATE and DUAL

SYSDATE is a date function that returns the current date and time. It is customary to select SYSDATE from a dummy table called DUAL.

Practice Overview

- **Creating queries that require the use of numeric, character, and date functions**
- **Using concatenation with functions**
- **Writing case-insensitive queries to test the usefulness of character functions**
- **Performing calculations of years and months of service for an employee**
- **Determining the review date for an employee**

1-111

Copyright © Oracle Corporation, 1998. All rights reserved.

ORACLE®

Practice Overview

This practice is designed to give you a variety of exercises using different functions available for character, number, and date data types.

Remember that for nested functions, the results are evaluated from the innermost function to the outermost function.



Practice 3

1. Write a query to display the current date. Label the column Date.

```
Date
-----
28-OCT-97
```

2. Display the employee number, name, salary, and salary increase by 15% expressed as a whole number. Label the column New Salary. Save your SQL statement to a file named *p3q2.sql*.
3. Run your query in the file *p3q2.sql*.

```
EMPNO  ENAME      SAL  New Salary
-----  -
7839  KING        5000      5750
7698  BLAKE       2850      3278
7782  CLARK       2450      2818
7566  JONES       2975      3421
7654  MARTIN     1250      1438
7499  ALLEN       1600      1840
7844  TURNER      1500      1725
7900  JAMES        950      1093
7521  WARD        1250      1438
7902  FORD        3000      3450
7369  SMITH        800        920
7788  SCOTT       3000      3450
7876  ADAMS       1100      1265
7934  MILLER      1300      1495
14 rows selected.
```

4. Modify your query *p3q2.sql* to add an additional column that will subtract the old salary from the new salary. Label the column Increase. Rerun your query.

```
EMPNO  ENAME      SAL  New Salary  Increase
-----  -
7839  KING        5000      5750       750
7698  BLAKE       2850      3278       428
7782  CLARK       2450      2818       368
7566  JONES       2975      3421       446
```


1

Displaying Data from Multiple Tables

Copyright © Oracle Corporation, 1998. All rights reserved.

ORACLE®

Schedule:	Timing	Topic
	40 minutes	Lecture
	50 minutes	Practice
	90 minutes	Total

Objectives

At the end of this lesson, you should be able to:

- **Write SELECT statements to access data from more than one table using equality and nonequality joins**
- **View data that generally does not meet a join condition by using outer joins**
- **Join a table to itself**

Lesson Aim

This lesson covers how to obtain data from more than one table, using the different methods available.


Obtaining Data from Multiple Tables

EMP

EMPNO	ENAME	...	DEPTNO
7839	KING	...	10
7698	BLAKE	...	30
...			
7934	MILLER	...	10

DEPT

DEPTNO	DNAME	LOC
10	ACCOUNTING	NEW YORK
20	RESEARCH	DALLAS
30	SALES	CHICAGO
40	OPERATIONS	BOSTON



EMPNO	DEPTNO	LOC
7839	10	NEW YORK
7698	30	CHICAGO
7782	10	NEW YORK
7566	20	DALLAS
7654	30	CHICAGO
7499	30	CHICAGO
...		
14 rows selected.		

1-119

Copyright © Oracle Corporation, 1998. All rights reserved.

ORACLE®

Data from Multiple Tables

Sometimes you need to use data from more than one table. In the example above, the report displays data from two separate tables.

- EMPNO exists in the EMP table.
- DEPTNO exists in both the EMP and DEPT tables.
- LOC exists in DEPT table.

To produce the report, you need to link EMP and DEPT tables and access data from both of them.

Class Management Note

In the slide above, the DEPTNO column can come from either the EMP or the DEPT table.

What Is a Join?

Use a join to query data from more than one table.

```
SELECT  table.column, table.column
FROM    table1, table2
WHERE   table1.column1 = table2.column2;
```

- **Write the join condition in the WHERE clause.**
- **Prefix the column name with the table name when the same column name appears in more than one table.**

1-120

Copyright © Oracle Corporation, 1998. All rights reserved.

ORACLE®

Defining Joins

When data from more than one table in the database is required, a *join* condition is used. Rows in one table can be joined to rows in another table according to common values existing in corresponding columns, that is, primary and foreign key columns.

To display data from two or more related tables, write a simple join condition in the WHERE clause. In the syntax:

table.column denotes the table and column from which data is retrieved
table1.column1 = table2.column2 is the condition that joins (or relates) the tables together

Guidelines

- When writing a SELECT statement that joins tables, precede the column name with the table name for clarity and to enhance database access.
- If the same column name appears in more than one table, the column name must be prefixed with the table name.
- To join *n* tables together, you need a minimum of *(n-1)* join conditions. Therefore, to join four tables, a minimum of three joins are required. This rule may not apply if your table has a concatenated primary key, in which case more than one column is required to uniquely identify each row.

For more information, see

Oracle8 Server SQL Language Reference Manual, “SELECT.”



Cartesian Product

- **A Cartesian product is formed when:**
 - **A join condition is omitted**
 - **A join condition is invalid**
 - **All rows in the first table are joined to all rows in the second table**
- **To avoid a Cartesian product, always include a valid join condition in a WHERE clause.**

Cartesian Product

When a join condition is invalid or omitted completely, the result is a *Cartesian product* in which all combinations of rows will be displayed. All rows in the first table are joined to all rows in the second table.

A Cartesian product tends to generate a large number of rows, and its result is rarely useful. You should always include a valid join condition in a WHERE clause, unless you have a specific need to combine all rows from all tables.

Generating a Cartesian Product

EMP (14 rows)

EMPNO	ENAME	...	DEPTNO
7839	KING	...	10
7698	BLAKE	...	30
...			
7934	MILLER	...	10

DEPT (4 rows)

DEPTNO	DNAME	LOC
10	ACCOUNTING	NEW YORK
20	RESEARCH	DALLAS
30	SALES	CHICAGO
40	OPERATIONS	BOSTON

“Cartesian
product: →
14*4=56 rows”

ENAME	DNAME
-----	-----
KING	ACCOUNTING
BLAKE	ACCOUNTING
...	
KING	RESEARCH
BLAKE	RESEARCH
...	
56 rows selected.	

1-122

Copyright © Oracle Corporation, 1998. All rights reserved.

ORACLE®

Cartesian Product

A Cartesian product is generated if a join condition is omitted. The example on the slide displays employee name and department name from EMP and DEPT tables. Because no WHERE clause has been specified, all rows (14 rows) from the EMP table are joined with all rows (4 rows) in the DEPT table, thereby generating 56 rows in the output.

```
SQL> SELECT ename, dname
       2 FROM emp, dept;
```





```
ENAME      DNAME
-----
KING        ACCOUNTING
BLAKE       ACCOUNTING
...
KING        RESEARCH
BLAKE       RESEARCH
...
56 rows selected.
```

Class Management Note

Demo: l4cart.sql

Types of Joins

Equijoin Non-equijoin Outer join Self join



1-123

Copyright © Oracle Corporation, 1998. All rights reserved.

ORACLE®

Types of Joins

There are two main types of join conditions:

- Equijoins
- Non-equijoins

Additional join methods include the following:

- Outer joins
- Self joins
- Set operators

Note: Set operators are not covered in this course. They are covered in Advanced SQL and SQL*Plus.

Class Management Note

Do not get into details of all the types of joins. Explain each join one by one as is done in the following slides.

What Is an Equijoin?

EMP

EMPNO	ENAME	DEPTNO
7839	KING	10
7698	BLAKE	30
7782	CLARK	10
7566	JONES	20
7654	MARTIN	30
7499	ALLEN	30
7844	TURNER	30
7900	JAMES	30
7521	WARD	30
7902	FORD	20
7369	SMITH	20
...		
14 rows selected.		

DEPT

DEPTNO	DNAME	LOC
10	ACCOUNTING	NEW YORK
30	SALES	CHICAGO
10	ACCOUNTING	NEW YORK
20	RESEARCH	DALLAS
30	SALES	CHICAGO
30	SALES	CHICAGO
30	SALES	CHICAGO
30	SALES	CHICAGO
30	SALES	CHICAGO
20	RESEARCH	DALLAS
20	RESEARCH	DALLAS
...		
14 rows selected.		

Primary key Foreign key

1-124

Copyright © Oracle Corporation, 1998. All rights reserved.

ORACLE®

Equijoins

To determine the name of an employee's department, you compare the value in the DEPTNO column in the EMP table with the DEPTNO values in the DEPT table. The relationship between the EMP and DEPT tables is an *equijoin*—that is, values in the DEPTNO column on both tables must be equal. Frequently, this type of joins involve primary and foreign key complements.

Note: Equijoins are also called simple joins or inner joins.

Class Management Note

Explain the use of decision matrix for simplifying writing joins. For example, if you want to display the name and department number of all the employees who are in the same department as Smith, you can start by making the following decision tree:

Columns to Display	Originating Table	Condition
ename	emp	ename = 'SMITH'
deptno	emp	emp.deptno =
dname	dept	dept.deptno

Retrieving Records with Equijoins

```
SQL> SELECT emp.empno, emp.ename, emp.deptno,  
2          dept.deptno, dept.loc  
3 FROM emp, dept  
4 WHERE emp.deptno=dept.deptno;
```

```
EMPNO ENAME DEPTNO DEPTNO LOC  
-----  
7839 KING      10      10 NEW YORK  
7698 BLAKE      30      30 CHICAGO  
7782 CLARK      10      10 NEW YORK  
7566 JONES      20      20 DALLAS  
...  
14 rows selected.
```

1-125

Copyright © Oracle Corporation, 1998. All rights reserved.

ORACLE®

Retrieving Records with Equijoins

In the above example:

- The SELECT clause specifies the column names to retrieve:
 - employee name, employee number, and department number, which are columns in the EMP table
 - department number, department name, and location, which are columns in the DEPT table
- The FROM clause specifies the two tables that the database must access:
 - EMP table
 - DEPT table
- The WHERE clause specifies how the tables are to be joined:
EMP.DEPTNO=DEPT.DEPTNO

Because the DEPTNO column is common to both tables, it must be prefixed by the table name to avoid ambiguity.

Qualifying Ambiguous Column Names

- **Use table prefixes to qualify column names that are in multiple tables.**
- **Improve performance by using table prefixes.**
- **Distinguish columns that have identical names but reside in different tables by using column aliases.**

1-126

Copyright © Oracle Corporation, 1998. All rights reserved.

ORACLE®

Qualifying Ambiguous Column Names

You need to qualify the names of the columns in the WHERE clause with the table name to avoid ambiguity. Without the table prefixes, the DEPTNO column could be from either the DEPT table or the EMP table. It is necessary to add the table prefix to execute your query.

If there are no common column names between the two tables, there is no need to qualify the columns. However, you will gain improved performance by using the table prefix because you tell the Oracle8 Server exactly where to go to find columns.

The requirement to qualify ambiguous column names is also applicable to columns that may be ambiguous in other clauses such as SELECT clause or ORDER BY clause.



Class Management Note

Demo: *l4loc.sql*

Additional Search Conditions

Using the AND Operator

EMP

EMPNO	ENAME	DEPTNO
-----	-----	-----
7839	KING	10
7698	BLAKE	30
7782	CLARK	10
7566	JONES	20
7654	MARTIN	30
7499	ALLEN	30
7844	TURNER	30
7900	JAMES	30
7521	WARD	30
7902	FORD	20
7369	SMITH	20
...		
14 rows selected.		

DEPT

DEPTNO	DNAME	LOC
-----	-----	-----
10	ACCOUNTING	NEW YORK
30	SALES	CHICAGO
10	ACCOUNTING	NEW YORK
20	RESEARCH	DALLAS
30	SALES	CHICAGO
30	SALES	CHICAGO
30	SALES	CHICAGO
30	SALES	CHICAGO
30	SALES	CHICAGO
20	RESEARCH	DALLAS
20	RESEARCH	DALLAS
...		
14 rows selected.		

1-127

Copyright © Oracle Corporation, 1998. All rights reserved.

ORACLE®

Additional Search Conditions

In addition to the join, you may have additional criteria for your WHERE clause. For example, to display employee King's employee number, name, department number, and department location, you need an additional condition in the WHERE clause.

```
SQL> SELECT empno, ename, emp.deptno, loc
2 FROM emp, dept;
3 WHERE emp.deptno = dept.deptno
4 AND INITCAP(ename) = 'King';
```

EMPNO	ENAME	DEPTNO	LOC
-----	-----	-----	-----
7839	KING	10	NEW YORK

Using Table Aliases

Simplify queries by using table aliases.

```
SQL> SELECT emp.empno, emp.ename, emp.deptno,  
2          dept.deptno, dept.loc  
3 FROM    emp, dept  
4 WHERE   emp.deptno=dept.deptno;
```

```
SQL> SELECT e.empno, e.ename, e.deptno,  
2          d.deptno, d.loc  
3 FROM    emp e, dept d  
4 WHERE   e.deptno=d.deptno;
```

Table Aliases

Qualifying column names with table names can be very time consuming, particularly if table names are lengthy. You can use table *aliases* instead of table names. Just as a column alias gives a column another name, a table alias gives a table another name. Table aliases help to keep SQL code smaller, therefore using less memory.

Notice how table aliases are identified in the FROM clause in the example. The table name is specified in full, followed by a space and then the table alias. The EMP table has been given an alias of E, whereas the DEPT table has an alias of D.

Guidelines

- Table aliases can be up to 30 characters in length, but the shorter they are the better.
- If a table alias is used for a particular table name in the FROM clause, then that table alias must be substituted for the table name throughout the SELECT statement.
- Table aliases should be meaningful.
- The table alias is valid only for the current SELECT statement.

Joining More Than Two Tables

CUSTOMER

NAME	CUSTID
-----	-----
JOCKSPORTS	100
TKB SPORT SHOP	101
VOLLYRITE	102
JUST TENNIS	103
K+T SPORTS	105
SHAPE UP	106
WOMENS SPORTS	107
...	...
9 rows selected.	

ORD

CUSTID	ORDID
-----	-----
101	610
102	611
104	612
106	601
102	602
106	
106	
...	...
21 rows selected.	

ITEM

ORDID	ITEMID
-----	-----
610	3
611	1
612	1
601	1
602	1
...	...
64 rows selected.	

1-129

Copyright © Oracle Corporation, 1998. All rights reserved.

ORACLE®

Additional Search Conditions

Sometimes you may need to join more than two tables. For example, to display the name, the orders placed, the item numbers, the total for each item, and the total for each order for customer TKB SPORT SHOP, you will have to join the CUSTOMER, ORD, and ITEM tables.

```
SQL> SELECT c.name, o.ordid, i.itemid, i.itemtot, o.total
2  FROM   customer c, ord o, item i
3  WHERE  c.custid = o.custid
4  AND    o.ordid = i.ordid
5  AND    c.name = 'TKB SPORT SHOP';
```

NAME	ORDID	ITEMID	ITEMTOT	TOTAL
-----	-----	-----	-----	-----
TKB SPORT SHOP	610	3	58	101.4
TKB SPORT SHOP	610	1	35	101.4
TKB SPORT SHOP	610	2	8.4	101.4

Non-Equi Joins

EMP

EMPNO	ENAME	SAL
7839	KING	5000
7698	BLAKE	2850
7782	CLARK	2450
7566	JONES	2975
7654	MARTIN	1250
7499	ALLEN	1600
7844	TURNER	1500
7900	JAMES	950
...		
14 rows selected.		

SALGRADE

GRADE	LOSAL	HISAL
1	700	1200
2	1201	1400
3	1401	2000
4	2001	3000
5	3001	9999

“salary in the EMP table is between low salary and high salary in the SALGRADE table”

Non-Equi Joins

The relationship between the EMP table and the SALGRADE table is a non-equi join, meaning that no column in the EMP table corresponds directly to a column in the SALGRADE table. The relationship between the two tables is that the SAL column in the EMP table is between the LOSAL and HISAL column of the SALGRADE table. The relationship is obtained using an operator other than equal (=).

Retrieving Records with Non-Equi Joins

```
SQL>  SELECT    e.ename, e.sal, s.grade
      2  FROM      emp e, salgrade s
      3  WHERE     e.sal
      4  BETWEEN   s.losal AND s.hisal;
```

ENAME	SAL	GRADE
-----	-----	-----
JAMES	950	1
SMITH	800	1
ADAMS	1100	1
...		
14 rows selected.		

Non-Equi Joins (continued)

The example above creates a non-equi join to evaluate an employee's salary grade. The salary must be *between* any pair of the low and high salary ranges.

It is important to note that all employees appear exactly once when this query is executed. No employee is repeated in the list. There are two reasons for this:

- None of the rows in the salary grade table contain grades that overlap. That is, the salary value for an employee can only lie between the low salary and high salary values of one of the rows in the salary grade table.
- All of the employees' salaries lie within the limits provided by the salary grade table. That is, no employee earns less than the lowest value contained in the LOSAL column or more than the highest value contained in the HISAL column

Note: Other operators such as <= and >= could be used, but BETWEEN is the simplest. Remember to specify the low value first and the high value last when using BETWEEN. Table aliases have been specified for performance reasons, not because of possible ambiguity.


Outer Joins

EMP

ENAME	DEPTNO
-----	-----
KING	10
BLAKE	30
CLARK	10
JONES	20
...	

DEPT

DEPTNO	DNAME
-----	-----
10	ACCOUNTING
30	SALES
10	ACCOUNTING
20	RESEARCH
...	
40	OPERATIONS



**No employee in the
OPERATIONS department**

Returning Records with No Direct Match with Outer Joins

If a row does not satisfy a join condition, the row will not appear in the query result. For example, in the equijoin condition of EMP and DEPT tables, department OPERATIONS does not appear because no one works in that department.

```
SQL> SELECT e.ename, e.deptno, d.dname
       2 FROM   emp e, dept d
       3 WHERE  e.deptno = d.deptno;
```

ENAME	DEPTNO	DNAME
-----	-----	-----
KING	10	ACCOUNTING
BLAKE	30	SALES
CLARK	10	ACCOUNTING
JONES	20	RESEARCH
...		
ALLEN	30	SALES
TURNER	30	SALES
JAMES	30	SALES
...		

14 rows selected.

Outer Joins

- You use an outer join to see rows that do not usually meet the join condition.
- Outer join operator is the plus sign (+).

```
SELECT table.column, table.column
FROM   table1, table2
WHERE  table1.column(+) = table2.column;
```

```
SELECT table.column, table.column
FROM   table1, table2
WHERE  table1.column = table2.column(+);
```

Returning Records with No Direct Match with Outer Joins

The missing row(s) can be returned if an *outer join* operator is used in the join condition. The operator is a plus sign enclosed in parentheses (+), and it is *placed on the “side” of the join that is deficient in information*. This operator has the effect of creating one or more null rows, to which one or more rows from the nondeficient table can be joined.

In the syntax:

table1.column =
table2.column (+)

is the condition that joins (or relates) the tables together.
is the outer join symbol; it can be placed on either side of the WHERE clause condition, but not on both sides. Place the outer join symbol following the name of the table without the matching rows.

Class Management Note

Demo: *l4ejoin.sql*

Using Outer Joins

```
SQL> SELECT    e.ename, d.deptno, d.dname
  2  FROM      emp e, dept d
  3  WHERE     e.deptno(+) = d.deptno
  4  ORDER BY  e.deptno;
```

```
ENAME          DEPTNO DNAME
-----
KING            10  ACCOUNTING
CLARK           10  ACCOUNTING
...
                40  OPERATIONS
15 rows selected.
```

Returning Records with No Direct Match with Outer Joins

The example above displays numbers and names for all the departments. The OPERATIONS department, which does not have any employees, is also displayed.

Outer Join Restrictions

- The outer join operator can appear only on *one* side of the expression—the side that has information missing. It returns those rows from one table that have no direct match in the other table.
- A condition involving an outer join cannot use the IN operator or be linked to another condition by the OR operator.

Class Management Note

Demo: *l4ojoin.sql*

Self Joins

EMP (WORKER)

EMPNO	ENAME	MGR
-----	-----	----
7839	KING	
7698	BLAKE	7839
7782	CLARK	7839
7566	JONES	7839
7654	MARTIN	7698
7499	ALLEN	7698

EMP (MANAGER)

EMPNO	ENAME
-----	-----
7839	KING
7839	KING
7839	KING
7698	BLAKE
7698	BLAKE



“MGR in the WORKER table is equal to EMPNO in the MANAGER table”

Joining a Table to Itself

Sometimes you need to join a table to itself. To find the name of each employee's manager you need to join the EMP table to itself. For example, to find the name of Blake's manager, you need to:

- Find Blake in the EMP table by looking at the ENAME column.
- Find the manager number for Blake by looking at the MGR column. Blake's manager number is 7839.
- Find the name of the manager with EMPNO 7839 by looking at the ENAME column. King's employee number is 7839. So, King is Blake's manager.

In this process, you look in the table twice. The first time you look in the table to find Blake in the ENAME column and MGR value of 7839. The second time you look in the EMPNO column to find 7839 and the ENAME column to find King.

Joining a Table to Itself

```
SQL> SELECT worker.ename || ' works for ' || manager.ename  
2 FROM emp worker, emp manager  
3 WHERE worker.mgr = manager.empno;
```

```
WORKER.ENAME || 'WORKSFOR' || MANAG  
-----  
BLAKE works for KING  
CLARK works for KING  
JONES works for KING  
MARTIN works for BLAKE  
...  
13 rows selected.
```

1-136

Copyright © Oracle Corporation, 1998. All rights reserved.

ORACLE®

Joining a Table to Itself (continued)

The above example joins the EMP table to itself. To simulate two tables in the FROM clause, there are two aliases, namely WORKER and MANAGER, for the same table, EMP.

In this example, the WHERE clause contains the join that means “where a worker’s manager number matches the employee number for the manager.”

Class Management Note





Point out the following to the students:

- The column heading in the result of the above query seems meaningless. A meaningful column alias should have been used instead.
- There are only 13 rows in the output, but there are 14 rows in the EMP table. This occurs because employee King, who is the president, does not have a manager.

Summary

```
SELECT    table.column, table.column
FROM      table1, table2
WHERE     table1.column1 = table2.column2;
```

Equijoin Non-equijoin Outer join Self join



Summary

There are multiple ways to join tables. The common thread, though, is that you want to link them through a condition in the WHERE clause. The method you choose will be based on the required result and the data structures that you are using.

```
SELECT    table.column, table.column
FROM      table1, table2
WHERE     table1.column1 = table2.column2;
```

Practice Overview

- **Joining tables using an equijoin**
- **Performing outer and self joins**
- **Adding additional conditions**

Practice Overview

This practice is intended to give you practical experience in extracting data from more than one table. You will be required to join and restrict rows in the **WHERE** clause.

1

Aggregating Data Using Group Functions

Copyright © Oracle Corporation, 1998. All rights reserved. ORACLE®

Schedule:	Timing	Topic
	35 minutes	Lecture
	40 minutes	Practice
	75 minutes	Total

Objectives

At the end of this lesson, you should be able to:

- **Identify the available group functions**
- **Describe the use of group functions**
- **Group data using the GROUP BY clause**
- **Include or exclude grouped rows by using the HAVING clause**

Lesson Aim

This lesson further addresses functions. It focuses on obtaining summary information, such as averages, for groups of rows. It discusses how to group rows in a table into smaller sets and how to specify search criteria for groups of rows.

What Are Group Functions?

Group functions operate on sets of rows to give one result per group.

EMP

DEPTNO	SAL
10	2450
10	5000
10	1300
20	800
20	1100
20	3000
20	3000
20	2975
30	1600
30	2850
30	1250
30	950
30	1500
30	1250

“maximum
salary in
the EMP table”

MAX (SAL)

5000

1-147

Copyright © Oracle Corporation, 1998. All rights reserved.

ORACLE®

Group Functions

Unlike single-row functions, group functions operate on sets of rows to give one result per group. These sets may be the whole table or the table split into groups.

Using AVG and SUM Functions

You can use AVG and SUM for numeric data.

```
SQL> SELECT  AVG(sal), MAX(sal),  
2           MIN(sal), SUM(sal)  
3 FROM      emp  
4 WHERE     job LIKE 'SALES%';
```

AVG(SAL)	MAX(SAL)	MIN(SAL)	SUM(SAL)
1400	1600	1250	5600

Group Functions

You can use AVG, SUM, MIN, and MAX functions against columns that can store numeric data. The example above displays the average, highest, lowest, and sum of monthly salaries for all salesmen.

Using MIN and MAX Functions

You can use MIN and MAX for any datatype.

```
SQL> SELECT  MIN(hiredate), MAX(hiredate)
           2  FROM    emp;
```

```
MIN(HIRED) MAX(HIRED)
-----
17-DEC-80 12-JAN-83
```

Group Functions (continued)

You can use MAX and MIN functions for any datatype. The example above displays the most junior and most senior employee.

The example below displays the employee name that is first and the employee name that is the last in an alphabetized list of all employees.

```
SQL> SELECT  MIN(ename), MAX(ename)
           2  FROM    emp;
```

```
MIN(ENAME) MAX(ENAME)
```

Note: AVG and SUM functions can be used only with numeric datatypes.

```
ADAMS      WARD
```

Using the COUNT Function

COUNT(*) returns the number of rows in a table.

```
SQL> SELECT  COUNT ( * )  
2  FROM      emp  
3  WHERE     deptno = 30;
```

```
COUNT ( * )  
-----  
6
```

The COUNT Function

The COUNT function has two formats:

- COUNT(*)
- COUNT(*expr*).

COUNT(*) returns the number of rows in a table, including duplicate rows and rows containing null values.

In contrast, COUNT(*expr*) returns the number of nonnull rows in the column identified by *expr*.

The example above displays the number of employees in department 30.

Class Management Note

Demo: *l5count1.sql*, *l5count2.sql*

Using the COUNT Function

COUNT(*expr*) returns the number of nonnull rows.

```
SQL> SELECT    COUNT ( comm )
      2  FROM      emp
      3  WHERE    deptno = 30;
```

```
COUNT ( COMM )
-----
                4
```

The COUNT Function (continued)

The example above displays the number of employees in department 30 who can earn a commission. Notice that the result gives the total number of rows to be four because two employees in department 30 cannot earn a commission and contain a null value in the COMM column.

Example

Display the number of departments in the EMP table.

```
SQL> SELECT    COUNT ( deptno )
      2  FROM      emp;
```

```
COUNT ( DEPTNO )
-----
                14
```

Display the number of distinct departments in the EMP table.

```
SQL> SELECT    COUNT ( DISTINCT ( deptno ) )
      2  FROM      emp;
```

```
COUNT ( DISTINCT ( DEPTNO ) )
-----
                3
```

Group Functions and Null Values

Group functions ignore null values in the column.

```
SQL> SELECT AVG(comm)
      2 FROM emp;
```

```
AVG (COMM)
-----
          550
```

Group Functions and Null Values

All group functions except COUNT (*) ignore null values in the column. In the above example, the average is calculated based *only* on the rows in the table where a valid value is stored in the COMM column. The average is calculated as total commission being paid to all employees divided by the number of employees receiving commission (4).

Using the NVL Function with Group Functions

The NVL function forces group functions to include null values.

```
SQL> SELECT AVG(NVL(comm,0))  
2 FROM emp;
```

```
AVG(NVL(COMM,0))  
-----  
157.14286
```

Group Functions and Null Values (continued)

The NVL function forces group functions to include null values. In the above example, the average is calculated based on *all* rows in the table regardless of whether null values are stored in the COMM column. The average is calculated as total commission being paid to all employees divided by the total number of employees in the company (14).

Creating Groups of Data

EMP

DEPTNO	SAL
10	2450
10	5000
10	1300
20	800
20	1100
20	3000
20	3000
20	2975
30	1600
30	2850
30	1250
30	950
30	1500
30	1250

2916.6667

2175

1566.6667

“average
salary
in EMP
table
for each
department”

DEPTNO	AVG (SAL)
10	2916.6667
20	2175
30	1566.6667

1-156

Copyright © Oracle Corporation, 1998. All rights reserved.

ORACLE®

Groups of Data

Until now, all group functions have treated the table as one large group of information. At times, you need to divide the table of information into smaller groups. This can be done by using the GROUP BY clause.

Creating Groups of Data:

GROUP BY Clause

```
SELECT      column, group_function
FROM        table
[WHERE      condition]
[GROUP BY   group_by_expression]
[ORDER BY   column];
```

Divide rows in a table into smaller groups by using the **GROUP BY** clause.

The GROUP BY Clause

You can use the GROUP BY clause to divide the rows in a table into groups. You can then use the group functions to return summary information for each group.

In the syntax:

<i>group_by_expression</i>	specifies columns whose values determine the basis for grouping rows
----------------------------	--

Guidelines

- If you include a group function in a SELECT clause, you cannot select individual results as well *unless* the individual column appears in the GROUP BY clause. You will receive an error message if you fail to include the column list.
- Using a WHERE clause, you can preexclude rows before dividing them into groups.
- You must include the *columns* in the GROUP BY clause.
- You cannot use the column alias in the GROUP BY clause.
- By default, rows are sorted by ascending order of the columns included in the GROUP BY list. You can override this by using the ORDER BY clause.

Using the GROUP BY Clause

All columns in the SELECT list that are not in group functions must be in the GROUP BY clause.

```
SQL> SELECT    deptno, AVG(sal)
  2  FROM      emp
  3  GROUP BY  deptno;
```

DEPTNO	AVG (SAL)
10	2916.6667
20	2175
30	1566.6667

The GROUP BY Clause (continued)

When using the GROUP BY clause, make sure that all columns in the SELECT list that are not in the group functions are included in the GROUP BY clause. The above example displays the department number and the average salary for each department. Here is how the SELECT statement above, containing a GROUP BY clause, is evaluated:

- The SELECT clause specifies the columns to be retrieved:
 - Department number column in the EMP table
 - The average of all the salaries in the group you specified in the GROUP BY clause
- The FROM clause specifies the tables that the database must access: the EMP table.
- The WHERE clause specifies the rows to be retrieved. Since there is no WHERE clause, by default all rows are retrieved.
- The GROUP BY clause specifies how the rows should be grouped. The rows are being grouped by department number, so the AVG function that is being applied to the salary column will calculate the *average salary for each department*.

Using the GROUP BY Clause

The GROUP BY column does not have to be in the SELECT list.

```
SQL> SELECT    AVG(sal)
      2  FROM      emp
      3  GROUP BY deptno;
```

```
AVG(SAL)
-----
2916.6667
      2175
1566.6667
```

1-159

Copyright © Oracle Corporation, 1998. All rights reserved.

ORACLE®

The GROUP BY Clause (continued)

The GROUP BY column does not have to be in the SELECT clause. For example, the above SELECT statement displays the average salaries for each department without displaying the respective department numbers. However, without the department numbers, the results do not look meaningful.

You can use the group function in the ORDER BY clause.

```
SQL> SELECT    deptno, AVG(sal)
      2  FROM      emp
      3  GROUP BY deptno
      4  ORDER BY  AVG(sal);
```

```
DEPTNO    AVG(SAL)
-----
10        1566.6667
20        2175
30        2916.6667
```

Class Management Note

Demonstrate the query with and without the DEPTNO in the SELECT statement.

Grouping by More Than One Column

EMP

DEPTNO	JOB	SAL
10	MANAGER	2450
10	PRESIDENT	5000
10	CLERK	1300
20	CLERK	800
20	CLERK	1100
20	ANALYST	3000
20	ANALYST	3000
20	MANAGER	2975
30	SALESMAN	1600
30	MANAGER	2850
30	SALESMAN	1250
30	CLERK	950
30	SALESMAN	1500
30	SALESMAN	1250

“sum salaries in the EMP table for each job, grouped by department”

DEPTNO	JOB	SUM(SAL)
10	CLERK	1300
10	MANAGER	2450
10	PRESIDENT	5000
20	ANALYST	6000
20	CLERK	1900
20	MANAGER	2975
30	CLERK	950
30	MANAGER	2850
30	SALESMAN	5600

1-160

Copyright © Oracle Corporation, 1998. All rights reserved.

ORACLE®

Groups Within Groups

Sometimes there is a need to see results for groups within groups. The slide above shows a report that displays the total salary being paid to each job title, within each department.

The EMP table is grouped first by department number and then within that grouping it is grouped by job title. For example, the two clerks in department 20 are grouped together and a single result (total salary) is produced for all salesmen within the group.

Class Management Note

Demo: *l5order1.sql, l5order2.sql*

Using the GROUP BY Clause on Multiple Columns

```
SQL> SELECT deptno, job, sum(sal)
2 FROM emp
3 GROUP BY deptno, job;
```

DEPTNO	JOB	SUM(SAL)
10	CLERK	1300
10	MANAGER	2450
10	PRESIDENT	5000
20	ANALYST	6000
20	CLERK	1900
...		

9 rows selected.

1-161

Copyright © Oracle Corporation, 1998. All rights reserved.

ORACLE®

Groups Within Groups (continued)

You can return summary results for groups and subgroups by listing more than one GROUP BY column. You can determine the default sort order of the results by the order of the columns in the GROUP BY clause. Here is how the SELECT statement above, containing a GROUP BY clause, is evaluated:

- The SELECT clause specifies the column to be retrieved:
 - Department number in the EMP table
 - Job title in the EMP table
 - The sum of all the salaries in the group you specified in the GROUP BY clause
- The FROM clause specifies the tables that the database must access: the EMP table.
- The GROUP BY clause specifies how you must group the rows:
 - First, the rows are grouped by department number.
 - Second, within the department number groups, the rows are grouped by job title.

So the SUM function is being applied to the salary column for all job titles within each department number group.

Illegal Queries Using Group Functions

Any column or expression in the SELECT list that is not an aggregate function must be in the GROUP BY clause.

```
SQL> SELECT deptno, COUNT(ename)
      2 FROM emp;
```

```
SELECT deptno, COUNT(ename)
      *
ERROR at line 1:
ORA-00937: not a single-group group function
```

1-162

Copyright © Oracle Corporation, 1998. All rights reserved.

ORACLE®

Illegal Queries Using Group Functions

Whenever you use a mixture of individual items (DEPTNO) and group functions (COUNT) in the same SELECT statement, you must include a GROUP BY clause that specifies the individual items (in this case, DEPTNO). If the GROUP BY clause is missing, then the error message “not a single-group group function” appears and an asterisk (*) points to the offending column. You can correct the above error by adding the GROUP BY clause.

```
SQL> SELECT deptno, COUNT(ename)
      2 FROM emp
      3 GROUP BY deptno;
```

```
DEPTNO COUNT ( ENAME )
-----
```

Any column or expression in the SELECT list that is not an aggregate function must be in the GROUP BY clause.

Class Management Note

5

Demo: l5error.sql

6



Illegal Queries Using Group Functions

- You cannot use the WHERE clause to restrict groups.
- You use the HAVING clause to restrict groups.

```
SQL> SELECT      deptno, AVG(sal)
  2  FROM        emp
  3  WHERE        AVG(sal) > 2000
  4  GROUP BY    deptno;
```

```
WHERE AVG(sal) > 2000
      *
```

```
ERROR at line 3:
```

```
ORA-00934: group function is not allowed here
```

1-163

Copyright © Oracle Corporation, 1998. All rights reserved.

ORACLE®

Illegal Queries Using Group Functions (continued)

The WHERE clause cannot be used to restrict groups. The above SELECT statement results in an error because it uses the WHERE clause to restrict the display of average salaries of those departments that have an average salary of greater than \$2000.

You can correct the above error by using the HAVING clause to restrict groups.

```
SQL> SELECT      deptno, AVG(sal)
  2  FROM        emp
  3  GROUP BY    deptno
  4  HAVING      AVG(sal) > 2000;
```

DEPTNO	AVG (SAL)
10	2916.6667
20	2175

Excluding Group Results

EMP

DEPTNO	SAL
10	2450
10	5000
10	1300
20	800
20	1100
20	3000
20	3000
20	2975
30	1600
30	2850
30	1250
30	950
30	1500
30	1250

5000

3000

2850

**“maximum
salary
per department
greater than
\$2900”**

DEPTNO	MAX (SAL)
10	5000
20	3000

1-164

Copyright © Oracle Corporation, 1998. All rights reserved.

ORACLE®

Restricting Group Results

In the same way that you use the WHERE clause to restrict the rows that you select, you use the HAVING clause to restrict groups. To find the maximum salary of each department, but show only the departments that have a maximum salary of more than \$2900, you need to do the following two things:

1. Find the average salary for each department by grouping by department number.
2. Restrict the groups to those departments with a maximum salary greater than \$2900.

Excluding Group Results: HAVING Clause

Use the HAVING clause to restrict groups

- Rows are grouped.
- The group function is applied.
- Groups matching the HAVING clause are displayed.

```
SELECT      column, group_function
FROM        table
[WHERE      condition]
[GROUP BY   group_by_expression]
[HAVING     group_condition]
[ORDER BY   column];
```

1-165

Copyright © Oracle Corporation, 1998. All rights reserved.

ORACLE®

The HAVING Clause

You use the HAVING clause to specify which groups are to be displayed. Therefore, you further restrict the groups on the basis of aggregate information.

In the syntax:

group_condition restricts the groups of rows returned to those groups for which the specified condition is TRUE

The Oracle8 Server performs the following steps when you use the HAVING clause:

- Rows are grouped.
- The group function is applied to the group.
- The groups that match the criteria in the HAVING clause are displayed.

The HAVING clause can precede the GROUP BY clause, but it is recommended that you place the GROUP BY clause first because it is more logical. Groups are formed and group functions are calculated before the HAVING clause is applied to the groups in the SELECT list.



Using the HAVING Clause

```
SQL> SELECT    deptno, max(sal)
  2  FROM      emp
  3  GROUP BY  deptno
  4  HAVING    max(sal)>2900;
```

DEPTNO	MAX(SAL)
10	5000
20	3000

1-166

Copyright © Oracle Corporation, 1998. All rights reserved.

ORACLE®

The HAVING Clause (continued)

The above example displays department numbers and maximum salary for those departments whose maximum salary is greater than \$2900.

You can use the GROUP BY clause without using a group function in the SELECT list.

If you restrict rows based on the result of a group function, you must have a GROUP BY clause as well as the HAVING clause.

```
SQL> SELECT    deptno
  2  FROM      emp
  3  GROUP BY  deptno
  4  HAVING    MAX(sal) > 2900;
```

DEPTNO
10
20

Using the HAVING Clause

```
SQL> SELECT      job, SUM(sal) PAYROLL
  2  FROM          emp
  3  WHERE         job NOT LIKE 'SALES%'
  3  GROUP BY     job
  4  HAVING        SUM(sal)>5000
  5  ORDER BY     SUM(sal);
```

JOB	PAYROLL
ANALYST	6000
MANAGER	8275

The HAVING Clause (continued)

The above example displays the job title and total monthly salary for each job title with a total payroll exceeding \$5000. The example excludes salesmen and sorts the list by the total monthly salary.

Class Management Note

Demo: *l5job1.sql, l5job2.sql*

Nesting Group Functions

Display the maximum average salary.

```
SQL> SELECT      max(avg(sal))  
  2  FROM        emp  
  3  GROUP BY deptno;
```

```
MAX(AVG(SAL))  
-----  
      2916.6667
```

Nesting Group Functions

Group functions can be nested. The above example displays the maximum average salary.

Summary

```
SELECT      column, group_function
FROM        table
[WHERE      condition]
[GROUP BY   group_by_expression]
[HAVING     group_condition]
[ORDER BY   column];
```

Summary

Seven group functions are available in SQL:

- AVG
- COUNT
- MAX
- MIN
- SUM
- STDDEV
- VARIANCE

You can create subgroups by using the GROUP BY clause. Groups can be excluded using the HAVING clause.

Place the HAVING and GROUP BY clauses after the WHERE clause in a statement. Place the ORDER BY clause last.

Oracle8 Server evaluates the clauses in the following order:

- If the statement contains a WHERE clause, the server establishes the candidate rows.
- The server identifies the groups specified in the GROUP BY clause.
- The HAVING clause further restricts result groups that do not meet the group criteria in the HAVING clause.

Practice Overview

- **Showing different queries that use group functions**
- **Grouping by rows to achieve more than one result**
- **Excluding groups by using the HAVING clause**

1-170

Copyright © Oracle Corporation, 1998. All rights reserved.

ORACLE®

Practice Overview

At the end of this practice, you should be familiar with using group functions and selecting groups of data.

Paper-Based Questions

For questions 1-3 circle either True or False.

Note: Column aliases are used for the queries.

1

Subqueries

Copyright © Oracle Corporation, 1998. All rights reserved. ORACLE®

Schedule:	Timing	Topic
	25 minutes	Lecture
	30 minutes	Practice
	55 minutes	Total

Objectives

At the end of this lesson, you should be able to:


- **Describe the types of problems that subqueries can solve**
- **Define subqueries**
- **List the types of subqueries**
- **Write single-row and multiple-row subqueries**

Lesson Aim

In this lesson you will learn about more advanced features of the SELECT statement. You can write subqueries in the WHERE clause of another SQL statement to obtain values based on an unknown conditional value. This lesson covers single-row subqueries and multiple-row subqueries.

Using a Subquery to Solve a Problem

“Who has a salary greater than Jones’s?”



1-177

Copyright © Oracle Corporation, 1998. All rights reserved.

ORACLE®

Using a Subquery to Solve a Problem

Suppose you want to write a query to find out who earns a salary greater than Jones's salary.

To solve this problem, you need *two* queries: one query to find what Jones earns and a second query to find who earns more than that amount.

You can solve this problem by combining the two queries, placing one query *inside* the other query.

An inner query or the *subquery* returns a value that is used by the outer query or the main query. Using a subquery is equivalent to performing two sequential queries and using the result of the first query as the search value in the second query.

Subqueries

```
SELECT    select_list
FROM      table
WHERE     expr operator
          (SELECT    select_list
           FROM      table);
```

- The subquery (inner query) executes once before the main query.
- The result of the subquery is used by the main query (outer query).

Subqueries

A subquery is a SELECT statement that is embedded in a clause of another SELECT statement. You can build powerful statements out of simple ones by using subqueries. They can be very useful when you need to select rows from a table with a condition that depends on the data in the table itself.

You can place the subquery in a number of SQL clauses:

- WHERE clause
- HAVING clause
- FROM clause

In the syntax:

operator includes a comparison operator such as >, =, or IN

Note: Comparison operators fall into two classes: single-row operators (>, =, >=, <, <>, <=) and multiple-row operators (IN, ANY, ALL) .

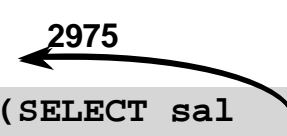
The subquery is often referred to as a nested SELECT, sub-SELECT, or inner SELECT statement. The subquery generally executes first, and its output is used to complete the query condition for the main or outer query.

Class Management Note

Additionally, subqueries can be placed in the CREATE VIEW statement, CREATE TABLE statement, UPDATE clause, INTO clause of an INSERT statement, and SET clause of an UPDATE statement.

Using a Subquery

```
SQL> SELECT  ename
      2  FROM    emp
      3  WHERE   sal > 2975
      4
      5          (SELECT sal
      6          FROM    emp
            WHERE   empno=7566) ;
```

A curved arrow points from the value '2975' to the subquery '(SELECT sal FROM emp WHERE empno=7566) ;' in the SQL statement.

ENAME

KING

FORD

SCOTT

Using a Subquery

In the slide, the inner query determines the salary of employee 7566. The outer query takes the result of the inner query and uses this result to display all the employees who earn more than this amount.

Class Management Note

Execute the subquery (inner query) on its own first to show the value that the subquery returns. Then execute the outer query using the result returned by the inner query. Finally, execute the entire query (containing the subquery) and show that the result is the same.

Guidelines for Using Subqueries

- **Enclose subqueries in parentheses.**
- **Place subqueries on the right side of the comparison operator.**
- **Do not add an ORDER BY clause to a subquery.**
- **Use single-row operators with single-row subqueries.**
- **Use multiple-row operators with multiple-row subqueries.**

Guidelines for Using Subqueries

- A subquery must be enclosed in parentheses.
- A subquery must appear on the right side of the comparison operator.
- Subqueries cannot contain an ORDER BY clause. You can have only one ORDER BY clause for a SELECT statement, and if specified it must be the last clause in the main SELECT statement.
- Two classes of comparison operators are used in subqueries: single-row operators and multiple-row operators.


Class Management Note

A subquery can execute multiple times in correlated subqueries, which are not included in this course. Refer students to the *Advanced SQL and SQL*Plus* course for this topic.

Students may ask how many subqueries can be written. The Oracle8 Server imposes no limit on the number of subqueries. The limit is related to the buffer size that the query uses.

Types of Subqueries


- **Single-row subquery**



- **Multiple-row subquery**



- **Multiple-column subquery**



1-181

Copyright © Oracle Corporation, 1998. All rights reserved.

ORACLE®

Types of Subqueries

- Single-row subqueries: Queries that return only one row from the inner SELECT statement
- Multiple-row subqueries: Queries that return more than one row from the inner SELECT statement
- Multiple-column subqueries: Queries that return more than one column from the inner SELECT statement

Single-Row Subqueries

- Return only one row
- Use single-row comparison operators

Operator	Meaning
=	Equal to
>	Greater than
>=	Greater than or equal to
<	Less than
<=	Less than or equal to
<>	Not equal to

1-182

Copyright © Oracle Corporation, 1998. All rights reserved.

ORACLE®

Single Row Subqueries

A *single-row subquery* is one that returns one row from the inner SELECT statement. This type of subquery uses a single-row operator. The slide gives a list of single-row operators.

Example

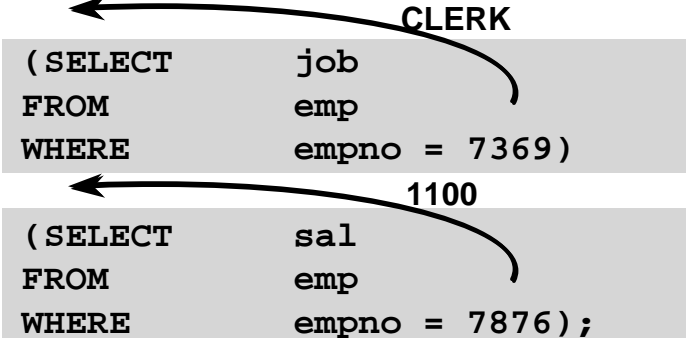
Display the employees whose job title is the same as that of employee 7369.

```
SQL> SELECT  ename, job
2 FROM      emp
3 WHERE     job =
4           (SELECT  job
5                FROM    emp
6                WHERE   empno = 7369);
```

ENAME	JOB
-----	-----
JAMES	CLERK
SMITH	CLERK
ADAMS	CLERK
MILLER	CLERK

Executing Single-Row Subqueries

```
SQL> SELECT  ename, job
2  FROM      emp
3  WHERE     job =
4             ( SELECT  job
5               FROM    emp
6               WHERE   empno = 7369 )
7  AND       sal >
8             ( SELECT  sal
9               FROM    emp
10              WHERE   empno = 7876 );
```



ENAME	JOB
-----	-----
MILLER	CLERK

Executing Single Row Subqueries

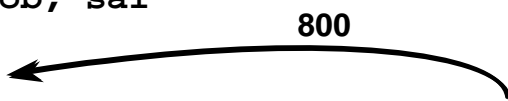
A SELECT statement can be considered as a query block. The example above displays employees whose job title is the same as that of employee 7369 and whose salary is greater than that of employee 7876.

The example consists of three query blocks: the outer query and two inner queries. The inner query blocks are executed first, producing the query results: CLERK and 1100, respectively. The outer query block is then processed and uses the values returned by the inner queries to complete its search conditions.

Both the inner queries return single values (CLERK and 1100, respectively), so this SQL statement is called a single-row subquery.

Using Group Functions in a Subquery

```
SQL> SELECT  ename, job, sal
2 FROM      emp
3 WHERE     sal =
4           ( SELECT      MIN(sal)
5             FROM        emp );
```

A curved arrow points from the value '800' to the 'sal =' comparison in the WHERE clause of the main query.

ENAME	JOB	SAL
-----	-----	-----
SMITH	CLERK	800

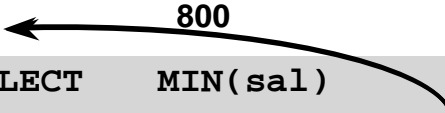
Using Group Functions in a Subquery

You can display data from a main query by using a group function in a subquery to return a single row. The subquery is in parentheses and is placed after the comparison operator.

The example above displays the employee name, job title, and salary of all employees whose salary is equal to the minimum salary. The MIN group function returns a single value (800) to the outer query.

HAVING Clause with Subqueries

- The Oracle8 Server executes subqueries first.
- The Oracle8 Server returns results into the main query's HAVING clause.

```
SQL> SELECT      deptno, MIN(sal)
  2  FROM        emp
  3  GROUP BY    deptno
  4  HAVING      MIN(sal) >  (SELECT      MIN(sal)
  5                                     FROM        emp
  6                                     WHERE      deptno = 20);
  7
```

1-185

Copyright © Oracle Corporation, 1998. All rights reserved.

ORACLE®

HAVING Clause With Subqueries

You can use subqueries not only in the WHERE clause, but also in the HAVING clause. The Oracle8 Server executes the subquery, and the results are returned into the main query's HAVING clause.

The SQL statement on the slide displays all the departments that have a minimum salary greater than that of department 20.

	DEPTNO	MIN (SAL)
	10	1300
Example	30	950

Find the job with the lowest average salary.

```
SQL> SELECT      job, AVG(sal)
  2  FROM        emp
  3  GROUP BY    job
  4  HAVING      AVG(sal) = (SELECT      MIN(AVG(sal))
                             FROM        EMP
                             GROUP BY job);
```

What Is Wrong with This Statement?

```
SQL> SELECT empno, ename
2 FROM emp
3 WHERE sal =
4         (SELECT MIN(sal)
5         FROM emp
6         GROUP BY deptno);
```

```
ERROR:
ORA-01427: single-row subquery returns more than
one row

no rows selected
```

1-186

Copyright © Oracle Corporation, 1998. All rights reserved.

ORACLE®

Errors with Subqueries

One common error with subqueries is more than one row returned for a single-row subquery.

In the SQL statement above, the subquery contains a GROUP BY (deptno) clause, which implies that the subquery will return multiple rows, one for each group it finds. In this case, the result of the subquery will be 800, 1300, and 950.

The outer query takes the results of the subquery (800, 950, 1300) and uses these results in its WHERE clause. The WHERE clause contains an equal (=) operator, a single-row comparison operator expecting only one value. The = operator cannot accept more than one value from the subquery and hence generates the error.

To correct this error, change the = operator to IN.

Will This Statement Work?

```
SQL> SELECT ename, job
2 FROM emp
3 WHERE job =
4         ( SELECT job
5           FROM emp
6           WHERE ename='SMYTHE' );
```

```
no rows selected
```

Subquery returns no values

Errors with Subqueries (continued)

Another common error with subqueries is no rows being returned by the inner query.

In the SQL statement above, the subquery contains a WHERE (ename='SMYTHE') clause. Presumably, the intention is to find the employee whose name is Smythe. The statement seems to be correct but selects no rows when executed.

The problem is that Smythe is misspelled. There is no employee named Smythe. So the subquery returns no rows. The outer query takes the results of the subquery (null) and uses these results in its WHERE clause. The outer query finds no employee with a job title equal to null and so returns no rows.

Class Management Note

Earlier versions of SQL*Plus will generate an error ORA-1246 instead of the output shown above. You may want to show the correct SQL statement by changing the spelling of Smythe to Smith.

Multiple-Row Subqueries

- Return more than one row
- Use multiple-row comparison operators

Operator	Meaning
IN	Equal to any member in the list
ANY	Compare value to each value returned by the subquery
ALL	Compare value to every value returned by the subquery

Multiple-Row subqueries

Subqueries that return more than one row are called *multiple-row subqueries*. You use a multiple-row operator, instead of a single-row operator, with a multiple row subquery. The multiple-row operator expects one or more values.

Example

Find the employees who earn the same salary as the minimum salary for departments.

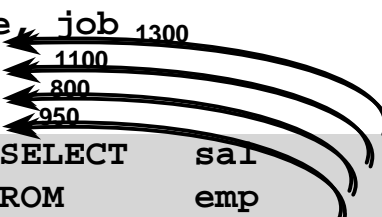
```
SQL> SELECT      ename, sal, deptno
  2 FROM          emp
  3 WHERE          sal IN (SELECT  MIN(sal)
  4 FROM          emp
  5 GROUP BY      deptno);
```

The inner query is executed first, producing a query result containing three rows: 800, 950, 1300. The main query block is then processed and uses the values returned by the inner query to complete its search condition. In fact, the main query would look like the following to the Oracle8 Server:

```
SQL> SELECT      ename, sal, deptno
  2 FROM          emp
  3 WHERE          sal IN (800, 950, 1300);
```

Using ANY Operator in Multiple-Row Subqueries

```
SQL> SELECT empno, ename, job 1300
2 FROM emp 1100
3 WHERE sal < ANY 800
4 (SELECT sal 950
5 FROM emp
6 WHERE job = 'CLERK')
7 AND job <> 'CLERK';
```



EMPNO	ENAME	JOB
7654	MARTIN	SALESMAN
7521	WARD	SALESMAN

Multiple Row Subqueries

The ANY operator (and its synonym SOME operator) compares a value to *each* value returned by a subquery. The example above displays employees whose salary is less than any clerk and who are not clerks. The maximum salary that a clerk earns is \$1300. The SQL statement displays all the employees who are not clerks but earn less than \$1300.

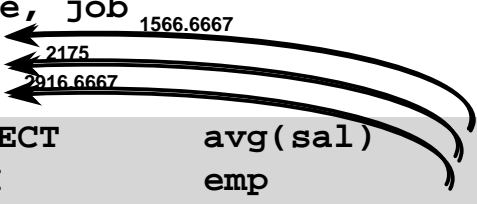
<ANY means less than the maximum. >ANY means more than the minimum. =ANY is equivalent to IN.

Class Management Note

When using SOME or ANY, you often use the DISTINCT keyword to prevent rows from being selected several times.

Using ALL Operator in Multiple-Row Subqueries

```
SQL> SELECT empno, ename, job
2 FROM emp
3 WHERE sal > ALL
4 (SELECT avg(sal)
5 FROM emp
6 GROUP BY deptno)
```



EMPNO	ENAME	JOB
7839	KING	PRESIDENT
7566	JONES	MANAGER
7902	FORD	ANALYST
7788	SCOTT	ANALYST

Multiple Row Subqueries

The ALL operator compares a value to *every* value returned by a subquery. The example above displays employees whose salary is greater than the average salaries of all the departments. The highest average salary of a department is \$2916.66, so the query returns those employees whose salary is greater than \$2916.66.

>ALL means more than the maximum and <ALL means less than the minimum.

The NOT operator can be used with IN, ANY, and ALL operators.

Summary

Subqueries are useful when a query is based on unknown values.

```
SELECT    select_list
FROM      table
WHERE     expr operator
          (SELECT select_list
           FROM table);
```

Summary

A subquery is a SELECT statement that is embedded in a clause of another SQL statement. Subqueries are useful when a query is based on unknown criteria.

Subqueries have the following characteristics:

- Can pass one row of data to a main statement that contains a single-row operator, such as =, <>, >, >=, <, or <=
- Can pass multiple rows of data to a main statement that contains a multiple-row operator, such as IN
- Are processed first by the Oracle8 Server, and the WHERE or HAVING clause uses the results
- Can contain group functions

Practice Overview

- **Creating subqueries to query values based on unknown criteria**
- **Using subqueries to find out what values exist in one set of data and not in another**

1-192

Copyright © Oracle Corporation, 1998. All rights reserved.

ORACLE®

Practice Overview

In this practice, you will write complex queries using nested SELECT statements.

Paper-Based Questions



You may want to consider creating the inner query first for these questions. Make sure that it runs and produces the data you anticipate before coding the outer query.

1

Multiple-Column Subqueries

Copyright © Oracle Corporation, 1998. All rights reserved.

ORACLE®

Schedule:	Timing	Topic
	20 minutes	Lecture
	20 minutes	Practice
	40 minutes	Total

Objectives


At the end of this lesson, you should be able to:

- **Write a multiple-column subquery**
- **Describe and explain the behavior of subqueries when null values are retrieved**
- **Write a subquery in a FROM clause**

Lesson Aim

In this lesson, you will learn how to write multiple-column subqueries and subqueries in the FROM clause of a SELECT statement.

Multiple-Column Subqueries



Main query compares to Values from a multiple-row and multiple-column subquery

MANAGER 10

SALESMAN 30

MANAGER 10

CLERK 20

1-197

Copyright © Oracle Corporation, 1998. All rights reserved.

ORACLE®

Multiple-Column Subqueries

So far you have written single-row subqueries and multiple-row subqueries where only one column was compared in the WHERE clause or HAVING clause of the SELECT statement. If you want to compare two or more columns, you must write a compound WHERE clause using logical operators. Multiple-column subqueries enable you to combine duplicate WHERE conditions into a single WHERE clause.

Syntax

```
SELECT    column, column, ...
FROM      table
WHERE     (column, column, ...) IN
                                (SELECT column, column, ...
                                FROM    table
                                WHERE   condition);
```

Using Multiple-Column Subqueries

Display the name, department number, salary, and commission of any employee whose salary and commission matches both the commission and salary of any employee in department 30.

```
SQL> SELECT  ename, deptno, sal, comm
  2  FROM      emp
  3  WHERE      (sal, NVL(comm,-1)) IN
  4              (SELECT sal, NVL(comm,-1)
  5                  FROM      emp
  6                  WHERE      deptno = 30);
```

1-198

Copyright © Oracle Corporation, 1998. All rights reserved.

ORACLE®

Using Multiple-Column Subqueries

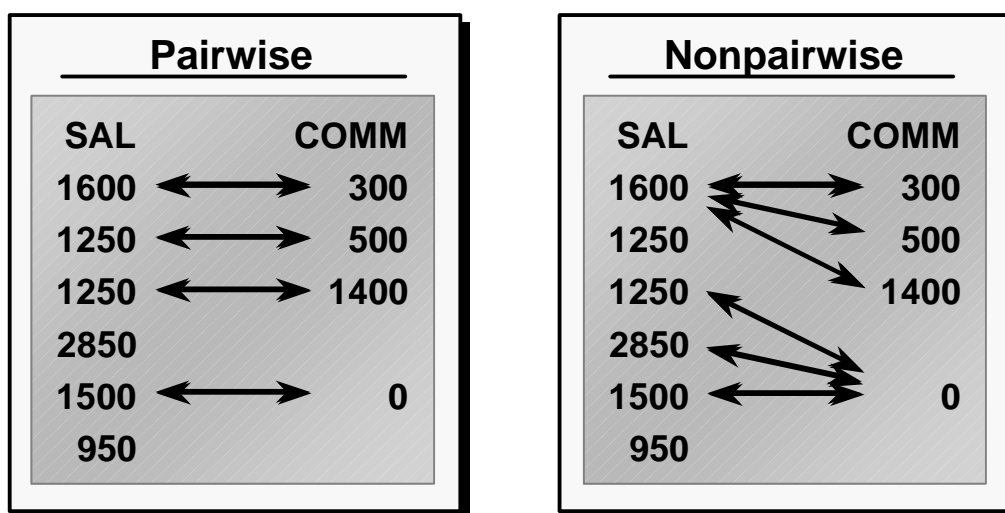
The example above is that of a multiple-column subquery because the subquery returns more than one column. It compares the SAL column and the COMM column. It displays the name, department number, salary, and commission of any employee whose salary and commission matches *both* the commission and salary of any employee in department 30.

The output of the above SQL statement will be as follows:

ENAME	DEPTNO	SAL	COMM
JAMES	30	950	
WARD	30	1250	500
MARTIN	30	1250	1400
TURNER	30	1500	0
ALLEN	30	1600	300
BLAKE	30	2850	

6 rows selected.

Column Comparisons



Pairwise Versus Nonpairwise Comparisons

Column comparisons in a multiple-column subquery can be pairwise comparisons or nonpairwise comparisons. In the example on the previous slide, a pairwise comparison was executed in the WHERE clause. Each candidate row in the SELECT statement must have *both* the same salary and the same commission of an employee in department 30.

If you want a nonpairwise comparison (a cross product), you must use a WHERE clause with multiple conditions.

Nonpairwise Comparison Subquery

Display the name, department number, salary, and commission of any employee whose salary and commission matches the commission and salary of any employee in department 30.

```
SQL> SELECT  ename, deptno, sal, comm
2  FROM      emp
3  WHERE      sal IN          (SELECT sal
4                               FROM    emp
5                               WHERE    deptno = 30)
6  AND
7            NVL(comm,-1) IN  (SELECT NVL(comm,-1)
8                               FROM    emp
9                               WHERE    deptno = 30);
```

1-200

Copyright © Oracle Corporation, 1998. All rights reserved.

ORACLE®

Nonpairwise Comparison Subquery

The above example does a nonpairwise comparison of the columns. It displays the name, department number, salary, and commission of any employee whose salary and commission match the salary and commission of any employee in department 30.

The output of the above SQL statement will be as follows:

ENAME	DEPTNO	SAL	COMM
JAMES	30	950	
BLAKE	30	2850	
TURNER	30	1500	0
ALLEN	30	1600	300
WARD	30	1250	500
MARTIN	30	1250	1400

The results of the last two queries were identical even though the comparison conditions were different. The results were obtained because of the specific data in the EMP table.

6 rows selected.

Modifying the EMP Table

- Assume that salary and commission for Clark are modified.
- Salary is changed to \$1500 and commission to \$300.

ENAME	SAL	COMM
-----	-----	-----
...		
CLARK	1500	300
...		
ALLEN	1600	300
TURNER	1500	0
...		
14 rows selected.		

Example

Assume that the salary and commission of employee Clark is modified so that he has the same salary as an employee in department 30 and the same commission as a different employee in department 30.

The salary for Clark is now equal to that of Turner (\$1500) and the commission for Clark to be equal to that of Allen (\$300).

Now run a pairwise and nonpairwise comparison to determine the number of rows returned by each query.

Note: The syntax for updating data in a table will be discussed in lesson 9.

Pairwise Subquery

```
SQL> SELECT  ename, deptno, sal, comm
2  FROM      emp
3  WHERE      (sal, NVL(comm,-1)) IN
4
5              (SELECT sal, NVL(comm,-1)
6              FROM      emp
              WHERE      deptno = 30);
```

ENAME	DEPTNO	SAL	COMM
JAMES	30	950	
WARD	30	1250	500
MARTIN	30	1250	1400
TURNER	30	1500	0
ALLEN	30	1600	300
BLAKE	30	2850	

6 rows selected.

1-202

Copyright © Oracle Corporation, 1998. All rights reserved.

ORACLE®

Pairwise Subquery

The output of the pairwise subquery still remains the same and returns six rows.

Nonpairwise Subquery

```
SQL> SELECT  ename,deptno, sal, comm
2  FROM      emp
3  WHERE      sal IN          (SELECT sal
4                               FROM    emp
5                               WHERE    deptno = 30)
6  AND
7           NVL(comm,-1) IN  (SELECT NVL(comm,-1)
8                               FROM    emp
9                               WHERE    deptno = 30);
```

ENAME	DEPTNO	SAL	COMM
JAMES	30	950	
BLAKE	30	2850	
TURNER	30	1500	0
CLARK	10	1500	300
...			

7 rows selected.

Nonpairwise Subquery

The results of the nonpairwise subquery include the employee Clark. Clark's salary is the same as that of Turner and his commission is the same as that of Allen.

Null Values in a Subquery

```
SQL> SELECT  employee.ename
      2  FROM    emp employee
      3  WHERE   employee.empno NOT IN
                (SELECT manager.mgr
                 FROM    emp manager);

no rows selected.
```

Returning Nulls in the Resulting Set of a Subquery

The above SQL statement attempts to display all the employees who do not have any subordinates. Logically, this SQL statement should have returned eight rows. However, the SQL statement does not return any rows. One of the values returned by the inner query is a null value and hence the entire query returns no rows. The reason is that all conditions that compare a null value result in a null. So whenever null values are likely to be part of the resultant set of a subquery, do not use the NOT IN operator. The NOT IN operator is equivalent to \neq ALL.

Notice that the null value as part of the resultant set of a subquery will not be a problem if you are using the IN operator. The IN operator is equivalent to $=$ ANY. For example, to display the employees who have subordinates, use the following SQL statement:

```
SQL> SELECT  employee.ename
      2  FROM    emp employee
      3  WHERE   employee.empno IN (SELECT manager.mgr
      4                                     FROM    emp manager);
```

```
ENAME
-----
KING
...
6 rows selected.
```

Using a Subquery in the FROM Clause

```
SQL> SELECT  a.ename, a.sal, a.deptno, b.salavg
  2 FROM      emp a, (SELECT  deptno, avg(sal) salavg
  3              FROM      emp
  4              GROUP BY deptno) b
  5 WHERE      a.deptno = b.deptno
  6 AND        a.sal > b.salavg;
```

ENAME	SAL	DEPTNO	SALAVG
-----	-----	-----	-----
KING	5000	10	2916.6667
JONES	2975	20	2175
SCOTT	3000	20	2175
...			
6 rows selected.			

Using a Subquery in the FROM Clause

You can use a subquery in the FROM clause of a SELECT statement, which is very similar to how views are used. The example above displays employee names, salaries, department numbers, and average salaries for all the employees who make more than the average salary in their department.

Summary

- **A multiple-column subquery returns more than one column.**
- **Column comparisons in a multiple-column comparisons can be pairwise or nonpairwise.**
- **A multiple-column subquery can also be used in the FROM clause of a SELECT statement.**

Practice Overview

Creating multiple-column subqueries

1-207

Copyright © Oracle Corporation, 1998. All rights reserved.

ORACLE®

Practice Overview

In this practice, you will write multiple-value subqueries.

1

Manipulating Data

Copyright © Oracle Corporation, 1998. All rights reserved.

ORACLE®

Schedule:	Timing	Topic
	40 minutes	Lecture
	30 minutes	Practice
	70 minutes	Total

Objectives

At the end of this lesson, you should be able to:

- **Describe each DML statement**
- **Insert rows into a table**
- **Update rows in a table**
- **Delete rows from a table**
- **Control transactions**

Lesson Aim

In this lesson, you will learn how to insert rows into a table, update existing rows in a table, and delete existing rows from a table. You will also learn how to control transactions with the COMMIT, SAVEPOINT, and ROLLBACK statements.

Data Manipulation Language

- A DML statement is executed when you:
 - Add new rows to a table
 - Modify existing rows in a table
 - Remove existing rows from a table
- A *transaction* consists of a collection of DML statements that form a logical unit of work.

Data Manipulation Language

Data manipulation language (DML) is a core part of SQL. When you want to add, update, or delete data in the database, you execute a DML statement. A collection of DML statements that form a logical unit of work is called a *transaction*.

Class Management Note

DML statements can be issued directly in SQL*Plus or SQL*DBA, performed automatically by tools such as Developer/2000 or programmed with tools such as the 3GL precompilers.

Every table has INSERT, UPDATE, and DELETE privileges associated with it. These privileges are automatically granted to the creator of the table, but in general they must be explicitly granted to other users.

Starting with Oracle 7.2, you can place a subquery in the place of the table name in an UPDATE statement, essentially the same way a view is used. For example:

```
UPDATE    (SELECT * FROM dept)
SET       DEPTNO = 90
WHERE     DEPTNO = 80;
```


Adding a New Row to a Table

50	DEVELOPMENT	DETROIT
----	-------------	---------

New row

DEPT

DEPTNO	DNAME	LOC
10	ACCOUNTING	NEW YORK
20	RESEARCH	DALLAS
30	SALES	CHICAGO
40	OPERATIONS	BOSTON

“...insert a new row
into DEPT table...”



DEPT

DEPTNO	DNAME	LOC
10	ACCOUNTING	NEW YORK
20	RESEARCH	DALLAS
30	SALES	CHICAGO
40	OPERATIONS	BOSTON
50	DEVELOPMENT	DETROIT

Adding a New Row to a Table

The graphic above adds a new department to the DEPT table.

The INSERT Statement

- Add new rows to a table by using the INSERT statement.

```
INSERT INTO          table [(column [, column...])]
VALUES               (value [, value...]);
```

- Only one row is inserted at a time with this syntax.

Adding a New Row to a Table

You can add new rows to a table by issuing the INSERT statement.

In the syntax:

<i>table</i>	is the name of the table
<i>column</i>	is the name of the column in the table to populate
<i>value</i>	is the corresponding value for the column

Note: This statement with the VALUES clause adds only one row at a time to a table.

Inserting New Rows

- Insert a new row containing values for each column.
- Optionally list the columns in the **INSERT** clause.

```
SQL> INSERT INTO dept (deptno, dname, loc)
2 VALUES (50, 'DEVELOPMENT', 'DETROIT');
1 row created.
```

- List values in the default order of the columns in the table.
- Enclose character and date values within single quotation marks.

1-214

Copyright © Oracle Corporation, 1998. All rights reserved.

ORACLE®

Adding a New Row to a Table

Because you can insert a new row that contains values for each column, the column list is not required in the INSERT clause. However, the values must be listed according to the default order of the columns in the table.

```
SQL> DESCRIBE dept
```

Name	Null?	Type
DEPTNO	NOT NULL	NUMBER(2)
DNAME		VARCHAR2(14)
LOC		VARCHAR2(13)

For clarity, use the column list in the INSERT clause.

Enclose character and date values within single quotation marks; do not enclose numeric values within single quotation marks.



Inserting Rows with Null Values

- **Implicit method: Omit the column from the column list.**

```
SQL> INSERT INTO      dept (deptno, dname )
      2  VALUES      (60, 'MIS');
1 row created.
```

- **Explicit method: Specify the NULL keyword.**

```
SQL> INSERT INTO      dept
      2  VALUES      (70, 'FINANCE', NULL);
1 row created.
```

1-215

Copyright © Oracle Corporation, 1998. All rights reserved.

ORACLE®

Methods for Inserting Null Values

Method	Description
Implicit	Omit the column from the column list
Explicit	Specify the NULL keyword in the VALUES list
Be sure that the targeted column allows null values by verifying the Null? status from the SQL*Plus DESCRIBE command.	
Specify the empty string (' ') in the VALUES list; for character strings and dates only	

The Oracle8 Server automatically enforces all datatypes, data ranges, and data integrity constraints. Any column that is not listed explicitly obtains a null value in the new row.

Class Management Note

Common errors that can occur during user input:

- Mandatory value missing for a NOT NULL column
- Duplicate value violates uniqueness constraint
- Foreign key constraint violated
- CHECK constraint violated
- Datatype mismatch
- Value too wide to fit in column

Inserting Special Values

The **SYSDATE** function records the current date and time.

```
SQL> INSERT INTO      emp (empno, ename, job,
 2                      mgr, hiredate, sal, comm,
 3                      deptno)
 4  VALUES             (7196, 'GREEN', 'SALESMAN',
 5                      7782, SYSDATE, 2000, NULL,
 6                      10);
1 row created.
```

1-216

Copyright © Oracle Corporation, 1998. All rights reserved.

ORACLE®

Inserting Special Values by Using SQL Functions

You can use pseudocolumns to enter special values in your table.

The example above records information for employee Green in the EMP table. It supplies the current date and time in the HIREDATE column. It uses the *SYSDATE* function for current date and time.

You can also use the *USER* function when inserting rows in a table. The *USER* function records the current username.

Confirming Additions to the Table

```
SQL> SELECT empno, ename, job, hiredate, comm
 2  FROM      emp
 3  WHERE empno = 7196;
```

EMPNO	ENAME	JOB	HIREDATE	COMM
7196	GREEN	SALESMAN	01-DEC-97	

Inserting Specific Date Values

- Add a new employee.

```
SQL> INSERT INTO emp
  2  VALUES      (2296,'AROMANO','SALESMAN',7782,
  3                TO_DATE('FEB 3,97','MON DD, YY'),
  4                1300, NULL, 10);
1 row created.
```

- Verify your addition.

EMPNO	ENAME	JOB	MGR	HIREDATE	SAL	COMM	DEPTNO
2296	AROMANO	SALESMAN	7782	03-FEB-97	1300		10

1-217

Copyright © Oracle Corporation, 1998. All rights reserved.

ORACLE®

Inserting Specific Date and Time Values

The format DD-MON-YY is usually used to insert a date value. With this format, recall that the century defaults to the current century. Because the date also contains time information, the default time is midnight (00:00:00).

If a date is required to be entered in another century and a specific time is also required, use the TO_DATE function.

The example on the slide records information for employee Aromano in the EMP table. It sets the HIREDATE column to be February 3, 1997.

If the RR format is set, the century may not be the current one.



Inserting Values by Using Substitution Variables

Create an interactive script by using SQL*Plus substitution parameters.

```
SQL> INSERT INTO      dept (deptno, dname, loc)
  2  VALUES          (&department_id,
  3                  '&department_name', '&location');

```

```
Enter value for department_id: 80
Enter value for department_name: EDUCATION
Enter value for location: ATLANTA

1 row created.

```

1-218

Copyright © Oracle Corporation, 1998. All rights reserved.

ORACLE®

Inserting Values by Using Substitution Variables

You can produce an INSERT statement that allows the user to add values interactively by using SQL*Plus substitution variables.

The example above records information for a department in the DEPT table. It prompts the user for the department number, department name, and location.

For date and character values, the ampersand and the variable name are enclosed in single quotation marks.



Class Management Note

Be sure to mention the following points about the example:

- The names of the SQL*Plus substitution parameters do not have to match the corresponding column names.
- Substitution parameters are lexical variables. Whatever characters the user enters are substituted as text for the variable name.
- The SQL*Plus SET VERIFY command lists the substitution before executing the statement.

Creating a Script with Customized Prompts

- **ACCEPT** stores the value into a variable.
- **PROMPT** displays your customized text.

```
ACCEPT      department_id PROMPT 'Please enter the -  
            department number: '  
  
ACCEPT      department_name PROMPT 'Please enter -  
            the department name: '  
  
ACCEPT      location PROMPT 'Please enter the -  
            location: '  
  
INSERT INTO dept (deptno, dname, loc)  
VALUES      (&department_id, '&department_name',  
            &location);
```

1-219

Copyright © Oracle Corporation, 1998. All rights reserved.

ORACLE®

Creating a Script to Manipulate Data

You can save your command with substitution variables to a file and execute the file. Each time you execute the command, it will prompt you for new values. Customize the prompts by using the SQL*Plus ACCEPT command.

The example on the slide records information for a department in the DEPT table. It prompts the user for the department number, department name, and location by using customized prompt messages.

```
Please enter the department number: 90  
Please enter the department name: PAYROLL  
Please enter the location: HOUSTON
```

```
1 row created.
```

Do not prefix the SQL*Plus substitution parameter with the ampersand (&) when referencing it in the ACCEPT command. Use a dash (-) to continue a SQL*Plus command on the next line.

Class Management Note

Be sure to mention the following points about the script:

- Do not prefix the SQL*Plus substitution parameter with the ampersand in the ACCEPT command.
- Use a dash to continue a SQL*Plus command on the next line.
- Add a space after the colon in the PROMPT command.



Copying Rows from Another Table

- Write your INSERT statement with a subquery.

```
SQL> INSERT INTO managers(id, name, salary, hiredate)
      2          SELECT empno, ename, sal, hiredate
      3          FROM    emp
      4          WHERE   job = 'MANAGER';
3 rows created.
```

- Do not use the VALUES clause.
- Match the number of columns in the INSERT clause to those in the subquery.

1-220

Copyright © Oracle Corporation, 1998. All rights reserved.

ORACLE®

Copying Rows from Another Table

You can use the INSERT statement to add rows to a table where the values are derived from existing tables. In place of the VALUES clause, you use a subquery.

Syntax

```
INSERT INTO table [ column (, column) ]
              subquery;
```

where: *table* is the table name
column is the name of the column in the table to populate
subquery is the subquery that returns rows into the table

For more information, see
Oracle8 Server SQL Reference, Release 8.0, “SELECT,” Subqueries section.



The number of columns and their datatypes in the column list of the INSERT clause must match the number of values and their datatypes in the subquery.



Class Management Note


Do not get into too many details on copying rows from another table.

Changing Data in a Table

EMP

EMPNO	ENAME	JOB	...	DEPTNO
7839	KING	PRESIDENT		10
7698	BLAKE	MANAGER		30
7782	CLARK	MANAGER		10
7566	JONES	MANAGER		20
...				

“...update a row
in EMP table...”



EMP

EMPNO	ENAME	JOB	...	DEPTNO
7839	KING	PRESIDENT		10
7698	BLAKE	MANAGER		30
7782	CLARK	MANAGER		20
7566	JONES	MANAGER		20
...				

1-221

Copyright © Oracle Corporation, 1998. All rights reserved.

ORACLE®

Changing Data in a Table

The graphic above changes the department number for Clark from 10 to 20.

The UPDATE Statement

- **Modify existing rows with the UPDATE statement.**

```
UPDATE      table
SET         column = value [, column = value]
[WHERE      condition];
```

- **Update more than one row at a time, if required.**

1-222

Copyright © Oracle Corporation, 1998. All rights reserved.

ORACLE®

Updating Rows

You can modify existing rows by using the UPDATE statement.

In the above syntax:

<i>table</i>	is the name of the table
<i>column</i>	is the name of the column in the table to populate
<i>value</i>	is the corresponding value or subquery for the column
<i>condition</i>	identifies the rows to be updated and is composed of column names expressions, constants, subqueries, and comparison operators

Confirm the update operation by querying the table to display the updated rows.

For more information, see

Oracle8 Server SQL Reference, Release 8.0, "UPDATE."



Note: In general, use the primary key to identify a single row. Using other columns may unexpectedly cause several rows to be updated. For example, identifying a single row in the EMP table by name is dangerous because more than one employee may have the same name.

Class Management Note

Demo: *l9sel.sql*, *l9upd.sql*

Updating Rows in a Table

- Specific row or rows are modified when you specify the WHERE clause.

```
SQL> UPDATE    emp
   2  SET      deptno = 20
   3  WHERE    empno = 7782;
1 row updated.
```

- All rows in the table are modified if you omit the WHERE clause.

```
SQL> UPDATE    employee
   2  SET      deptno = 20;
14 rows updated.
```

1-223

Copyright © Oracle Corporation, 1998. All rights reserved.

ORACLE®

Updating Rows

The UPDATE statement modifies specific row(s), if the WHERE clause is specified. The example above transfers employee 7782 (Clark) to department 20.

If you omit the WHERE clause, all the rows in the table are modified.

```
SQL> SELECT    ename, deptno
   2  FROM      employee;
```

ENAME	DEPTNO
-----	-----
KING	20
BLAKE	20
CLARK	20
JONES	20
MARTIN	20
ALLEN	20
TURNER	20

Note: The EMPLOYEE table has the same data as the EMP table.

...

14 rows selected.

Updating with Multiple-Column Subquery

Update employee 7698's job and department to match that of employee 7499.

```
SQL> UPDATE   emp
   2  SET      (job, deptno) =
   3              (SELECT job, deptno
   4                  FROM    emp
   5                  WHERE   empno = 7499)
   6  WHERE    empno = 7698;
1 row updated.
```

1-224

Copyright © Oracle Corporation, 1998. All rights reserved.

ORACLE®

Updating Rows with a Multiple-Column Subquery

Multiple-column subqueries can be implemented in the SET clause of an UPDATE statement.

Syntax

```
UPDATE   table
SET      (column, column, ...) =
              (SELECT   column, column,
                  FROM    table
                  WHERE   condition)
WHERE condition;
```

Class Management Note

It may be worth showing participants that the results would be the same for the example above if two different subqueries were used in the SET clause as illustrated below:

```
UPDATE      emp
SET         job = (SELECT  job FROM emp
                  WHERE    empno = 7499),
            deptno = (SELECT deptno FROM dept
                     WHERE  dname = 'SALES')
WHERE      empno = 7698;
```

Updating Rows Based on Another Table

Use subqueries in UPDATE statements to update rows in a table based on values from another table.

```
SQL> UPDATE    employee
  2  SET        deptno = (SELECT    deptno
  3                                     FROM      emp
  4                                     WHERE     empno = 7788)
  5  WHERE      job    = (SELECT    job
  6                                     FROM      emp
  7                                     WHERE     empno = 7788);
2 rows updated.
```

Updating Rows Based on Another Table

You can use subqueries in UPDATE statements to update rows in a table. The example above updates the EMPLOYEE table based on the values from the EMP table. It changes the department number of all employees with employee 7788's job title to employee 7788's current department number.

Updating Rows: Integrity Constraint Error

```
SQL> UPDATE emp
2 SET deptno = 55
3 WHERE deptno = 10;
```

```
UPDATE emp
*
ERROR at line 1:
ORA-02291: integrity constraint (USR.EMP_DEPTNO_FK)
violated - parent key not found
```

Integrity Constraint Error

If you attempt to update a record with a value that is tied to an integrity constraint, you will experience an error.

In the example above, department number 55 does not exist in the parent table, DEPT, and so you receive the “parent key” violation ORA-02291.

Note: Integrity constraints assure that the data adheres to a predefined set of rules. Lesson 11 will cover integrity constraints in greater depth.


Removing a Row from a Table

DEPT

DEPTNO	DNAME	LOC
10	ACCOUNTING	NEW YORK
20	RESEARCH	DALLAS
30	SALES	CHICAGO
40	OPERATIONS	BOSTON
50	DEVELOPMENT	DETROIT
60	MIS	
...		

“...delete a row
from DEPT table...”

DEPT



DEPTNO	DNAME	LOC
10	ACCOUNTING	NEW YORK
20	RESEARCH	DALLAS
30	SALES	CHICAGO
40	OPERATIONS	BOSTON
60	MIS	
...		

1-227

Copyright © Oracle Corporation, 1998. All rights reserved.

ORACLE®

Removing a Row from a Table

The graphic above removes the DEVELOPMENT department from the DEPT table (assuming there are no constraints defined on the DEPT table).

Class Management Note

After all the rows have been eliminated with the DELETE statement, only the data structure of the table remains. A more efficient method of emptying a table is with the TRUNCATE statement.

You can use the TRUNCATE statement to quickly remove all rows from a table or cluster. Removing rows with the TRUNCATE statement is faster than removing them with the DELETE statement for the following reasons:

- The TRUNCATE statement is a data definition language (DDL) statement and generates no rollback information. It will be covered in lesson 10.
- Truncating a table does not fire the table's DELETE triggers.
- If the table is the parent of a referential integrity constraint, you cannot truncate the table. Disable the constraint before issuing the TRUNCATE statement.

The DELETE Statement

You can remove existing rows from a table by using the DELETE statement.

```
DELETE [FROM]   table
[WHERE          condition];
```

1-228

Copyright © Oracle Corporation, 1998. All rights reserved.

ORACLE®

Deleting Rows

You can remove existing rows by using the DELETE statement.

In the syntax:

table is the table name
condition identifies the rows to be deleted and is composed of column names, expressions, constants, subqueries, and comparison operators

For more information, see

Oracle8 Server SQL Reference, Release 8.0, "DELETE."



Class Management Note

The DELETE statement does not ask for confirmation. However, the delete operation is not made permanent until the data transaction is committed. Therefore, you can undo the operation with the ROLLBACK statement if you make a mistake.

Deleting Rows from a Table

- Specific row or rows are deleted when you specify the WHERE clause.

```
SQL> DELETE FROM      department
      2 WHERE          dname = 'DEVELOPMENT';
      1 row deleted.
```

- All rows in the table are deleted if you omit the WHERE clause.

```
SQL> DELETE FROM      department;
      4 rows deleted.
```

1-229

Copyright © Oracle Corporation, 1998. All rights reserved.

ORACLE®

Deleting Rows

You can delete specific row(s) by specifying the WHERE clause in the DELETE statement. The example above deletes the DEVELOPMENT department from the DEPARTMENT table. You can confirm the delete operation by displaying the deleted rows by using the SELECT statement.

```
SQL> SELECT  *
      2 FROM    department
      3 WHERE   dname = 'DEVELOPMENT';
no rows selected.
```

Example

Remove all employees who started after January 1, 1997.

```
SQL> DELETE FROM emp
      2 WHERE hiredate > TO_DATE('01-01-97', 'DD.MM.YY');
      1 row deleted.
Note: The DEPARTMENT table has the same data as the DEPT table.
```

Deleting Rows Based on Another Table

Use subqueries in DELETE statements to remove rows from a table based on values from another table.

```
SQL> DELETE FROM      employee
  2  WHERE              deptno =
  3                      ( SELECT    deptno
  4                          FROM      dept
  5                          WHERE      dname = 'SALES' );
  6 rows deleted.
```

Deleting Rows Based on Another Table

You can use subqueries to delete rows from a table based on values from another table. The example above deletes all the employees who are in department 30. The subquery searches the DEPT table to find the department number for the SALES department. The subquery then feeds the department number to the main query, which deletes rows of data from the EMPLOYEE table based on this department number.

Deleting Rows: Integrity Constraint Error

```
SQL> DELETE FROM dept
      2 WHERE deptno = 10;
```

```
DELETE FROM dept
      *
ERROR at line 1:
ORA-02292: integrity constraint (USR.EMP_DEPTNO_FK)
violated - child record found
```

1-231

Copyright © Oracle Corporation, 1998. All rights reserved.

ORACLE®

Integrity Constraint Error

If you attempt to delete a record with a value that is tied to an integrity constraint, you will experience an error.

The example above tries to delete department number 10 from the DEPT table, but it results in an error because department number is used as a foreign key in the EMP table. If the parent record you attempt to delete has child records, then you receive the “child record found” violation ORA-02292.

Class Management Note

If there are referential integrity constraints in use, you might receive an Oracle Server error message when you attempt to delete a row. However, if the referential integrity constraint contains the ON DELETE CASCADE option, then the selected row and its children are deleted from their respective tables.

Database Transactions

Contain one of the following statements:

- **DML statements that make up one consistent change to the data**
- **One DDL statement**
- **One DCL statement**

Database Transactions

The Oracle8 Server ensures data consistency based on transactions. Transactions give you more flexibility and control when changing data, and they assure data consistency in the event of user process failure or system failure.

Transactions consist of DML statements that make up one consistent change to the data. For example, a transfer of funds between two accounts should include the debit to one account and the credit to another account in the same amount. Both actions should either fail or succeed together. The credit should not be committed without the debit.

Transaction Types

Type	Description
Data manipulation language (DML)	Consists of any number of DML statements that the Oracle8 Server treats as a single entity or a logical unit of work
Data definition language (DDL)	Consists of only one DDL statement
Data control language (DCL)	Consists of only one DCL statement

Database Transactions

- **Begin when the first executable SQL statement is executed**
- **End with one of the following events:**
 - **COMMIT or ROLLBACK**
 - **DDL or DCL statement executes (automatic commit)**
 - **Certain errors, exit, or system crash**

1-233

Copyright © Oracle Corporation, 1998. All rights reserved.

ORACLE®

When Does a Transaction Start and End?

A transaction begins when the first executable SQL statement is encountered and terminates when one of the following occurs:

- A COMMIT or ROLLBACK statement is issued
- A DDL statement, such as CREATE, is issued
- A DCL statement is issued
- Certain errors are detected, such as deadlocks
- The user exits SQL*Plus
- A machine fails or the system crashes

After one transaction ends, the next executable SQL statement will automatically start the next transaction.

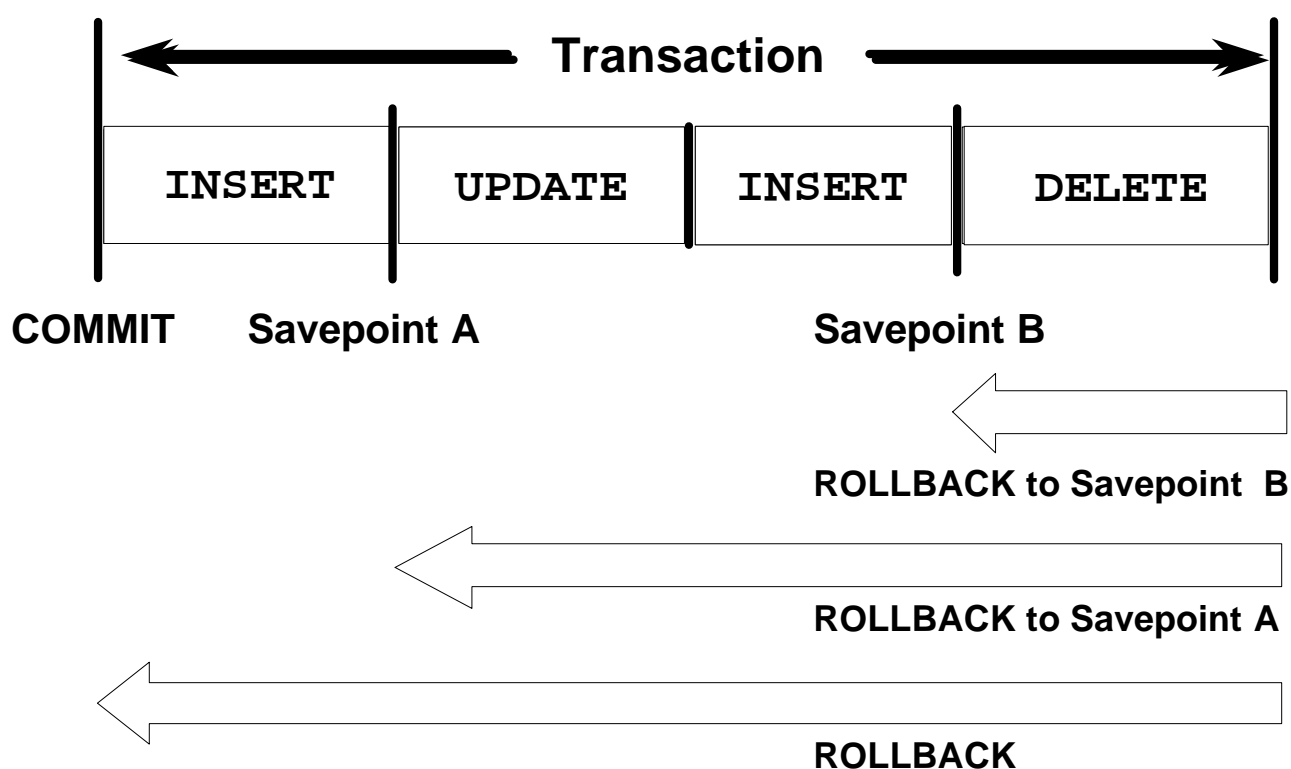
A DDL statement or a DCL statement is automatically committed and therefore implicitly ends a transaction.



Advantages of COMMIT and ROLLBACK

- **Ensure data consistency**
- **Preview data changes before making changes permanent**
- **Group logically related operations**

Controlling Transactions



1-235

Copyright © Oracle Corporation, 1998. All rights reserved.

ORACLE®

Explicit Transaction Control Statements

You can control the logic of transactions by using the **COMMIT**, **SAVEPOINT**, and **ROLLBACK** statements.

Statement	Description
COMMIT	Ends the current transaction by making all pending data changes permanent
SAVEPOINT <i>name</i>	Marks a savepoint within the current transaction
ROLLBACK [TO <i>SAVEPOINT name</i>] <small>Note: SAVEPOINT is not ANSI standard SQL.</small>	A ROLLBACK ends the current transaction by discarding all pending data changes. ROLLBACK TO SAVEPOINT <i>name</i> discards the savepoint and all subsequent changes

Implicit Transaction Processing

- **An automatic commit occurs under the following circumstances:**
 - A DDL statement is issued
 - A DCL statement is issued
 - A normal exit from SQL*Plus, without explicitly issuing COMMIT or ROLLBACK
- **An automatic rollback occurs under an abnormal termination of SQL*Plus or a system failure**

1-236

Copyright © Oracle Corporation, 1998. All rights reserved.

ORACLE®

Implicit Transaction Processing

Status	Circumstances
Automatic commit	DDL statement or DCL statement is issued Normal exit from SQL*Plus, without explicitly issuing COMMIT or ROLLBACK
Automatic rollback	Abnormal termination of SQL*Plus or system failure

Note: A third command is available in SQL*Plus. The SQL*Plus AUTOCOMMIT command can be toggled to be ON or OFF. If set to ON, each individual DML statement is committed as soon as it is executed. You cannot roll back the changes. If set to OFF, COMMIT can be issued explicitly. Also, COMMIT is issued when a DDL statement is issued or when you exit from SQL*Plus.

System Failures

When a transaction is interrupted by a system failure, the entire transaction is automatically rolled back. This prevents the error from causing unwanted changes to the data and returns the tables to their state at the time of the last commit. In this way, SQL*Plus protects the integrity of the tables.

State of the Data Before COMMIT or ROLLBACK

- The previous state of the data can be recovered because the database buffer is affected.
- The current user can review the results of the DML operations by using the SELECT statement.
- Other users *cannot* view the results of the DML statements by the current user.
- The affected rows are *locked*; other users cannot change the data within the affected rows.

1-237

Copyright © Oracle Corporation, 1998. All rights reserved.

ORACLE®

Committing Changes

Every data change made during the transaction is temporary until the transaction is committed.

State of the Data Before COMMIT or ROLLBACK

- Data manipulation operations primarily affect the database buffer; therefore, the previous state of the data can be recovered.
- The current user can review the results of the data manipulation operations by querying the tables.
- Other users cannot view the results of the data manipulation operations made by the current user. Oracle8 institutes read consistency to ensure that each user sees data as it existed at the last commit.
- The affected rows are locked; other users cannot change the data within the affected rows

Class Management Note

With the Oracle8 Server, data changes may actually be written to the database files before COMMIT, but they are still only temporary.

If a number of users are making changes simultaneously to the same table, then each user sees only his or her changes until other users commit their changes.

Other users see data as it is committed in the database (in other words, before changes).

By default, the Oracle8 Server has *row-level locking*. It is possible to alter the default locking mechanism.

State of the Data After COMMIT

- **Data changes are made permanent in the database.**
- **The previous state of the data is permanently lost.**
- **All users can view the results.**
- **Locks on the affected rows are released; those rows are available for other users to manipulate.**
- **All savepoints are erased.**

1-238

Copyright © Oracle Corporation, 1998. All rights reserved.

ORACLE®

Committing Changes

Make all pending changes permanent by using the COMMIT statement. Following a COMMIT:

- Data changes are written to the database.
- The previous state of the data is permanently lost.
- All users can view the results of the transaction.
- The locks on the affected rows are released; the rows are now available for other users to perform new data changes.
- All savepoints are erased.

Committing Data

- **Make the changes.**

```
SQL> UPDATE emp
2 SET deptno = 10
3 WHERE empno = 7782;
1 row updated.
```

- **Commit the changes.**

```
SQL> COMMIT;
Commit complete.
```

Committing Changes

The above example updates the EMP table and sets the department number for employee 7782 (Clark) to 10. It then makes the change permanent by issuing the COMMIT statement.

Example

Create a new ADVERTISING department with at least one employee. Make the data change permanent.

```
SQL> INSERT INTO department(deptno, dname, loc)
2 VALUES (50, 'ADVERTISING', 'MIAMI');
1 row created.
```

```
SQL> UPDATE employee
2 SET deptno = 50
3 WHERE empno = 7876;
1 row updated.
```

Class Management Note

```
SQL> COMMIT;
Commit complete.
```

Use this example to explain how COMMIT ensures that two related operations must occur together, or not at all. In this case, COMMIT prevents empty departments from being created.

State of the Data After ROLLBACK

Discard all pending changes by using the ROLLBACK statement.

- Data changes are undone.
- Previous state of the data is restored.
- Locks on the affected rows are released.

```
SQL> DELETE FROM      employee;  
14 rows deleted.  
SQL> ROLLBACK;  
Rollback complete.
```

Rolling Back Changes

Discard all pending changes by using the ROLLBACK statement. Following a ROLLBACK

- Data changes are undone.
- The previous state of the data is restored.
- The locks on the affected rows are released.

Example

While attempting to remove a record from the TEST table, you can accidentally empty the table. You can correct the mistake, then reissue the proper statement, and make the data change permanent.

```
SQL> DELETE FROM  test;  
25,000 rows deleted.  
SQL> ROLLBACK;  
Rollback complete.  
SQL> DELETE FROM  test  
2  WHERE          id = 100;  
1 row deleted.  
SQL> SELECT      *  
2  FROM          test  
3  WHERE          id = 100;  
No rows selected.  
SQL> COMMIT;  
Commit complete.
```

Rolling Back Changes to a Marker

- **Create a marker within a current transaction by using the `SAVEPOINT` statement.**
- **Roll back to that marker by using the `ROLLBACK TO SAVEPOINT` statement.**

```
SQL> UPDATE...  
SQL> SAVEPOINT update_done;  
Savepoint created.  
SQL> INSERT...  
SQL> ROLLBACK TO update_done;  
Rollback complete.
```

Rolling Back Changes to a Savepoint

You can create a marker within the current transaction by using the `SAVEPOINT` statement. The transaction therefore can be divided into smaller sections. You can then discard pending changes up to that marker by using the `ROLLBACK TO SAVEPOINT` statement.

If you create a second savepoint with the same name as an earlier savepoint, the earlier savepoint is deleted.

Class Management Note

Savepoints are especially useful in PL/SQL or a 3GL program in which recent changes can be undone conditionally based on runtime conditions.

Statement-Level Rollback

- **If a single DML statement fails during execution, only that statement is rolled back.**
- **Oracle8 implements an implicit savepoint.**
- **All other changes are retained.**
- **The user should terminate transactions explicitly by executing a COMMIT or ROLLBACK statement.**

1-242

Copyright © Oracle Corporation, 1998. All rights reserved.

ORACLE®

Statement-Level Rollback

Part of a transaction can be discarded by an implicit rollback if a statement execution error is detected. If a single DML statement fails during execution of a transaction, its effect is undone by a statement-level rollback, but the changes made by the previous DML statements in the transaction will not be discarded. They can be committed or rolled back explicitly by the user.

Oracle issues an implicit COMMIT before and after any data definition language (DDL) statement. So, even if your DDL statement does not execute successfully, you cannot roll back the previous statement because the server issued a commit.

Terminate your transactions explicitly by executing a COMMIT or ROLLBACK statement.



Class Management Note

The Oracle8 Server implements locks on data to provide data concurrency in the database. Those locks are released when certain events occur (such as a system failure) or the transaction is completed.

Implicit locks on the database are obtained when a DML statement is successfully executed. The Oracle8 Server by default locks data at the lowest level possible.

Manually acquire locks on the database tables by executing a LOCK TABLE statement or the SELECT statement with the FOR UPDATE clause.

Read Consistency

- **Read consistency guarantees a consistent view of the data at all times.**
- **Changes made by one user do not conflict with changes made by another user.**
- **Ensures that on the same data:**
 - **Readers do not wait for writers**
 - **Writers do not wait for readers**

Read Consistency

Database users make two types of access to the database

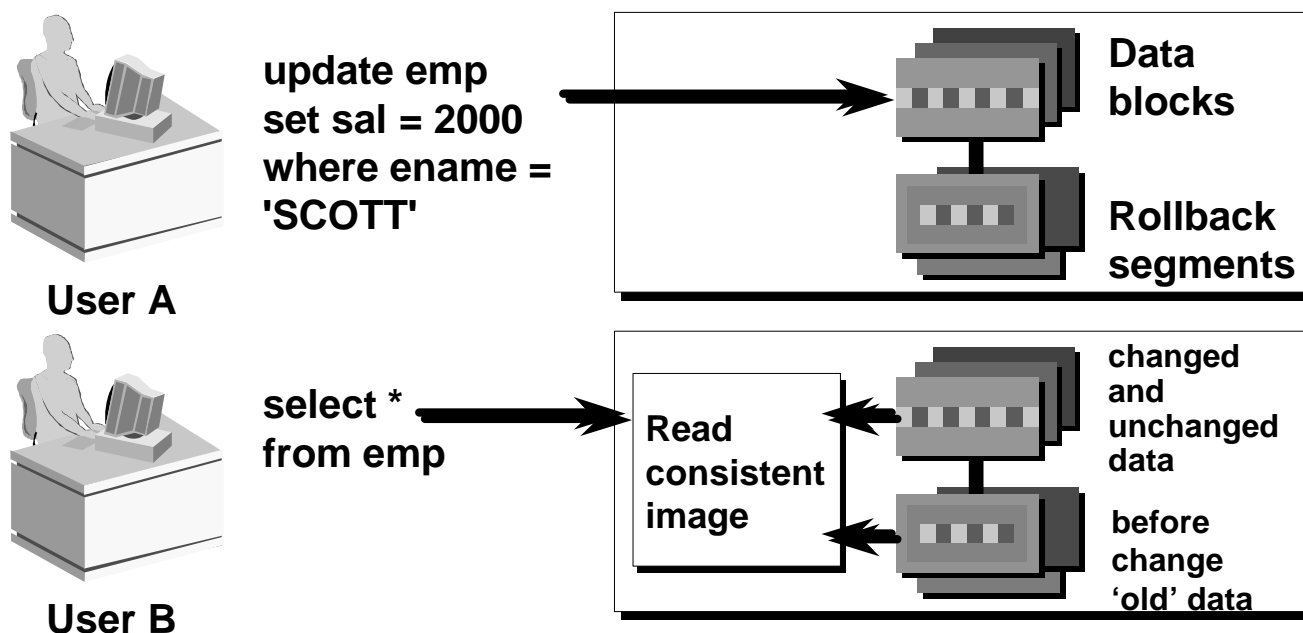
- Read operations (SELECT statement)
- Write operations (INSERT, UPDATE, DELETE statements)

You need read consistency so that the following occur:

- The database reader and writer are ensured a consistent view of the data
- Readers do not view data that is in the process of being changed
- Writers are ensured that the changes to the database are done in a consistent way
- Changes made by one writer do not disrupt or conflict with changes another writer is making

The purpose of read consistency is to ensure that each user sees data as it existed at the last commit, before a DML operation started.

Implementation of Read Consistency



1-244

Copyright © Oracle Corporation, 1998. All rights reserved.

ORACLE®

Implementation of Read Consistency

Read consistency is an automatic implementation. It keeps a partial copy of the database in rollback segments.

When an insert, update, or delete operation is made to the database, Oracle8 Server takes a copy of the data before it is changed and writes it to a *rollback segment*.

All readers, except the one who issued the change, still see the database as it existed before the changes started; they view the rollback segments' "snapshot" of the data.

Before changes are committed to the database, only the user who is modifying the data sees the database with the alterations, everyone else sees the snapshot in the rollback segment. This guarantees that readers of the data read consistent data that is not currently undergoing change.

When a DML statement is committed, the change made to the database becomes visible to anyone executing a SELECT statement.

The space occupied by the "old" data in the rollback segment file is freed for reuse.

If the transaction is rolled back, the changes are "undone."

- The original, older version, of the data in the rollback segment is written back to the table.
- All users see the database as it existed before the transaction began.

Class Management Note

Do not get into the details of data blocks and rollback segments. Explain them in basic terms to describe read consistency.

Locking

Oracle8 locks:

- **Prevent destructive interaction between concurrent transactions**
- **Require no user action**
- **Automatically use the lowest level of restrictiveness**
- **Are held for the duration of the transaction**
- **Have two basic modes:**
 - **Exclusive**
 - **Shared**

1-245

Copyright © Oracle Corporation, 1998. All rights reserved.

ORACLE®

What are Locks?

Locks are mechanisms that prevent destructive interaction between transactions accessing the same resource, either a user object (such as tables or rows) or system objects not visible to users (such as shared data structures and data dictionary rows).

How Oracle Locks Data

Locking in an Oracle database is fully automatic and requires no user action. Implicit locking occurs for all SQL statements. The Oracle default locking mechanism automatically uses the lowest applicable level of restrictiveness, thus providing the highest degree of concurrency yet also providing maximum data integrity. Oracle also allows the user to lock data manually.

Locking Modes

Oracle uses two modes of locking in a multiuser database.

Lock Mode	Description
<i>exclusive</i>	Prevents a resource from being shared The first transaction to lock a resource exclusively, is the only transaction that can alter the resource until the exclusive lock is released
<i>share lock</i>	Allows the resource to be shared Multiple users reading data can share the data, holding share locks to prevent concurrent access by a writer (who needs an exclusive lock) Several transactions can acquire share locks on the same resource

Summary

Statement	Description
INSERT	Adds a new row to the table
UPDATE	Modifies existing rows in the table
DELETE	Removes existing rows from the table
COMMIT	Makes all pending changes permanent
SAVEPOINT	Allows a rollback to the savepoint marker
ROLLBACK	Discards all pending data changes

Summary

Manipulate data in the Oracle database by using the INSERT, UPDATE, and DELETE statements. Control data changes by using the COMMIT, SAVEPOINT, and ROLLBACK statements.

Oracle8 guarantees a consistent view of data at all times.

Locking can be implicit or explicit.

Practice Overview

- **Inserting rows into the tables.**
- **Updating and deleting rows in the table.**
- **Controlling transactions.**

Practice Overview

In this practice, you will add rows to the MY_EMPLOYEE table, update, and delete data from the table, and control your transactions.

1

Creating and Managing Tables

Copyright © Oracle Corporation, 1998. All rights reserved.

ORACLE®

Schedule:	Timing	Topic
	30 minutes	Lecture
	20 minutes	Practice
	50 minutes	Total

Objectives

At the end of this lesson, you will be able to:

- **Describe the main database objects**
- **Create tables**
- **Describe the datatypes that can be used when specifying column definition**
- **Alter table definitions**
- **Drop, rename, and truncate tables**

1-252

Copyright © Oracle Corporation, 1998. All rights reserved.

ORACLE®

Lesson Aim

In this lesson, you will learn about main database objects and their relationships to each other. You will also learn how to create, alter, and drop tables.

Database Objects

Object	Description
Table	Basic unit of storage; composed of rows and columns
View	Logically represents subsets of data from one or more tables
Sequence	Generates primary key values
Index	Improves the performance of some queries
Synonym	Gives alternative names to objects

1-253

Copyright © Oracle Corporation, 1998. All rights reserved.

ORACLE®

Database Objects

An Oracle8 database can contain multiple data structures. Each structure should be outlined in the database design so that it can be created during the build stage of database development.

- Table: Stores data
- View: Subset of data from one or more tables
- Sequence: Generates primary key values
- Index: Improves the performance of some queries
- Synonym: Gives alternative names to objects

Oracle8 Table Structures

- Tables can be created at any time, even while users are using the database.
- You do not need to specify the size of any table. The size is ultimately defined by the amount of space allocated to the database as a whole. It is important, however, to estimate how much space a table will use over time.
- Table structure can be modified online.

Class Management Note

Tables can have up to 1000 columns and must conform to standard database object naming conventions.

Column definitions can be omitted when using the AS subquery clause. Tables are created without data unless a query is specified. Rows are usually added by using INSERT statements, SQL*Loader, or a form.

Naming Conventions

- **Must begin with a letter**
- **Can be 1–30 characters long**
- **Must contain only A–Z, a–z, 0–9, _, \$, and #**
- **Must not duplicate the name of another object owned by the same user**
- **Must not be an Oracle8 Server reserved word**

1-254

Copyright © Oracle Corporation, 1998. All rights reserved.

ORACLE®

Naming Rules

Name database tables and columns according to the standard rules for naming any Oracle8 database object.

- Table names and column names must begin with a letter and can be 1-30 characters long.
- Names must contain only the characters A-Z, a-z, 0-9, _ (underscore), \$, and # (legal characters, but their use is discouraged).
- Names must not duplicate the name of another object owned by the same Oracle8 Server user.
- Names must not be an Oracle8 reserved word.

Naming Guidelines

- Use descriptive names for tables and other database objects.
- Name the same entity consistently in different tables. For example, the department number column is called DEPTNO in both the EMP table and the DEPT table.

Note: Names are case insensitive. For example, EMP is treated as the same name as eMP or eMp.

For more information, see

Oracle8 Server SQL Reference, Release 8.0, “Object Names and Qualifiers.”



The CREATE TABLE Statement

- You must have :
 - CREATE TABLE privilege
 - A storage area

```
CREATE TABLE [schema.]table  
              (column datatype [DEFAULT expr];
```

- You specify:
 - Table name
 - Column name, column datatype, and column size

1-255

Copyright © Oracle Corporation, 1998. All rights reserved.

ORACLE®

The CREATE TABLE Statement

Create tables to store data by executing the SQL CREATE TABLE statement. This statement is one of the data definition language (DDL) statements, which you are covered in the next several lessons. DDL statements are a subset of SQL statements used to create, modify, or remove Oracle8 database structures. These statements have an immediate effect on the database, and they also record information in the data dictionary.

To create a table, a user must have the CREATE TABLE privilege and a storage area in which to create objects. The database administrator uses data control language (DCL) statements, which are covered in a later lesson, to grant privileges to users.

In the syntax:

<i>schema</i>	is the same as the owner's name
<i>table</i>	is the name of the table
DEFAULT <i>expr</i>	specifies a default value if a value is omitted in the INSERT statement
<i>column</i>	is the name of the column
<i>datatype</i>	is the column's datatype and length

For more information, see

Oracle8 Server SQL Reference, Release 8.0, "CREATE TABLE."



Referencing Another User's Tables

- **Tables belonging to other users are not in the user's schema.**
- **You should use the owner's name as a prefix to the table.**

Referencing Another User's Tables

A *schema* is a collection of objects. Schema objects are the logical structures that directly refer to the data in a database. Schema objects include tables, views, synonyms, sequences, stored procedures, indexes, clusters, and database links.

If a table does not belong to the user, the owner's name must be prefixed to the table.

The DEFAULT Option

- Specify a default value for a column during an insert.

```
... hiredate DATE DEFAULT SYSDATE, ...
```

- Legal values are literal value, expression, or SQL function.
- Illegal values are another column's name or pseudocolumn.
- The default datatype must match the column datatype.

1-257

Copyright © Oracle Corporation, 1998. All rights reserved.

ORACLE®

The DEFAULT Option

A column can be given a default value by using the DEFAULT option. This option prevents null values from entering the columns if a row is inserted without a value for the column. The default value can be a literal, an expression, or a SQL function, such as SYSDATE and USER, but the value cannot be the name of another column or a pseudocolumn, such as NEXTVAL or CURRVAL. The default expression must match the datatype of the column.

Creating Tables

- Create the table.

```
SQL> CREATE TABLE dept
  2      (deptno NUMBER(2),
  3      dname  VARCHAR2(14),
  4      loc    VARCHAR2(13));
Table created.
```

- Confirm table creation.

```
SQL> DESCRIBE dept
```

Name	Null?	Type
-----	-----	-----
DEPTNO	NOT NULL	NUMBER(2)
DNAME		VARCHAR2(14)
LOC		VARCHAR2(13)

Creating Tables

The example above creates the DEPT table, with three columns—namely, DEPTNO, DNAME, and LOC. It further confirms the creation of the table by issuing the DESCRIBE command.

Since creating a table is a DDL statement, an automatic commit takes place when this statement is executed.

Querying the Data Dictionary

- Describe tables owned by the user.

```
SQL> SELECT *  
2 FROM user_tables;
```

- View distinct object types owned by the user.

```
SQL> SELECT DISTINCT object_type  
2 FROM user_objects;
```

- View tables, views, synonyms, and sequences owned by the user.

```
SQL> SELECT *  
2 FROM user_catalog;
```

1-259

Copyright © Oracle Corporation, 1998. All rights reserved.

ORACLE®

Querying the Data Dictionary

You can query the data dictionary tables to view various database objects owned by you. The data dictionary tables frequently used are these:

- USER_TABLES
- USER_OBJECTS
- USER_CATALOG

Note: USER_CATALOG has a synonym called CAT. You can use this synonym instead of USER_CATALOG in SQL statements.

```
SQL> SELECT *  
2 FROM CAT;
```

Class Management Note

(For slide 10-10)

Oracle8 introduces large object (LOB) datatypes that can store large and unstructured data such as text, image, video, and spatial data, up to 4 gigabytes in size.

Datatypes

Datatype	Description
VARCHAR2(<i>size</i>)	Variable-length character data
CHAR(<i>size</i>)	Fixed-length character data
NUMBER(<i>p,s</i>)	Variable-length numeric data
DATE	Date and time values
LONG	Variable-length character data up to 2 gigabytes
CLOB	Single-byte character data up to 4 gigabytes
RAW and LONG RAW	Raw binary data
BLOB	Binary data up to 4 gigabytes
BFILE	Binary data stored in an external file; up to 4 gigabytes

1-260

Copyright © Oracle Corporation, 1998. All rights reserved.

ORACLE®

Datatypes

Datatype	Description
VARCHAR2(<i>size</i>)	Variable-length character data. A maximum <i>size</i> must be specified. Default and minimum <i>size</i> is 1, maximum <i>size</i> is 4000.
CHAR(<i>size</i>)	Fixed-length character data of length <i>size</i> bytes. Default and minimum <i>size</i> is 1, maximum <i>size</i> is 2000.
NUMBER(<i>p,s</i>)	Number having precision <i>p</i> and scale <i>s</i> ; the precision is the total number of decimal digits, and the scale is the number of digits to the right of the decimal point. The precision can range from 1 to 38 and the scale can range from -84 to 127.
DATE	Date and time values between January 1, 4712 B.C., and December 31, 9999 A.D.
LONG	Variable-length character data up to 2 gigabytes.
CLOB	Single-byte character data up to 4 gigabytes.
RAW(<i>size</i>)	Raw binary data of length <i>size</i> . Maximum <i>size</i> is 2000. A maximum <i>size</i> must be specified.
LONG RAW	Raw binary data of variable length up to 2 gigabytes.
BLOB	Binary data up to 4 gigabytes.
BFILE	Binary data stored in an external file; up to 4 gigabytes.

Creating a Table by Using a Subquery

- **Create a table and insert rows by combining the CREATE TABLE statement and *AS subquery* option.**

```
CREATE TABLE table  
      [column(, column...)]  
AS subquery;
```

- **Match the number of specified columns to the number of subquery columns.**
- **Define columns with column names and default values.**

1-261

Copyright © Oracle Corporation, 1998. All rights reserved.

ORACLE®

Creating a Table from Rows in Another Table

A second method to create a table is to apply the *AS subquery* clause to both create the table and insert rows returned from the subquery.

In the syntax:

<i>table</i>	is the name of the table
<i>column</i>	is the name of the column, default value, and integrity constraint
<i>subquery</i> new table	is the SELECT statement that defines the set of rows to be inserted into the

Guidelines

- The table will be created with the specified column names, and the rows retrieved by the SELECT statement will be inserted into the table.
- The column definition can contain only the column name and default value.
- If column specifications are given, the number of columns must equal the number of columns in the subquery SELECT list.
- If no column specifications are given, the column names of the table are the same as the column names in the subquery.

Creating a Table by Using a Subquery

```
SQL> CREATE TABLE dept30
2 AS
3 SELECT empno, ename, sal*12 ANNSAL, hiredate
4 FROM emp
5 WHERE deptno = 30;
Table created.
```

```
SQL> DESCRIBE dept30
```

Name	Null?	Type
-----	-----	-----
EMPNO	NOT NULL	NUMBER(4)
ENAME		VARCHAR2(10)
ANNSAL		NUMBER
HIREDATE		DATE

1-262

Copyright © Oracle Corporation, 1998. All rights reserved.

ORACLE®

Creating a Table from Rows in Another Table

The example above creates a table, DEPT30, that contains details of all the employees working in department 30. Notice that the data for the DEPT30 table is coming from the EMP table.

You can verify the existence of a database table and check column definitions by using the SQL*Plus DESCRIBE command.

Give a column alias, when selecting an expression.



The ALTER TABLE Statement

Use the ALTER TABLE statement to:

- Add a new column
- Modify an existing column
- Define a default value for the new column

```
ALTER TABLE table
ADD          (column datatype [DEFAULT expr]
              [, column datatype]...);
```

```
ALTER TABLE table
MODIFY       (column datatype [DEFAULT expr]
              [, column datatype]...);
```

1-263

Copyright © Oracle Corporation, 1998. All rights reserved.

ORACLE®

ALTER TABLE Statement

After you create your tables, you may need to change the table structures because you omitted a column or your column definition needs to be changed. You can do this by using the ALTER TABLE statement.

You can add columns to a table by using the ALTER TABLE statement with the ADD clause.

In the syntax:

<i>table</i>	is the name of the table
<i>column</i>	is the name of the new column
<i>datatype</i>	is the datatype and length of the new column
DEFAULT <i>expr</i>	specifies the default value for a new column

You can modify existing columns in a table by using the ALTER TABLE statement with the MODIFY clause.

Note: The slide gives the abridged syntax for ALTER TABLE. More about ALTER TABLE is covered in lesson 11.

Adding a Column

DEPT30

EMPNO	ENAME	ANNSAL	HIREDATE	New column JOB
7698	BLAKE	34200	01-MAY-81	
7654	MARTIN	15000	28-SEP-81	
7499	ALLEN	19200	20-FEB-81	
7844	TURNER	18000	08-SEP-81	
...				

“...add a new column into DEPT30 table...”



DEPT30

EMPNO	ENAME	ANNSAL	HIREDATE	JOB
7698	BLAKE	34200	01-MAY-81	
7654	MARTIN	15000	28-SEP-81	
7499	ALLEN	19200	20-FEB-81	
7844	TURNER	18000	08-SEP-81	
...				

Adding a Column

The graphic adds the JOB column to DEPT30 table. Notice that the new column becomes the last column in the table.

Adding a Column

- You use the ADD clause to add columns.

```
SQL> ALTER TABLE dept30
      2 ADD          (job VARCHAR2(9));
Table altered.
```

- The new column becomes the last column.

EMPNO	ENAME	ANNSAL	HIREDATE	JOB
7698	BLAKE	34200	01-MAY-81	
7654	MARTIN	15000	28-SEP-81	
7499	ALLEN	19200	20-FEB-81	
7844	TURNER	18000	08-SEP-81	
...				

6 rows selected.

1-265

Copyright © Oracle Corporation, 1998. All rights reserved.

ORACLE®

Guidelines for Adding a Column

- You can add or modify columns, but you cannot drop them from a table.
- You cannot specify where the column is to appear. The new column becomes the last column.

The example above adds a column named JOB to the DEPT30 table. The JOB column becomes the last column in the table.

Note: If a table already contains rows when a column is added, then the new column is initially null for all the rows.

You can define a NOT NULL column only if the table contains no rows because data cannot be specified for existing rows at the same time that the column is added.

Modifying a Column

- You can change a column's datatype, size, and default value.

```
ALTER TABLE dept30
MODIFY      (ename VARCHAR2(15));
Table altered.
```

- A change to the default value affects only subsequent insertions to the table.

Modifying a Column

You can modify a column definition by using the ALTER TABLE statement with the MODIFY clause. Column modification can include changes to a column's datatype, size, and default value.

Guidelines

- Increase the width or precision of a numeric column.
- Decrease the width of a column if the column contains only null values or if the table has no rows.
- Change the datatype if the column contains null values.
- Convert a CHAR column to the VARCHAR2 datatype or convert a VARCHAR2 column to the CHAR datatype if the column contains null values or if you do not change the size.
- A change to the default value of a column affects only subsequent insertions to the table.

Dropping a Table

- All data in the table is deleted.
- Any pending transactions are committed.
- All indexes are dropped.
- You *cannot* roll back this statement.

```
SQL> DROP TABLE dept30;  
Table dropped.
```

1-267

Copyright © Oracle Corporation, 1998. All rights reserved.

ORACLE®

Dropping a Table

The DROP TABLE statement removes the definition of an Oracle8 table. When you drop a table, the database loses all the data in the table and all the indexes associated with it.

Syntax

```
DROP TABLE table;
```

where: *table* is the name of the table

Guidelines

- All data is deleted from the table.
- Any views, synonyms, stored procedures, functions, and packages will remain but are invalid.
- Any pending transactions are committed.
- Only the creator of the table or a user with the DROP ANY TABLE privilege can remove a table.

The DROP TABLE statement, once executed, is irreversible. The Oracle8 Server does not question the action when you issue the DROP TABLE statement. If you own that table or have a high-level privilege, then the table is immediately removed. All DDL statements issue a commit, therefore making the transaction permanent.



Changing the Name of an Object

- To change the name of a table, view, sequence, or synonym, you execute the **RENAME** statement.

```
SQL> RENAME dept TO department;  
Table renamed.
```

- You must be the owner of the object.

1-268

Copyright © Oracle Corporation, 1998. All rights reserved.

ORACLE®

Renaming a Table

Additional DDL statements include the RENAME statement, which is used to rename a table, view, sequence, or a synonym.

Syntax

```
RENAME old_name TO new_name;
```

where: *old_name* is the old name of the table, view, sequence, or synonym

new_name is the new name of the table, view, sequence, or synonym

You must be the owner of the object you rename.



Truncating a Table

- The TRUNCATE TABLE statement:
 - Removes all rows from a table
 - Releases the storage space used by that table

```
SQL> TRUNCATE TABLE department;  
Table truncated.
```

- Cannot roll back row removal when using TRUNCATE
- Alternatively, remove rows by using the DELETE statement

1-269

Copyright © Oracle Corporation, 1998. All rights reserved.

ORACLE®

Truncating a Table

Another DDL statement is the TRUNCATE TABLE statement, which is used to remove all rows from a table and to release the storage space used by that table. When using the TRUNCATE TABLE statement, you cannot rollback row removal.

Syntax

```
TRUNCATE TABLE table;
```

where: *table* is the name of the table

You must be the owner of the table or have DELETE TABLE system privileges to truncate a table.

The DELETE statement can also remove all rows from a table, but it does not release storage space.



Adding Comments to a Table

- You can add comments to a table or column by using the **COMMENT** statement.

```
SQL> COMMENT ON TABLE emp
      2 IS 'Employee Information';
Comment created.
```

- Comments can be viewed through the data dictionary views.
 - ALL_COL_COMMENTS
 - USER_COL_COMMENTS
 - ALL_TAB_COMMENTS
 - USER_TAB_COMMENTS

1-270

Copyright © Oracle Corporation, 1998. All rights reserved.

ORACLE®

Adding a Comment to a Table

You can add a comment of up to 2000 bytes about a column, table, view, or snapshot by using the **COMMENT** statement. The comment is stored in the data dictionary and can be viewed in one of the following data dictionary views in the **COMMENTS** column:

- ALL_COL_COMMENTS
- USER_COL_COMMENTS
- ALL_TAB_COMMENTS
- USER_TAB_COMMENTS

Syntax

where: **COMMENT ON TABLE** *table* | **COLUMN** *table.column*
table is the name of the table
column IS 'text'; *text* is the name of the column in a table
text is the text of the comment

You can drop a comment from the database by setting it to empty string (' ').

```
SQL> COMMENT ON TABLE emp IS ' ';
```


Summary

Statement	Description
CREATE TABLE	Creates a table
ALTER TABLE	Modifies table structures
DROP TABLE	Removes the rows and table structure
RENAME	Changes the name of a table, view, sequence, or synonym
TRUNCATE	Removes all rows from a table and releases the storage space
COMMENT	Adds comments to a table or view

1-271

Copyright © Oracle Corporation, 1998. All rights reserved.

ORACLE®

CREATE TABLE

- You can create a table.
- Create a table based on another table by using a subquery.

ALTER TABLE

- Modify table structures.
- Change column widths, change column datatypes, and add columns.

DROP TABLE

- Remove rows and a table structure.
- Once executed, this statement cannot be rolled back.

RENAME

- Rename a table, view, sequence, or synonym.

TRUNCATE

- Remove all rows from a table and release the storage space used by the table.
- DELETE statement only removes rows.

COMMENT

- Add a comment to a table or a column.
- Query the data dictionary to view the comment.

Practice Overview

- **Creating new tables**
- **Creating a new table by using the CREATE TABLE AS syntax**
- **Modifying column definitions**
- **Verifying that the tables exist**
- **Adding comments to a tables**
- **Dropping tables**
- **Altering tables**

1-272

Copyright © Oracle Corporation, 1998. All rights reserved.

ORACLE®

Practice Overview

Create new tables containing constraints by using the CREATE TABLE statement. Confirm that the new table was added to the database. Create the syntax in the command file, and then execute the command file to create the table.